

Deep Learning Applications for collider physics

Lecture 2

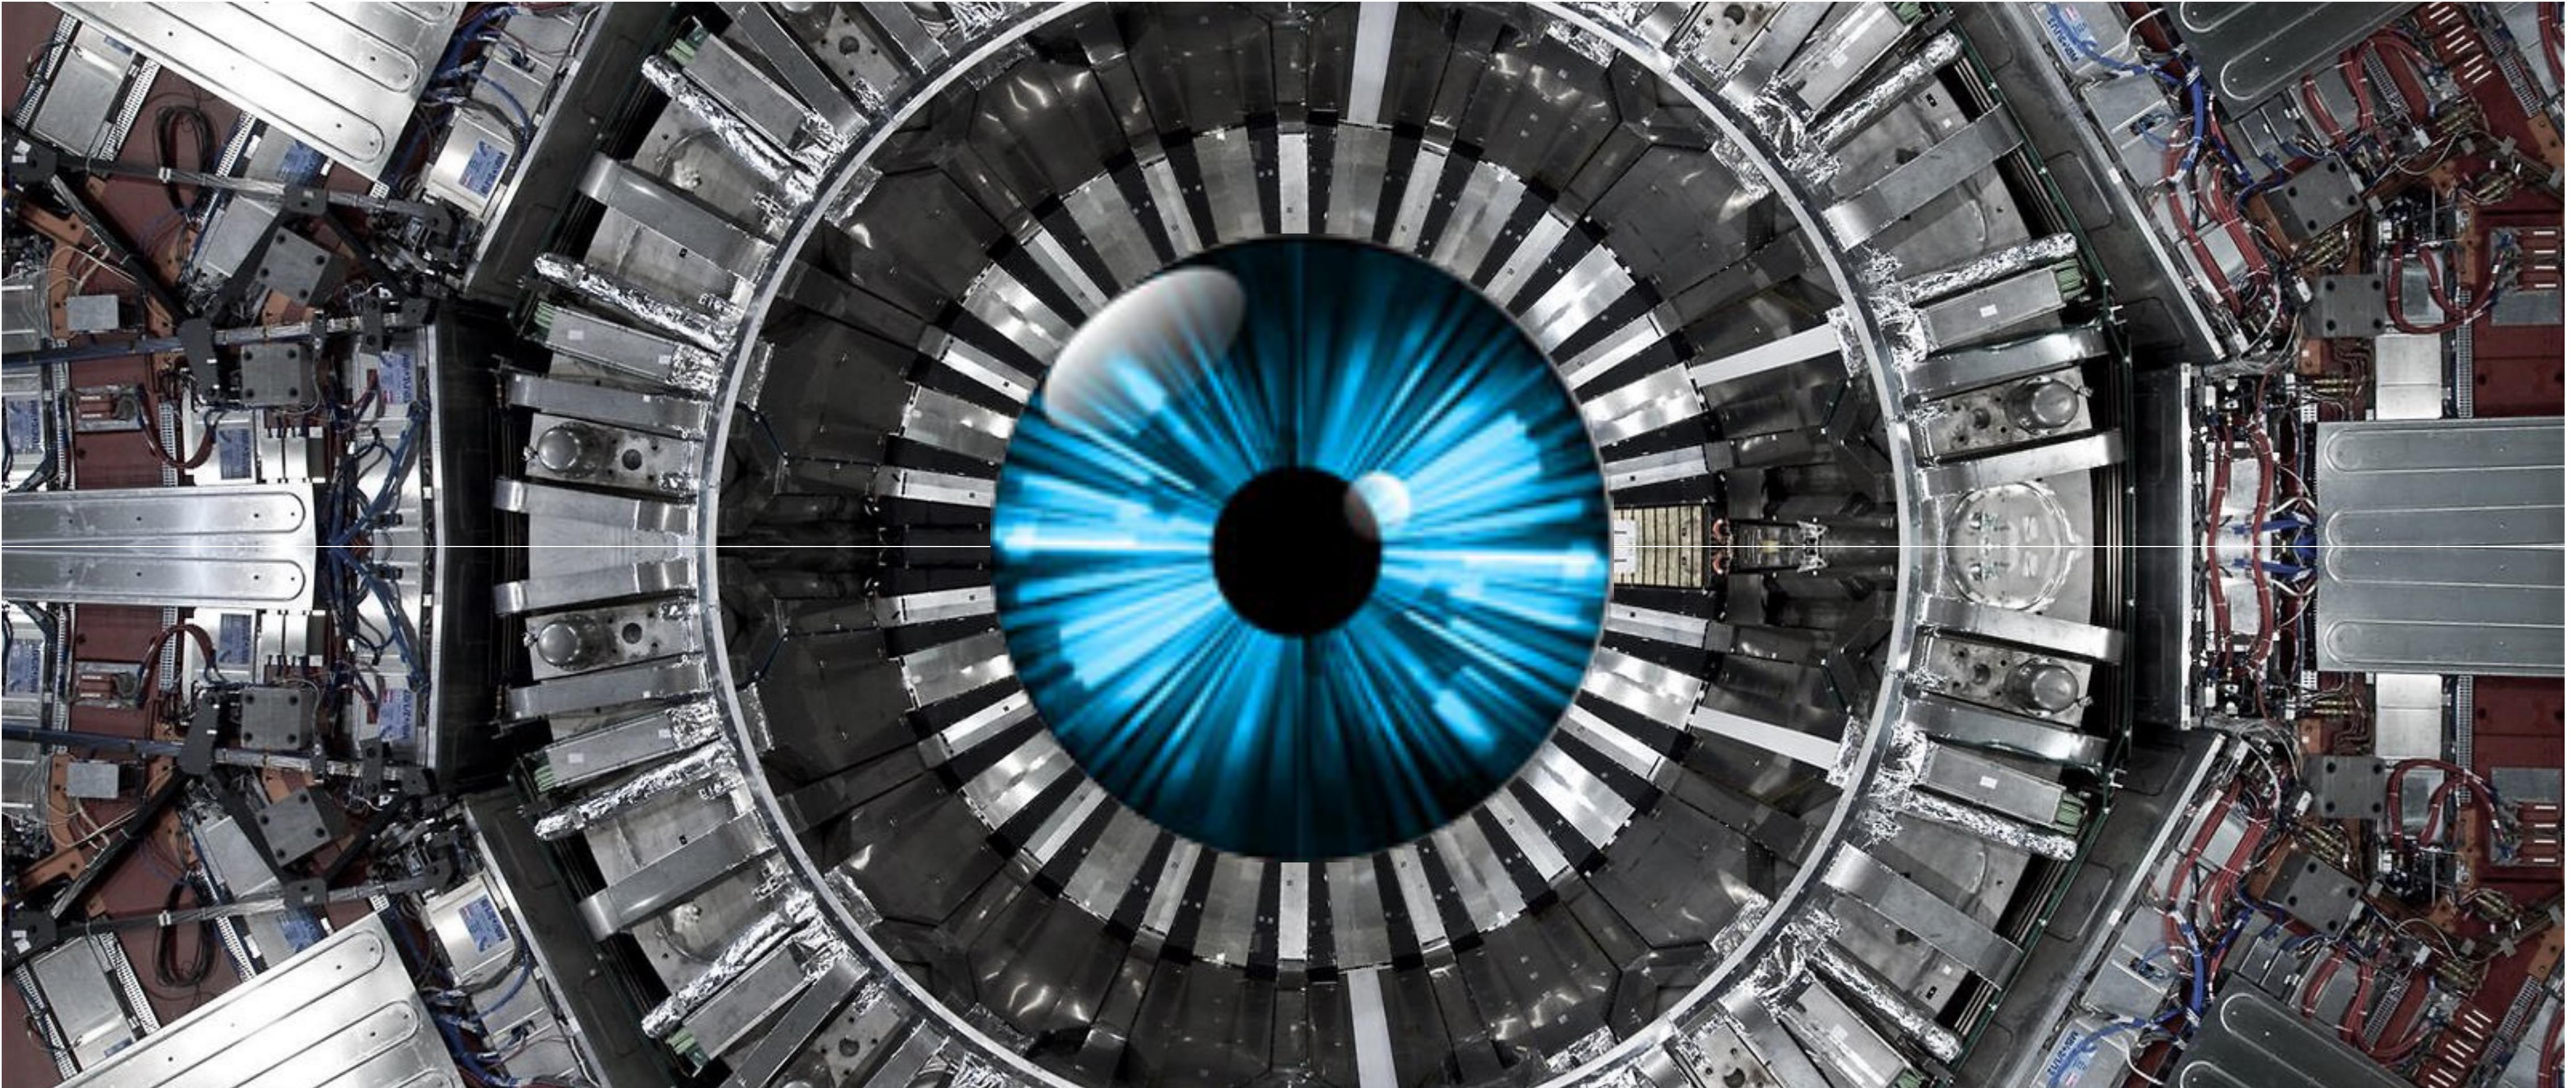
Maurizio Pierini





Plan for this week

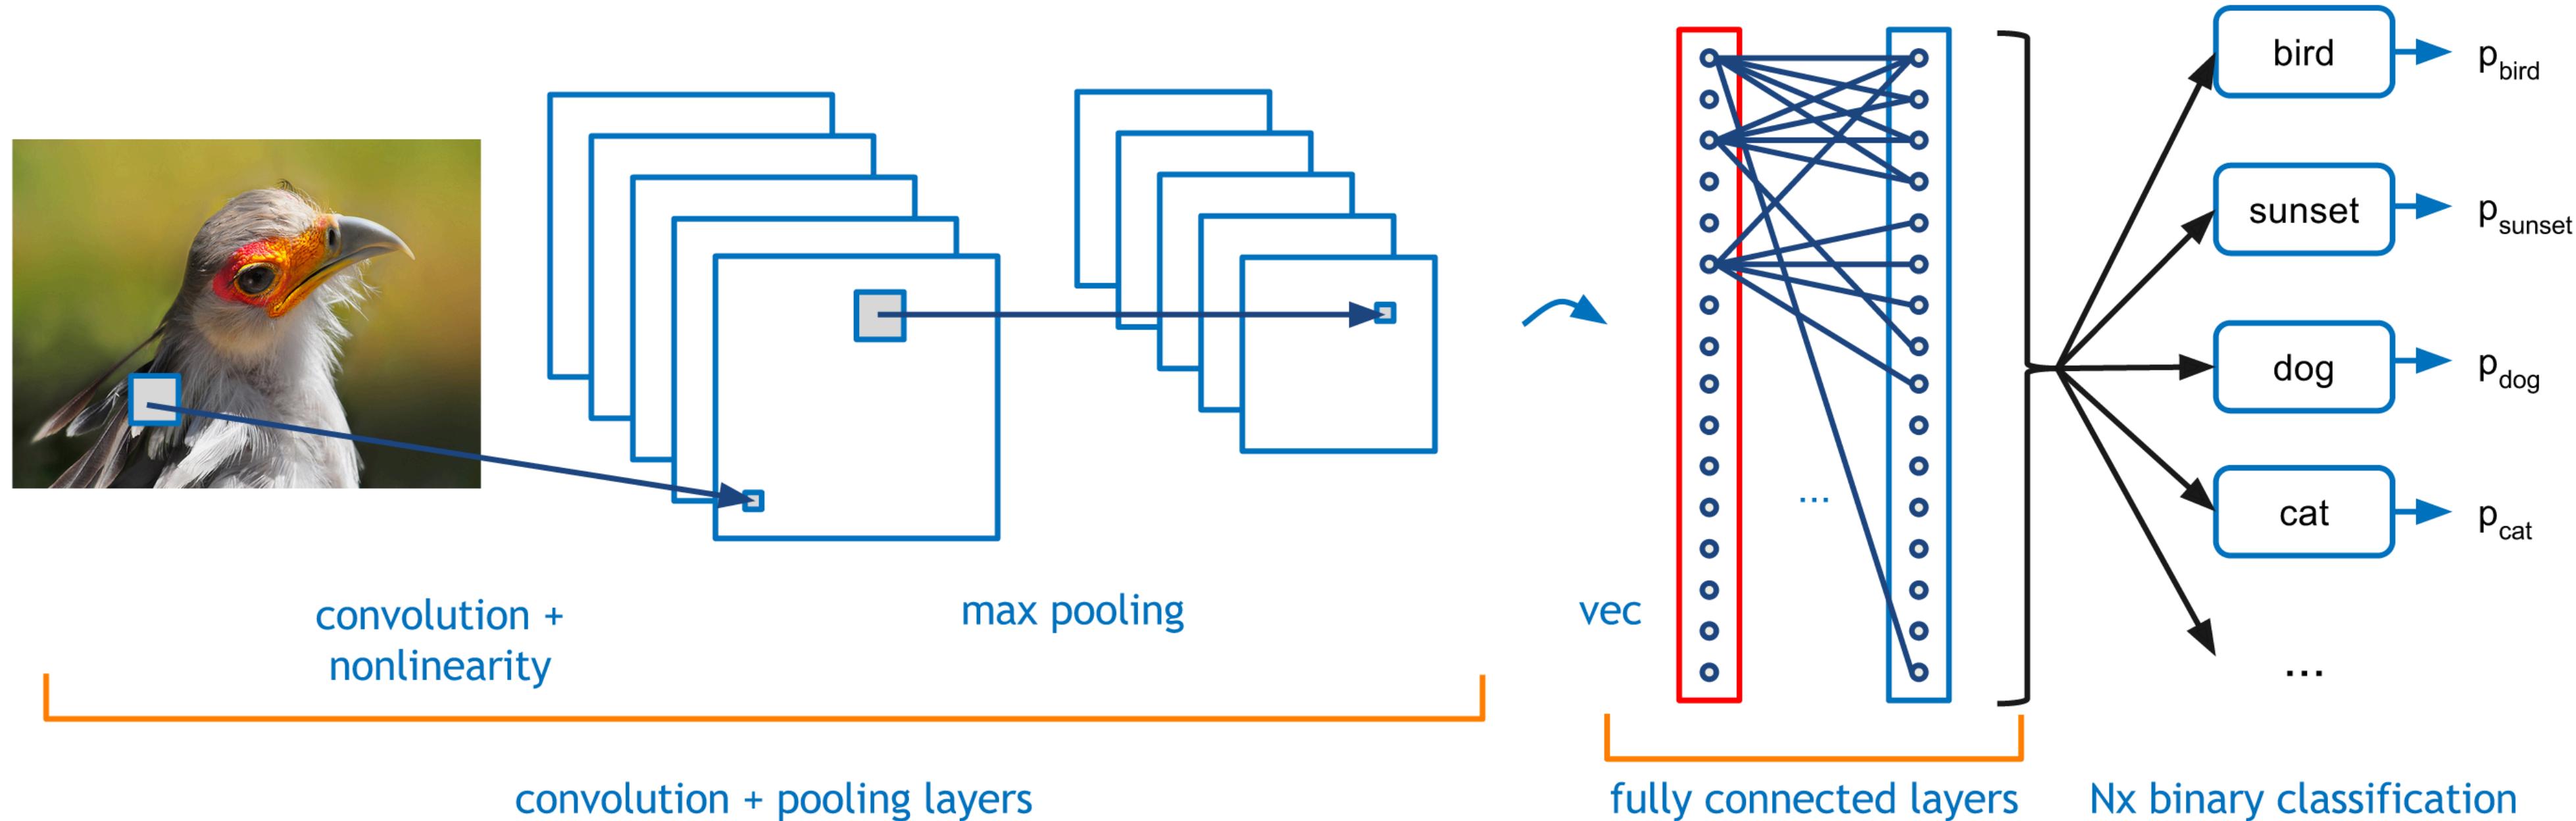
| | Day1 | Day2 | Day3 | Day4 | Day5 |
|----------|----------------------------|-------------------|-----------------|-------------------|-----------------------|
| Lecture | Introduction | ConvNN | RNNs | Graphs | Unsupervised Learning |
| Tutorial | Fully Connected Classifier | ConvNN Classifier | RNNs Classifier | Graphs Classifier | Anomaly Detection |



Particle Reconstruction & Computer Vision

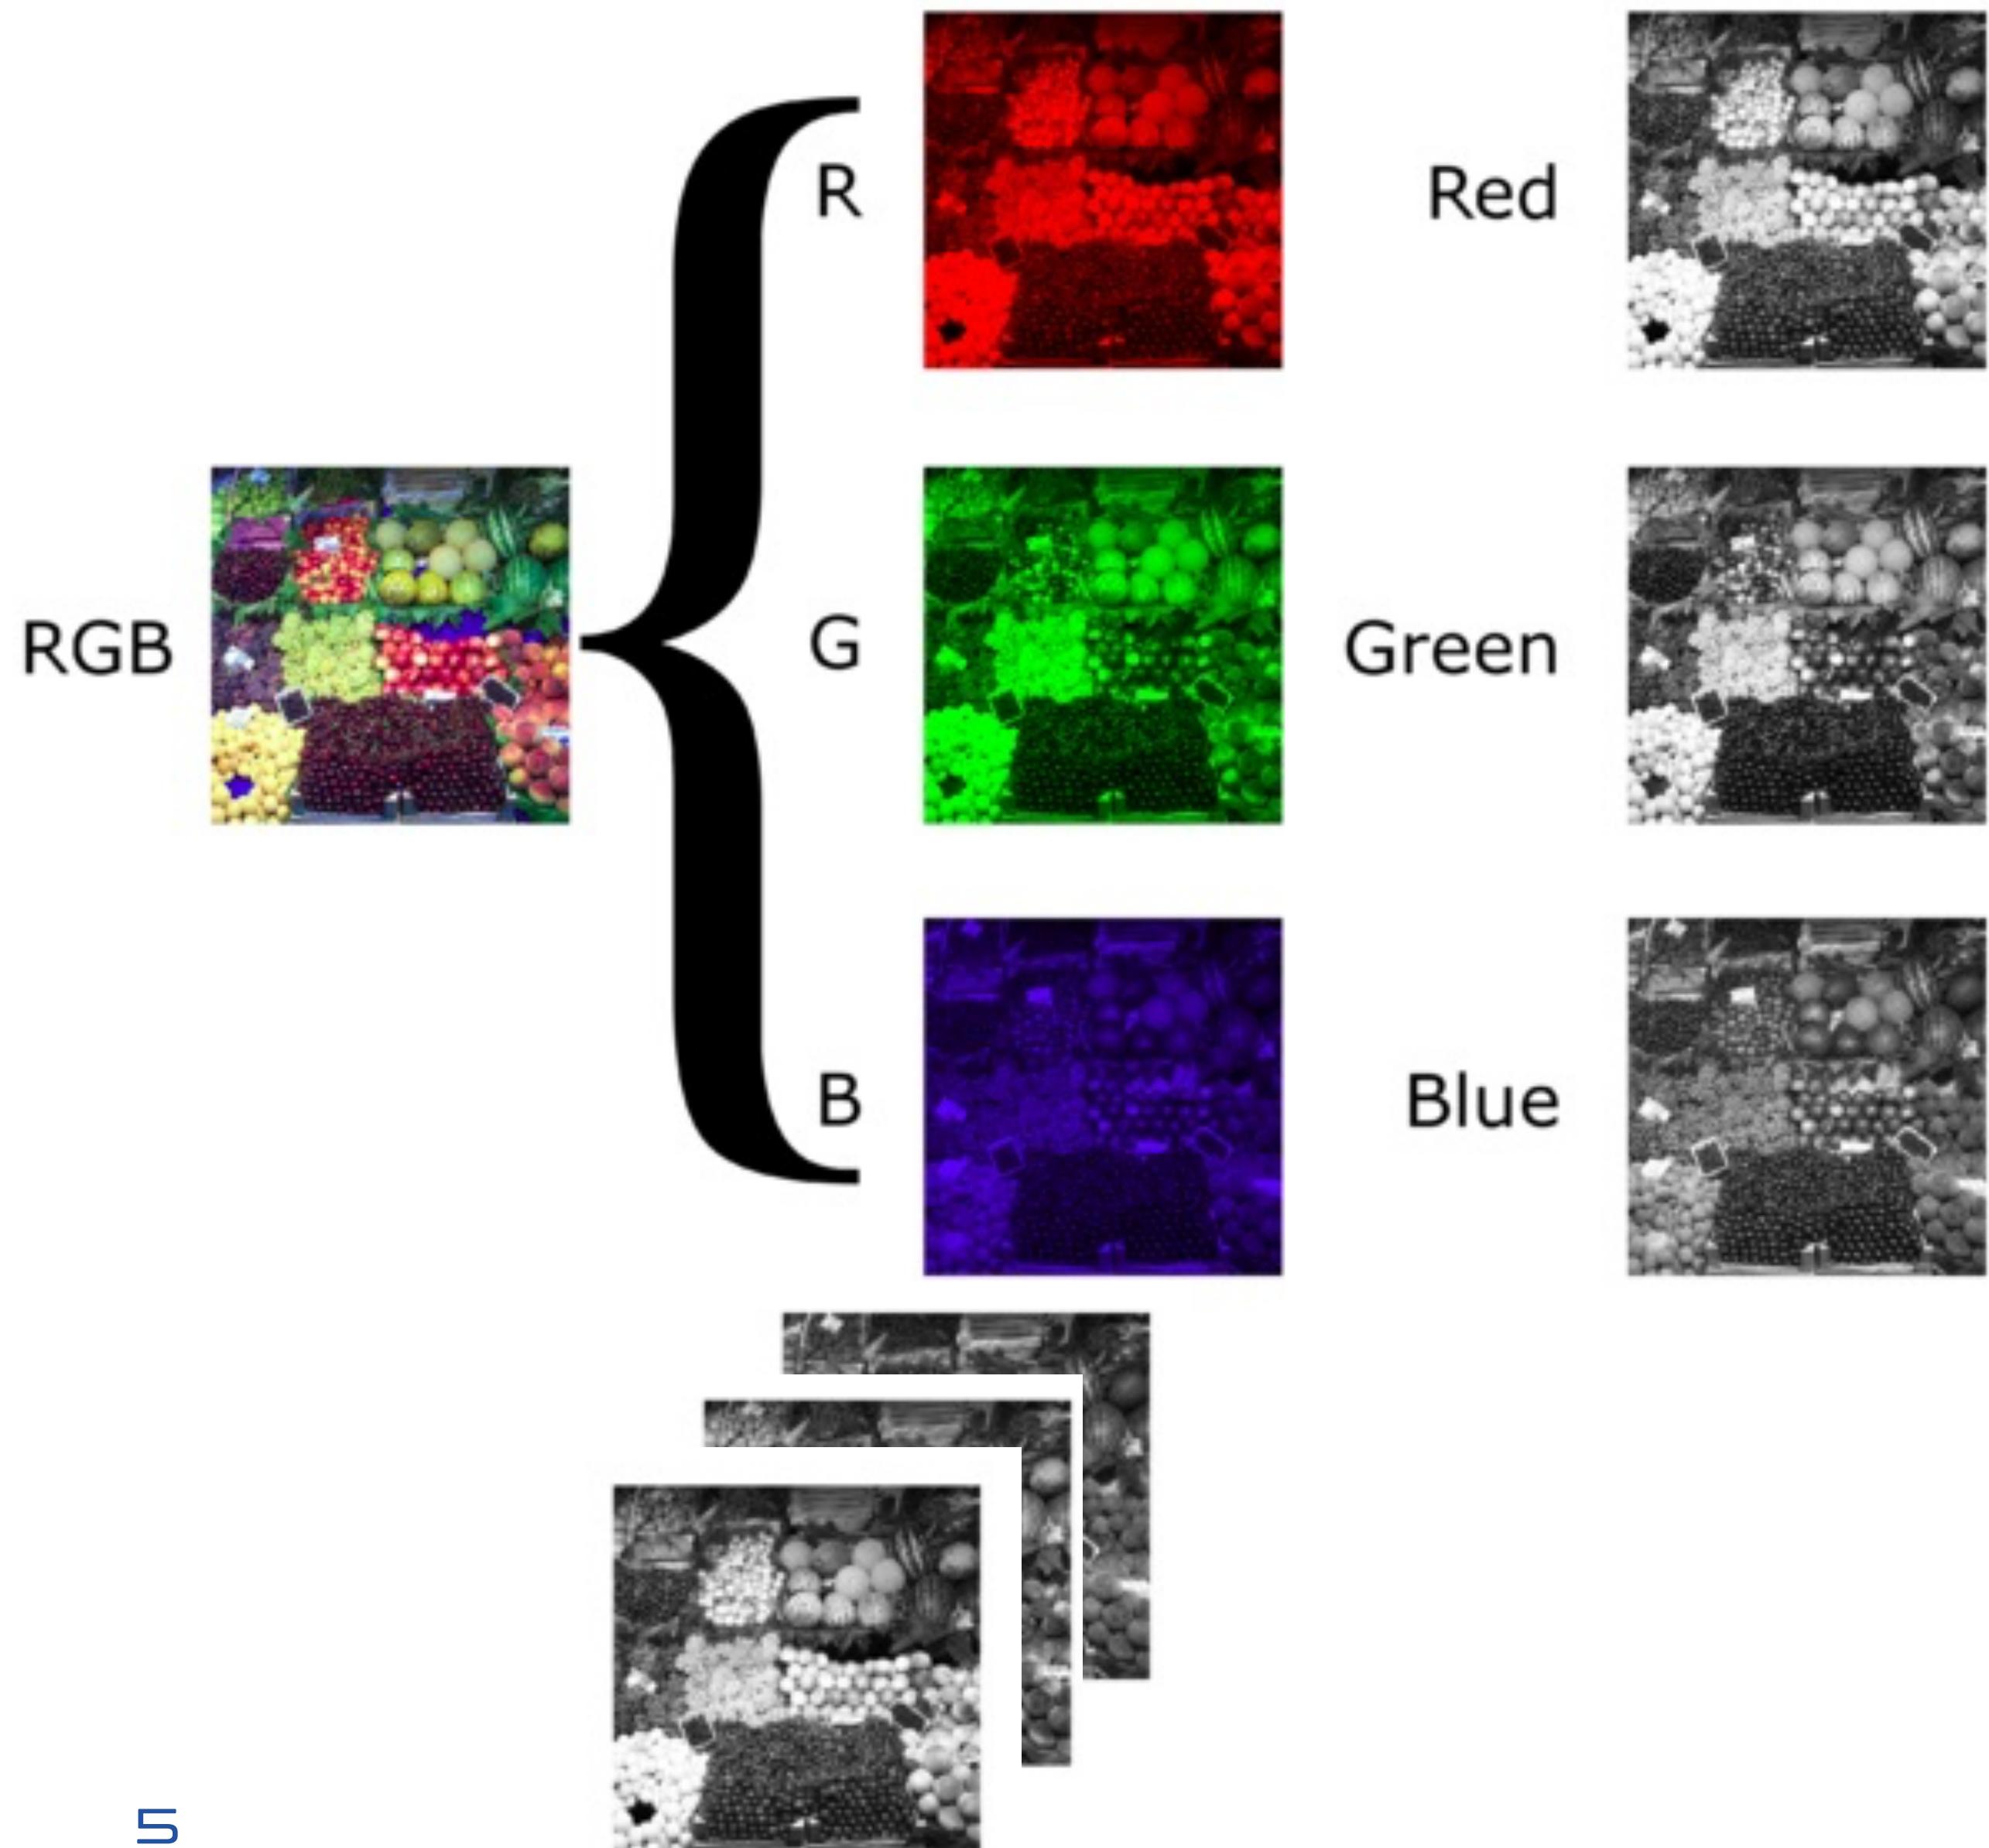
Convolutional Layer

- Special architectures read the raw information (e.g., images) and convert them into “smart variables” (high-level features) to accomplish the task
- Typical example: convolutional neural networks for image processing & computing vision



Digital images

- *Each image is a matrix of pixels*
- *Each pixel comes with a color*
- *In RGB scheme, these are three values $i \in [0,1]$*
- *Each image becomes a 3D tensor*
- *3 channels of 2D pixelated images*



Convolutional filter

- The main ingredient of ConvNN is a filter, a $k \times k'$ matrix of weights
- The filter scans the image and performs a scalar product of each image patch
- This results into a new matrix of values, with different dimensionality

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 3 | 5 | 6 | 2 | 4 | 5 |
| 7 | 4 | 7 | 3 | 6 | 3 | 4 |
| 9 | 1 | 2 | 1 | 9 | 6 | 0 |
| 9 | 2 | 1 | 1 | 7 | 3 | 5 |
| 8 | 0 | 4 | 7 | 6 | 8 | 0 |
| 8 | 3 | 4 | 5 | 5 | 3 | 4 |
| 7 | 9 | 4 | 6 | 5 | 2 | 6 |

| | | |
|----|----|----|
| 4 | -1 | 4 |
| -2 | 2 | -5 |
| 3 | 1 | -6 |

$$\begin{aligned} & 0 \times 4 - 3 \times 1 + 5 \times 4 + \\ & -7 \times 2 + 4 \times 2 - 7 \times 5 + \\ & 9 \times 3 + 1 \times 1 - 2 \times 6 = -8 \end{aligned}$$



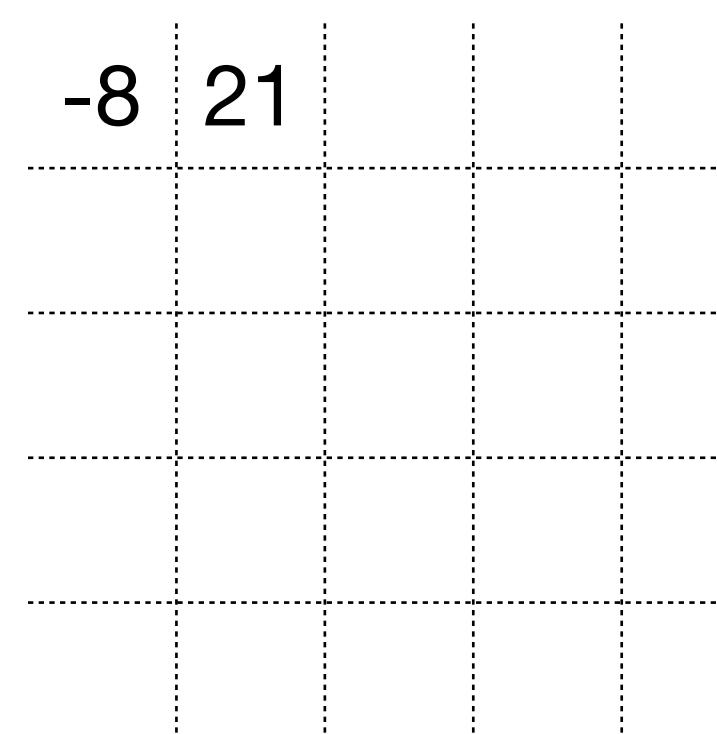
Filtro convolutional

- The main ingredient of ConvNN is a filter, a $k \times k'$ matrix of weights
- The filter scans the image and performs a scalar product of each image patch
- The scan is done shifting the filter by a stride (of q , 2, ... cells)
- This results into a new matrix of values, with different dimensionality

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 3 | 5 | 6 | 2 | 4 | 5 |
| 7 | 4 | 7 | 3 | 6 | 3 | 4 |
| 9 | 1 | 2 | 1 | 9 | 6 | 0 |
| 9 | 2 | 1 | 1 | 7 | 3 | 5 |
| 8 | 0 | 4 | 7 | 6 | 8 | 0 |
| 8 | 3 | 4 | 5 | 5 | 3 | 4 |
| 7 | 9 | 4 | 6 | 5 | 2 | 6 |

| | | |
|----|----|----|
| 4 | -1 | 4 |
| -2 | 2 | -5 |
| 3 | 1 | -6 |

$$\begin{aligned}
 & 3 \times 4 - 5 \times 1 + 6 \times 4 + \\
 & -4 \times 2 + 7 \times 2 - 3 \times 5 + \\
 & 1 \times 3 + 2 \times 1 - 1 \times 6 = 21
 \end{aligned}$$



Pooling

- *MaxPooling: Given an image and a filter of size $k \times k'$, scans the image and replaces each $k \times k'$ patch with its maximum*

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 3 | 5 | 6 | 2 | 4 | 5 |
| 7 | 4 | 7 | 3 | 6 | 3 | 4 |
| 9 | 1 | 2 | 1 | 9 | 6 | 0 |
| 9 | 2 | 1 | 1 | 7 | 3 | 5 |
| 8 | 0 | 4 | 7 | 6 | 8 | 0 |
| 8 | 3 | 4 | 5 | 5 | 3 | 4 |
| 7 | 9 | 4 | 6 | 5 | 2 | 6 |



| | | | | |
|---|---|---|---|---|
| 9 | 7 | 9 | 9 | 9 |
| 9 | 7 | 9 | 9 | 9 |
| 9 | 7 | 7 | 9 | 9 |
| 9 | 7 | 7 | 8 | 8 |
| 9 | 9 | 7 | 8 | 8 |

- *AveragePooling: Given an image and a filter of size $k \times k'$, scans the image and replaces each $k \times k'$ patch with its average*

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 3 | 5 | 6 | 2 | 4 | 5 |
| 7 | 4 | 7 | 3 | 6 | 3 | 4 |
| 9 | 1 | 2 | 1 | 9 | 6 | 0 |
| 9 | 2 | 1 | 1 | 7 | 3 | 5 |
| 8 | 0 | 4 | 7 | 6 | 8 | 0 |
| 8 | 3 | 4 | 5 | 5 | 3 | 4 |
| 7 | 9 | 4 | 6 | 5 | 2 | 6 |



| | | | | |
|-----|--|--|--|--|
| 4.2 | | | | |
| | | | | |
| | | | | |

| | | | | |
|--|--|--|-----|--|
| | | | 5.0 | |
| | | | | |
| | | | | |

- ...

Padding

- When the filter arrived at the edge, it might exceeds it (if n/k is not an integer)

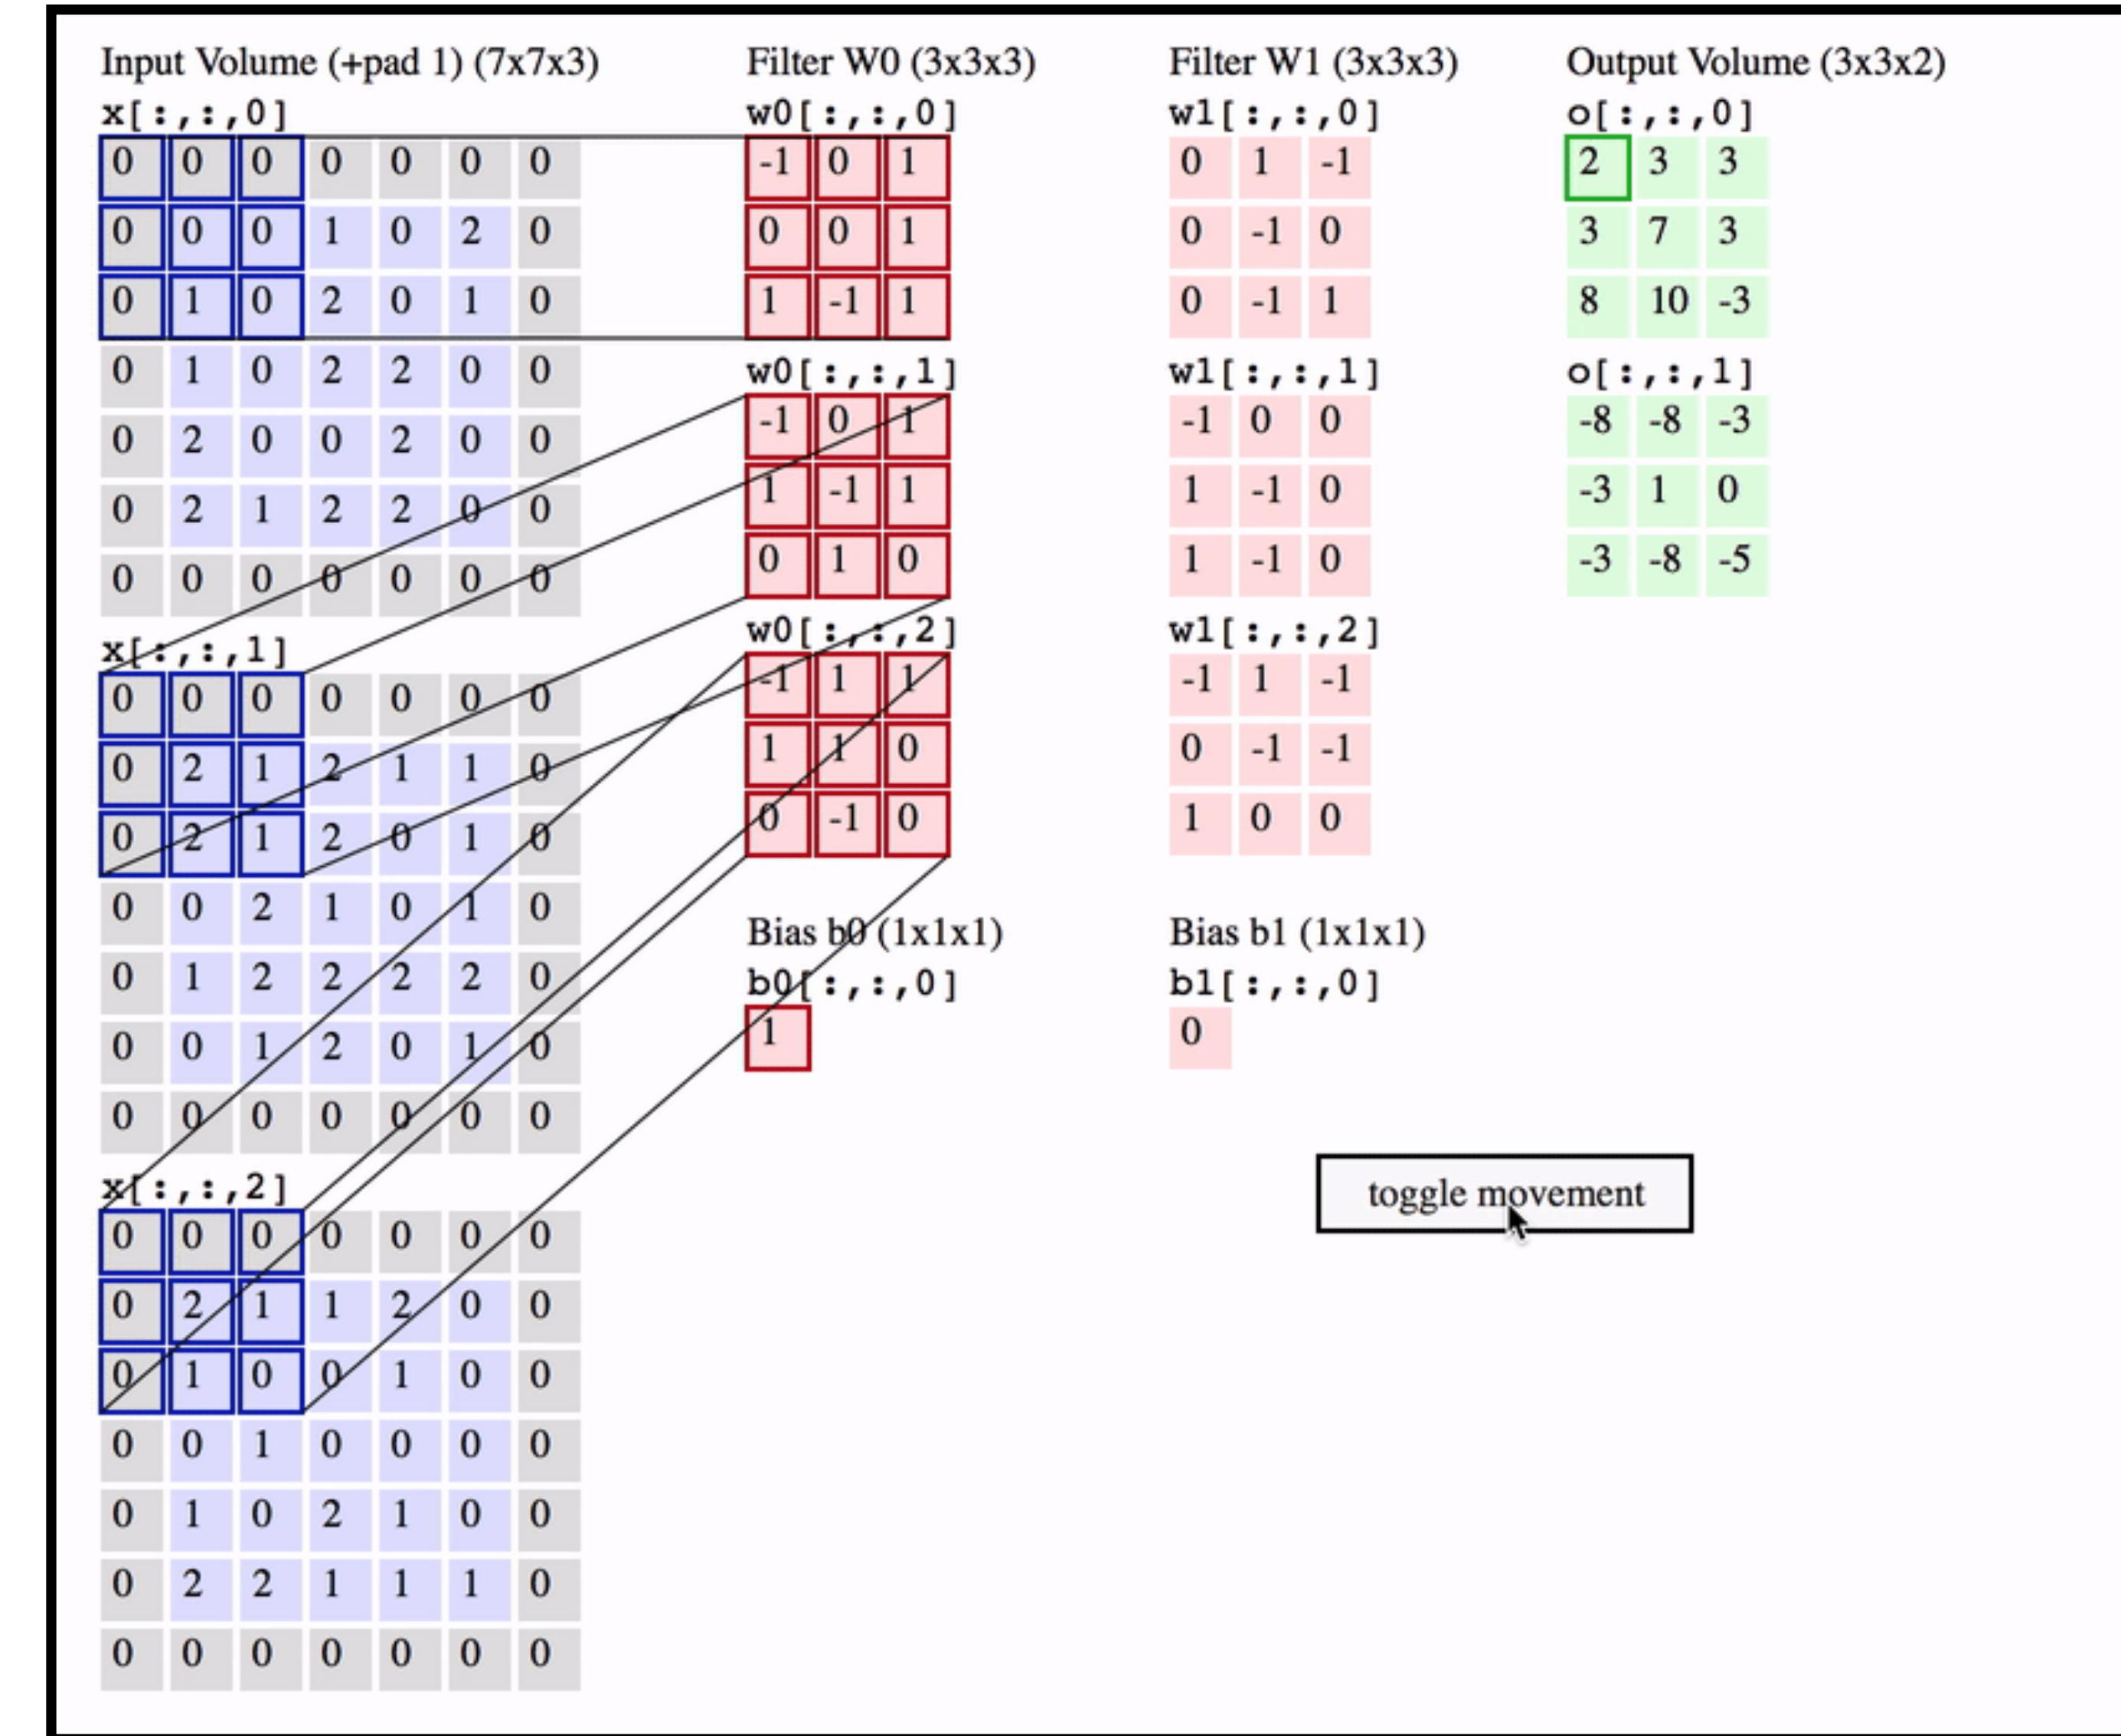
| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 3 | 5 | 6 | 2 | 4 | 4 |
| 7 | 4 | 7 | 3 | 6 | 3 | 3 |
| 9 | 1 | 2 | 1 | 9 | 6 | 6 |
| 9 | 2 | 1 | 1 | 7 | 3 | 3 |
| 8 | 0 | 4 | 7 | 6 | 8 | 8 |
| 8 | 3 | 4 | 5 | 5 | 3 | 3 |
| 8 | 3 | 4 | 5 | 5 | 3 | 3 |

- In this case, a padding rule needs to be specified

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 3 | 5 | 6 | 2 | 4 | 0 |
| 7 | 4 | 7 | 3 | 6 | 3 | 0 |
| 9 | 1 | 2 | 1 | 9 | 6 | 0 |
| 9 | 2 | 1 | 1 | 7 | 3 | 0 |
| 8 | 0 | 4 | 7 | 6 | 8 | 0 |
| 8 | 3 | 4 | 5 | 5 | 3 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

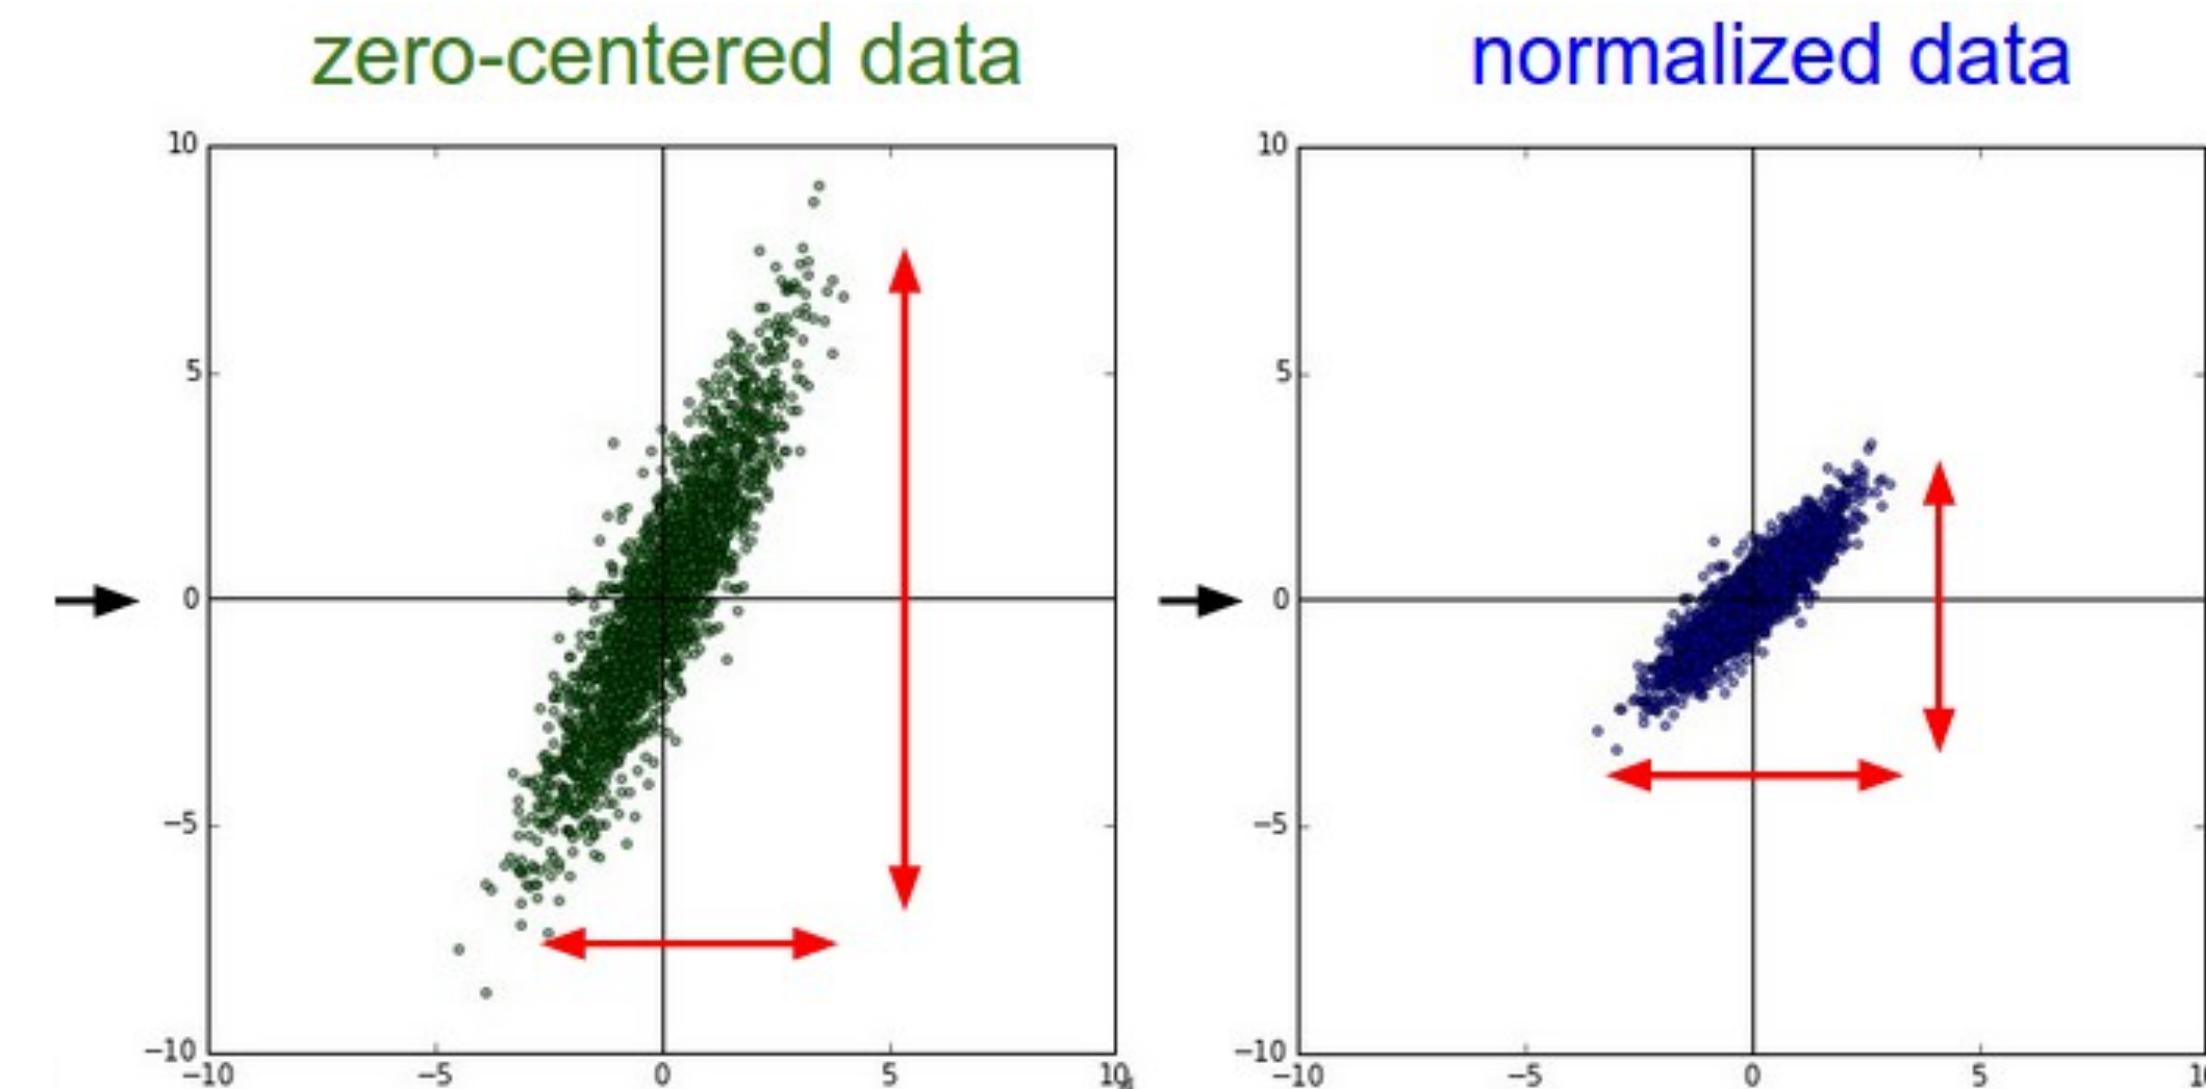
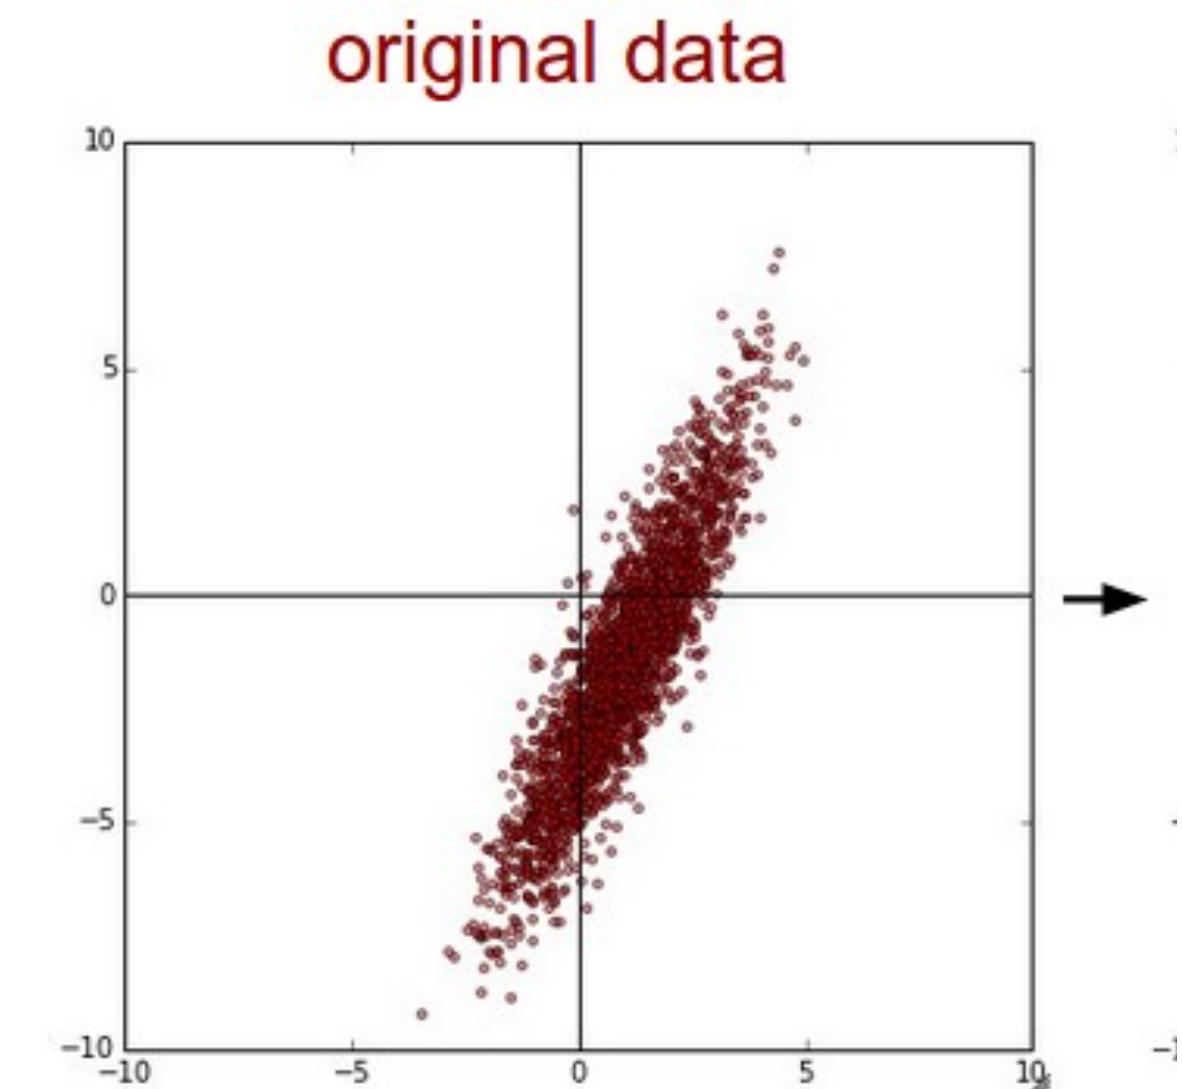
- Same: repeat the values at the boundary
- Zero: fill the extra columns with zeros

Convolutional Layer



BatchNormalization Layer

- It is good practice to give normalized inputs to a layer
- With all inputs having the same order of magnitude, all weights are equal important in the gradient
- Prevents explosion of the loss function
- This can be done automatically with BatchNormalization
- non-learnable shift and scale parameters, adjusted batch by batch



more complex structures

- Dense NN architectures can be made more complex

8|0|4|7|6|8|0 7|9|4|6|5|2|6 8|3|4|5|5|3|4



Concatenation

8|0|4|7|6|8|0|8|3|4|5|5|3|4|7|9|4|6|5|2|6

- Multiple inputs
- Multiple outputs
- Different networks branches

- This is possible thanks to layer-manipulation layers

- Add, Subtract, etc.

8|0|4|7|6|8|0
8|3|4|5|5|3|4
7|9|4|6|5|2|6

Flattening



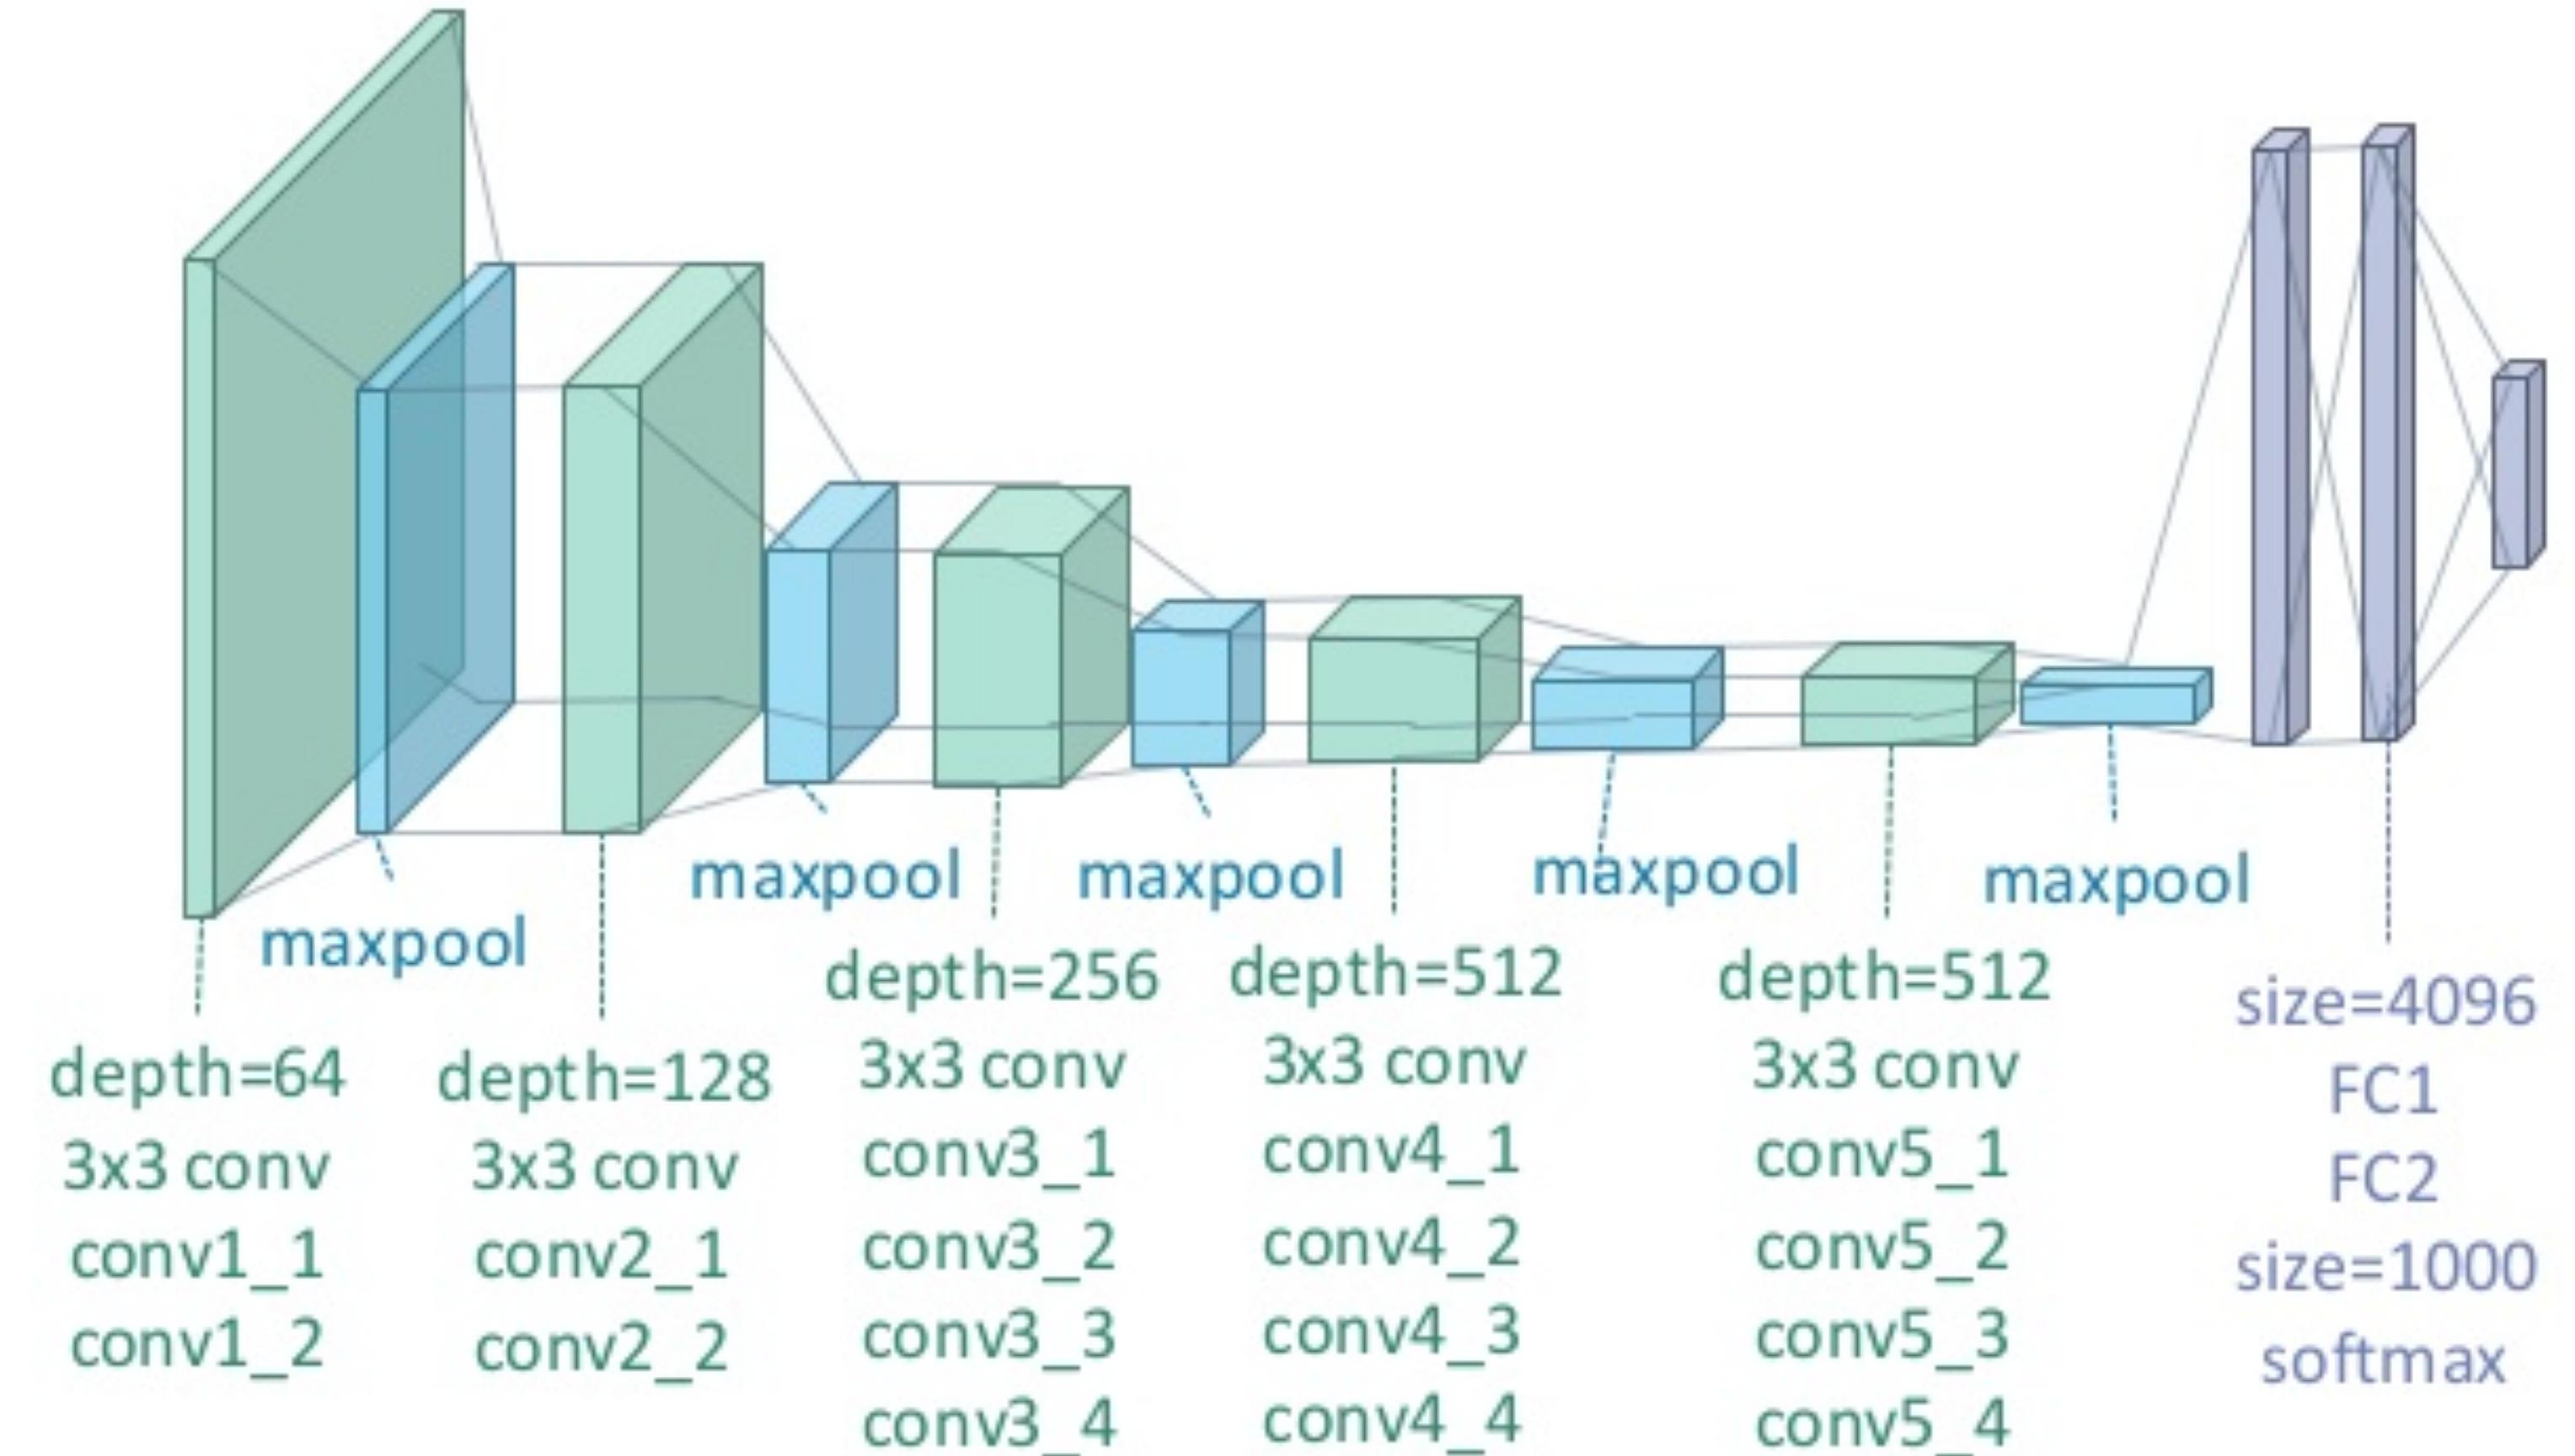
8|0|4|7|6|8|0|8|3|4|5|5|3|4|7|9|4|6|5|2|6

- All these operations are usually provided with NN training libraries

The full network

- A full ConvNN is a sequence of *Con2D+Pooling (+BatchNormalization+Dropout) layers*
- The Conv+Pooling layer reduces the 2D image representation
- The use of multiple filters on the image make the output grow on a third dimension
- Eventually, flattening occurs and the result is given to a dense layer

VGG 19



What does a ConvNN learn?

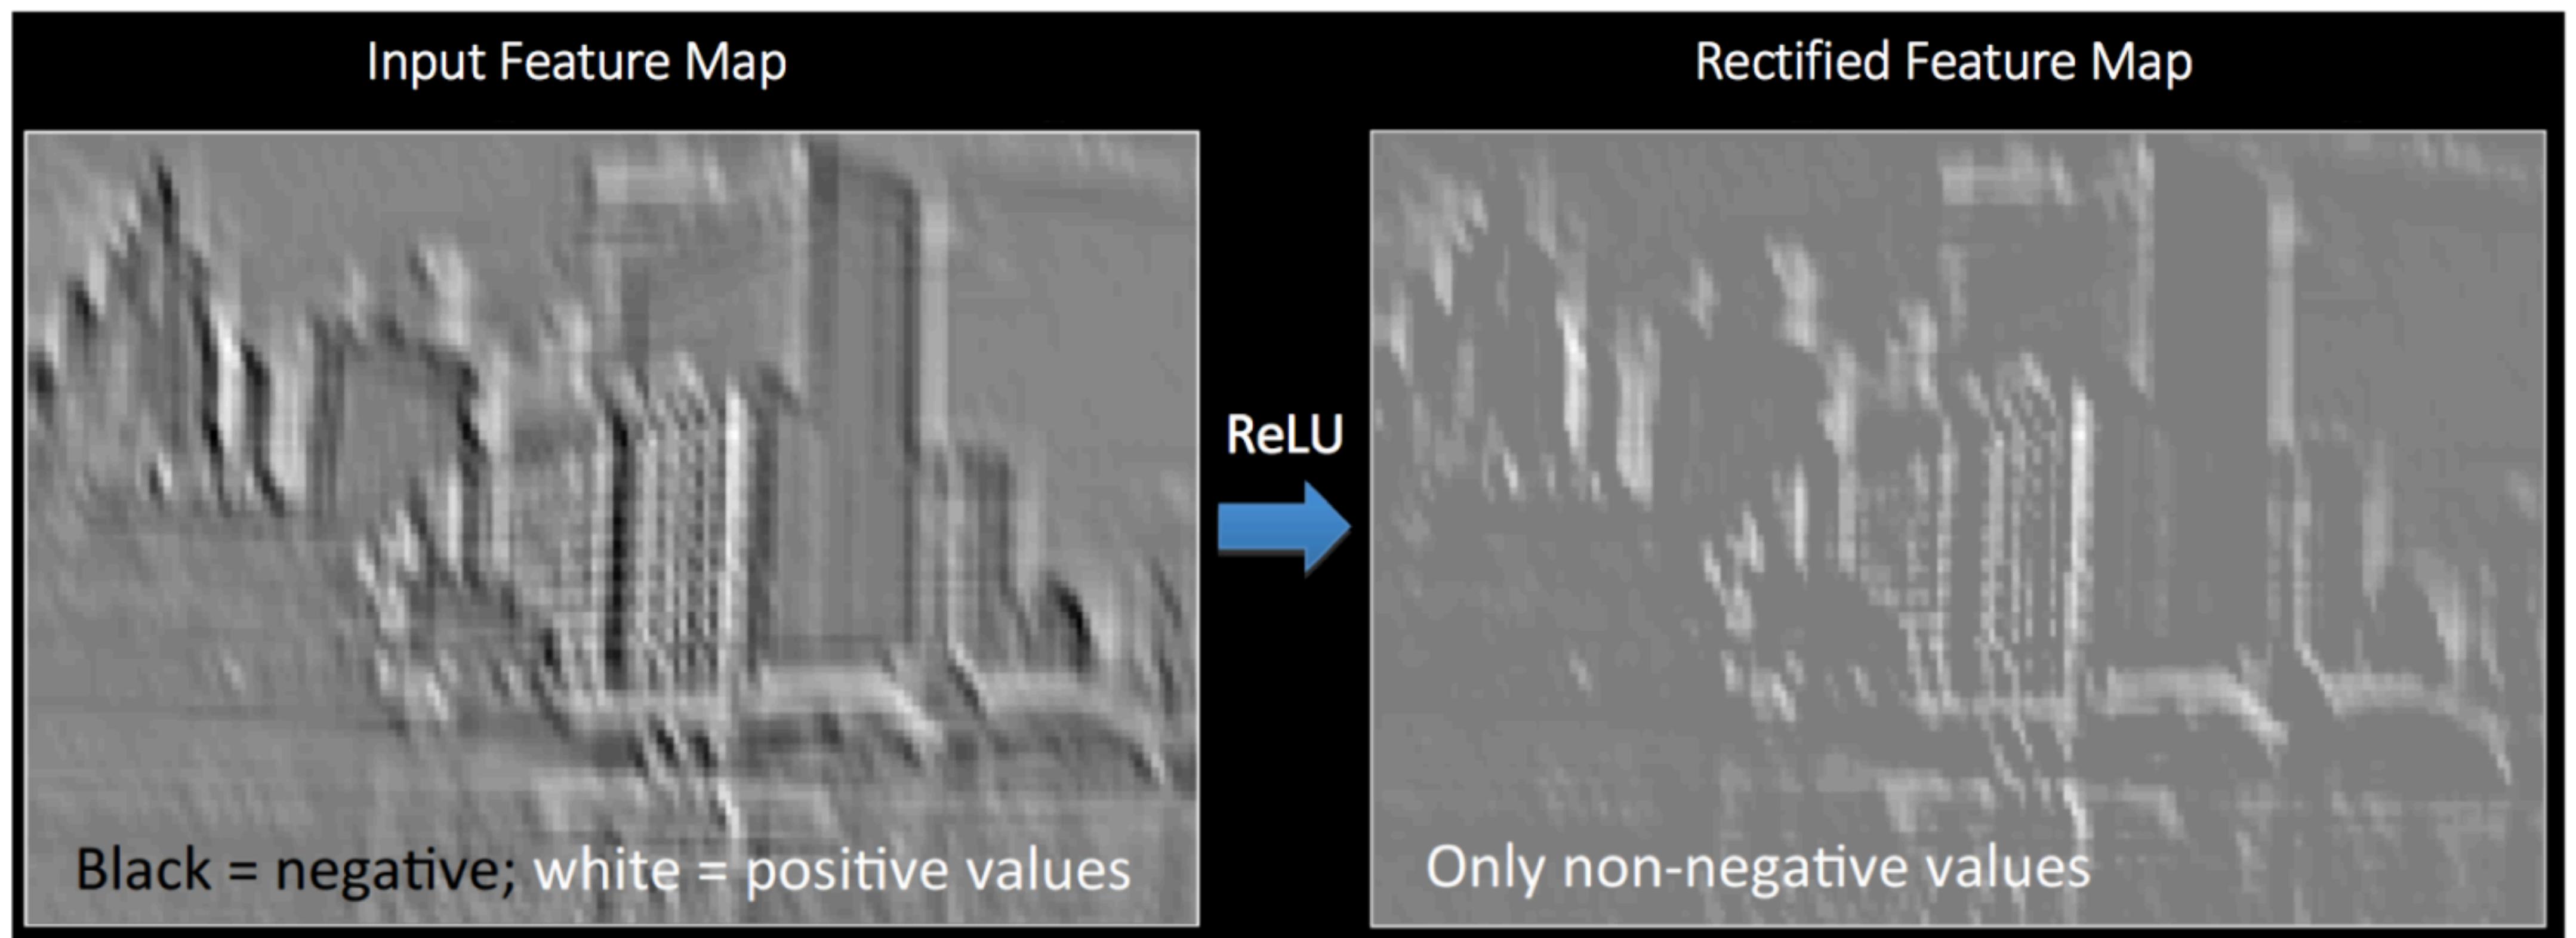
- *Each filter alters the image in a different way, picking up different aspects of the image*
- *edges oriented in various ways*
- *enhancing / blur of certain features*
- *It is interesting to check what each filter is doing (and to produce DeepDreams)*

| Operation | Filter | Convolved Image | Operation | Filter | Convolved Image |
|-----------|---|-----------------|-----------|--|-----------------|
| Identity | $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$ | | Sharpen | $\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$ | |
| | $\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$ | | | $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ | |
| | $\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$ | | | $\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$ | |
| | $\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$ | | | | |



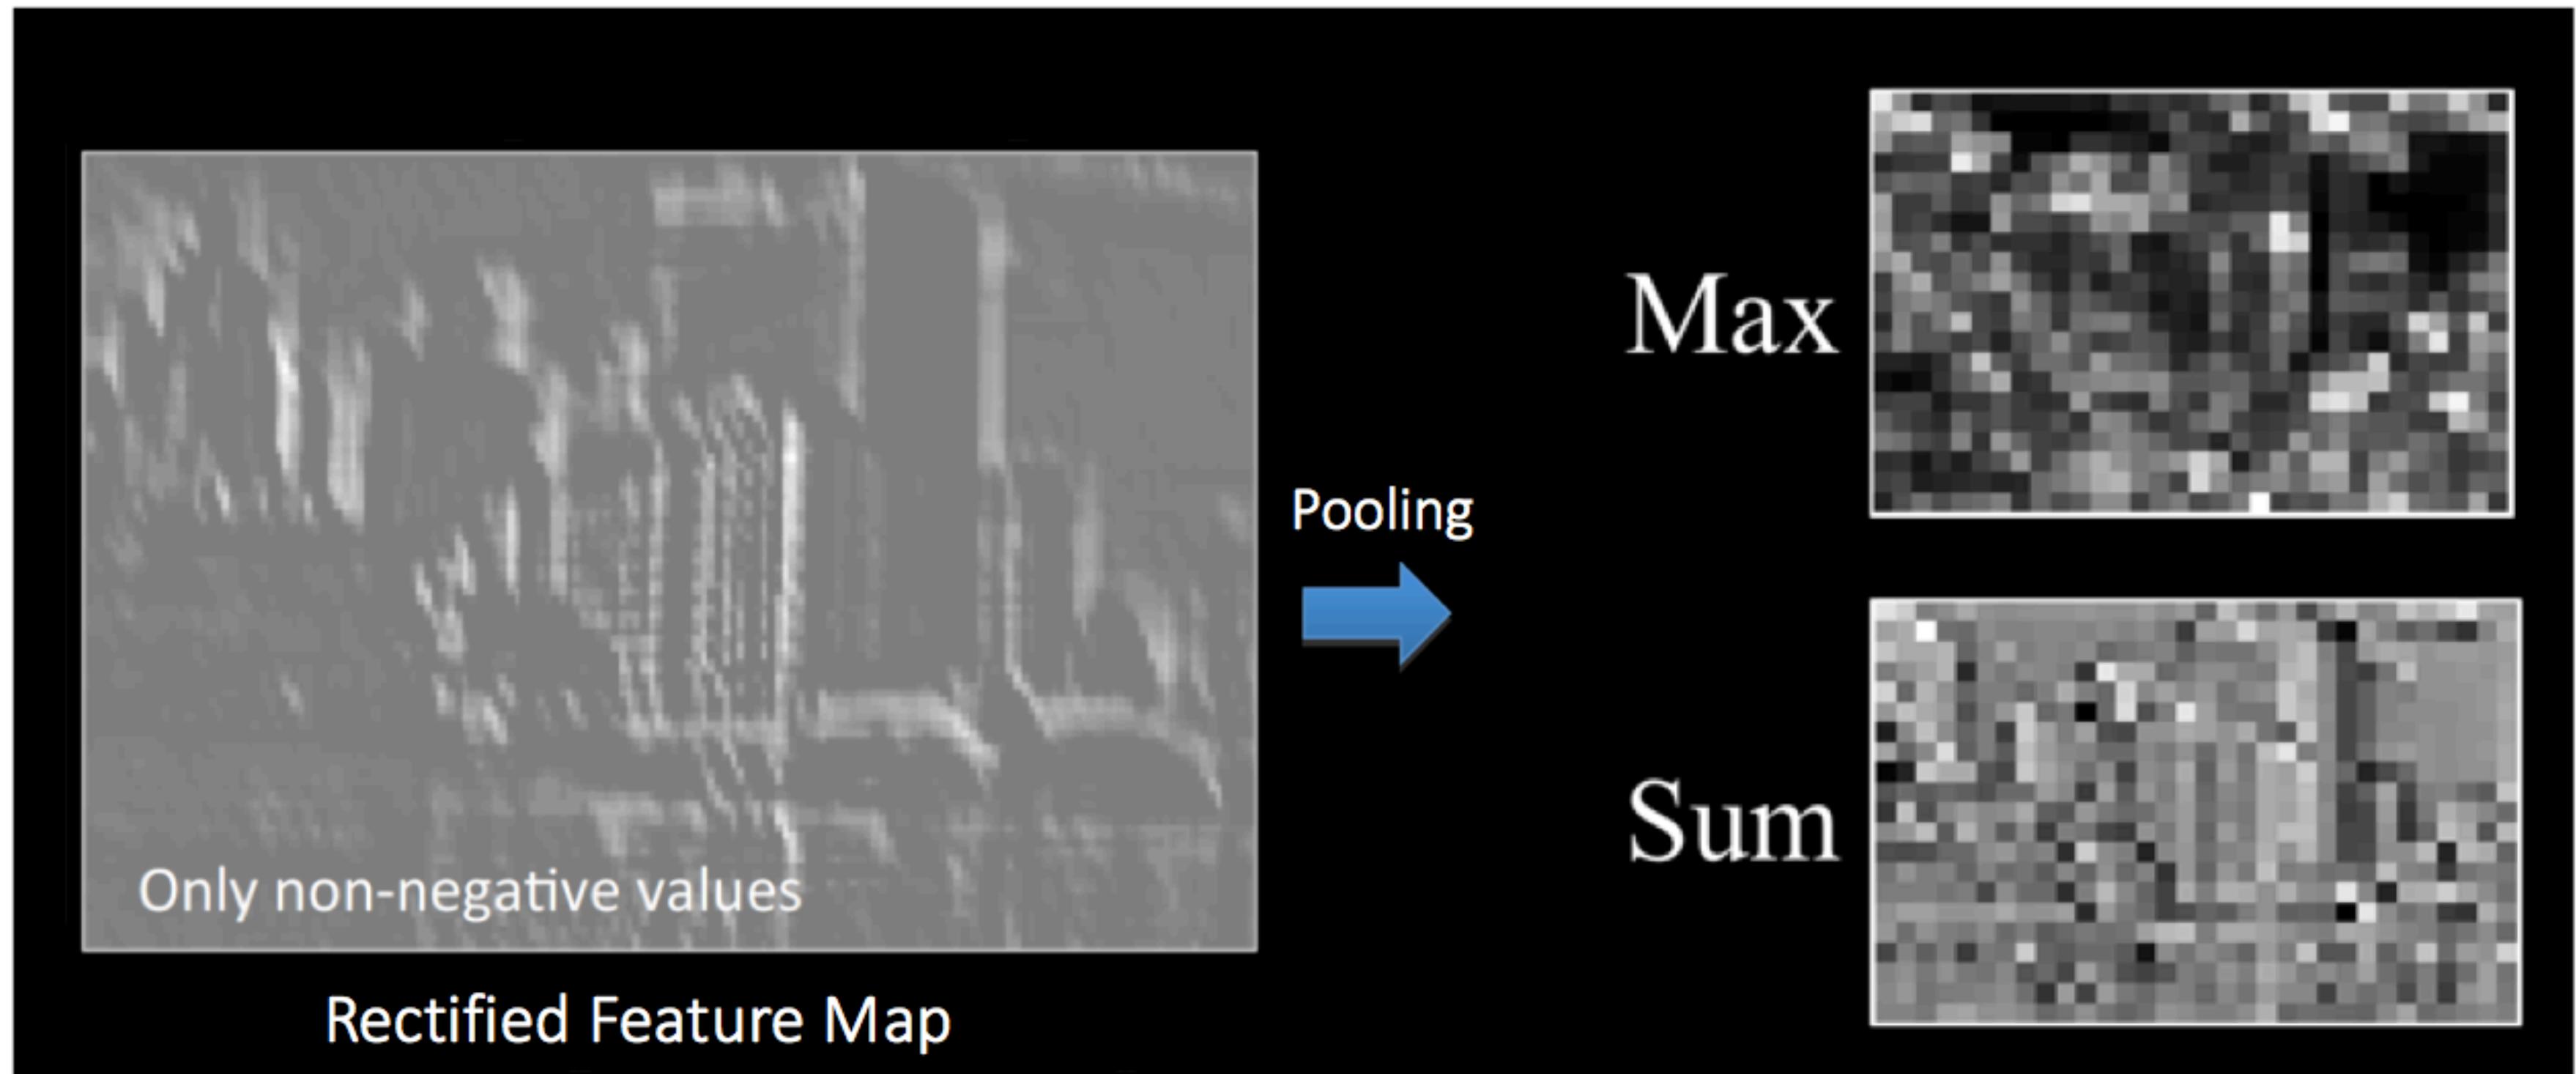
What does a ConvNet learn?

- *The use of non-linear activation functions plays a special role in enhancing features*



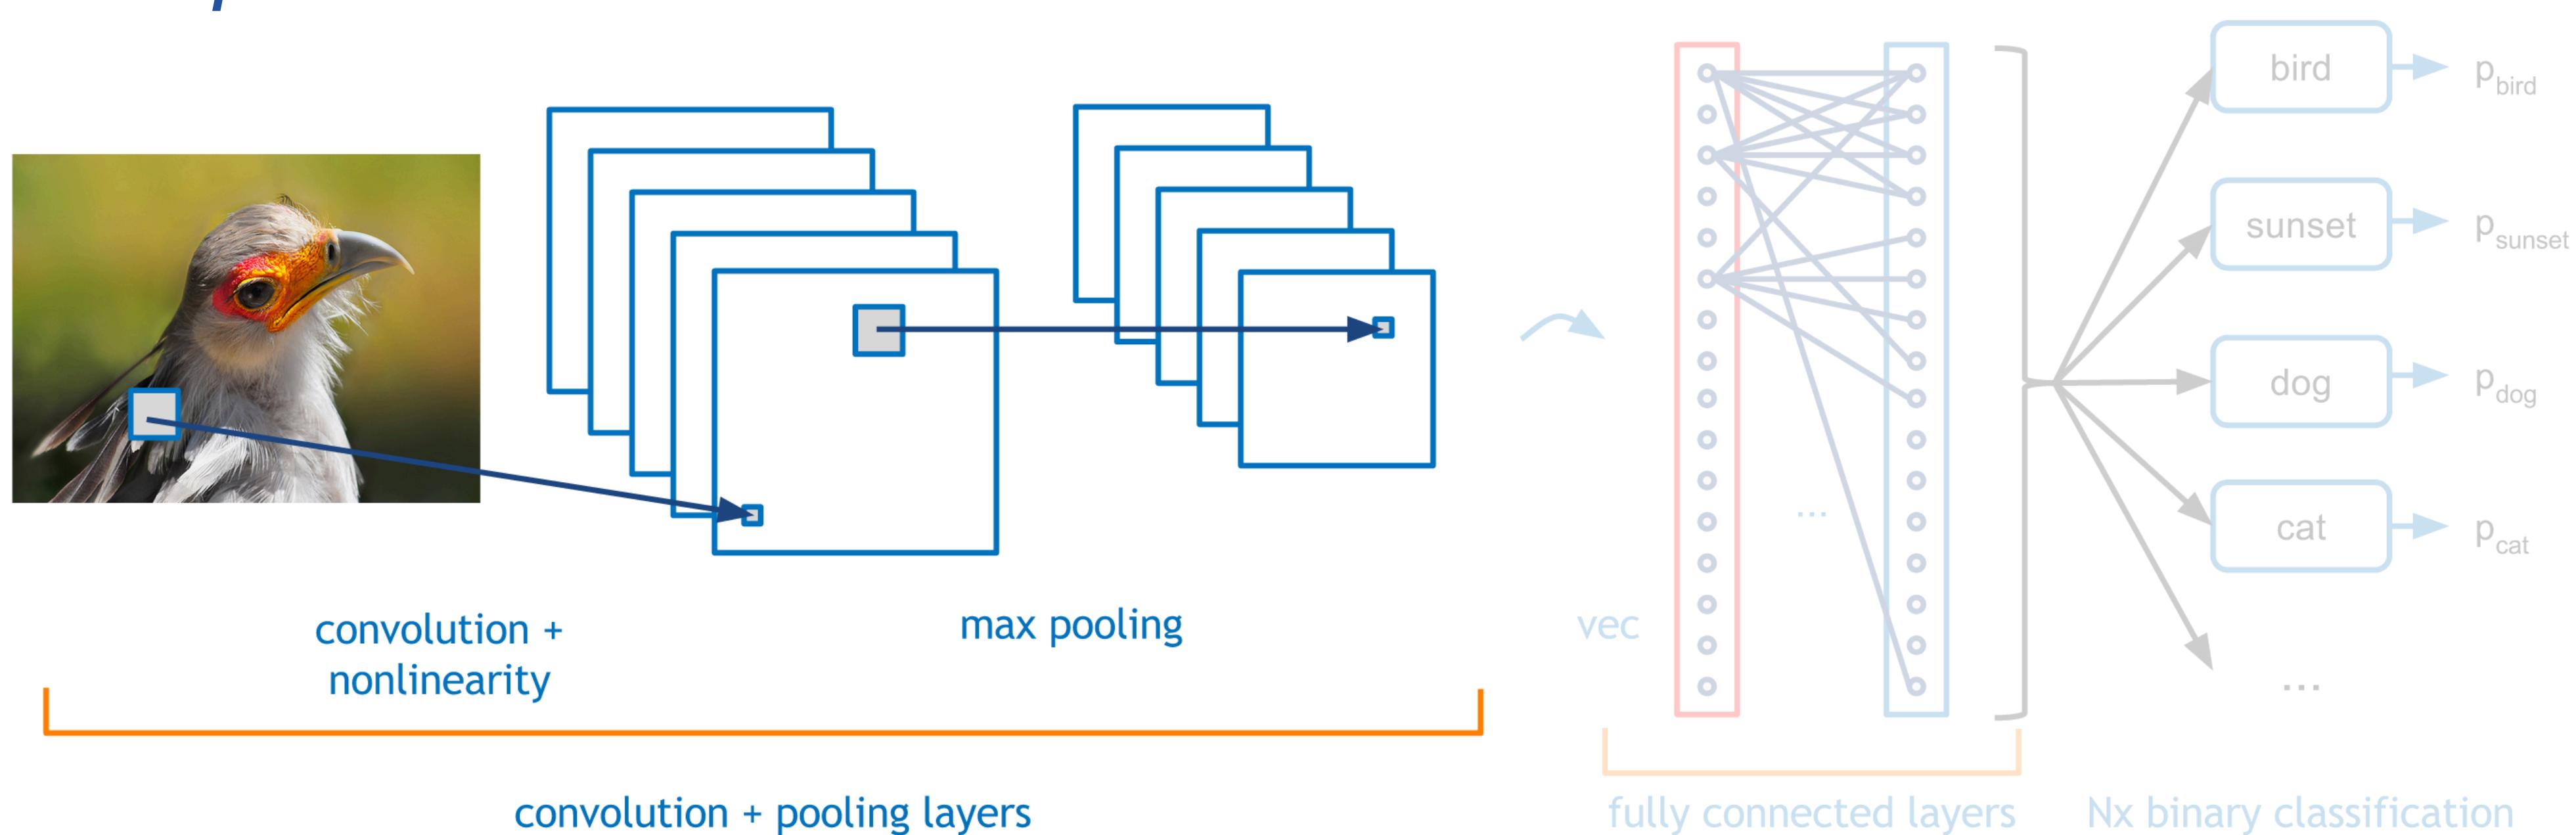
What does a ConvNet learn?

- *Pooling is important in smoothing out the image*
- *reduce parameters downstream (and prevent overfitting)*
- *makes processing independent to local features (distortion, translation)*
- *yields a scale invariant representation of the image (which could be good or bad)*



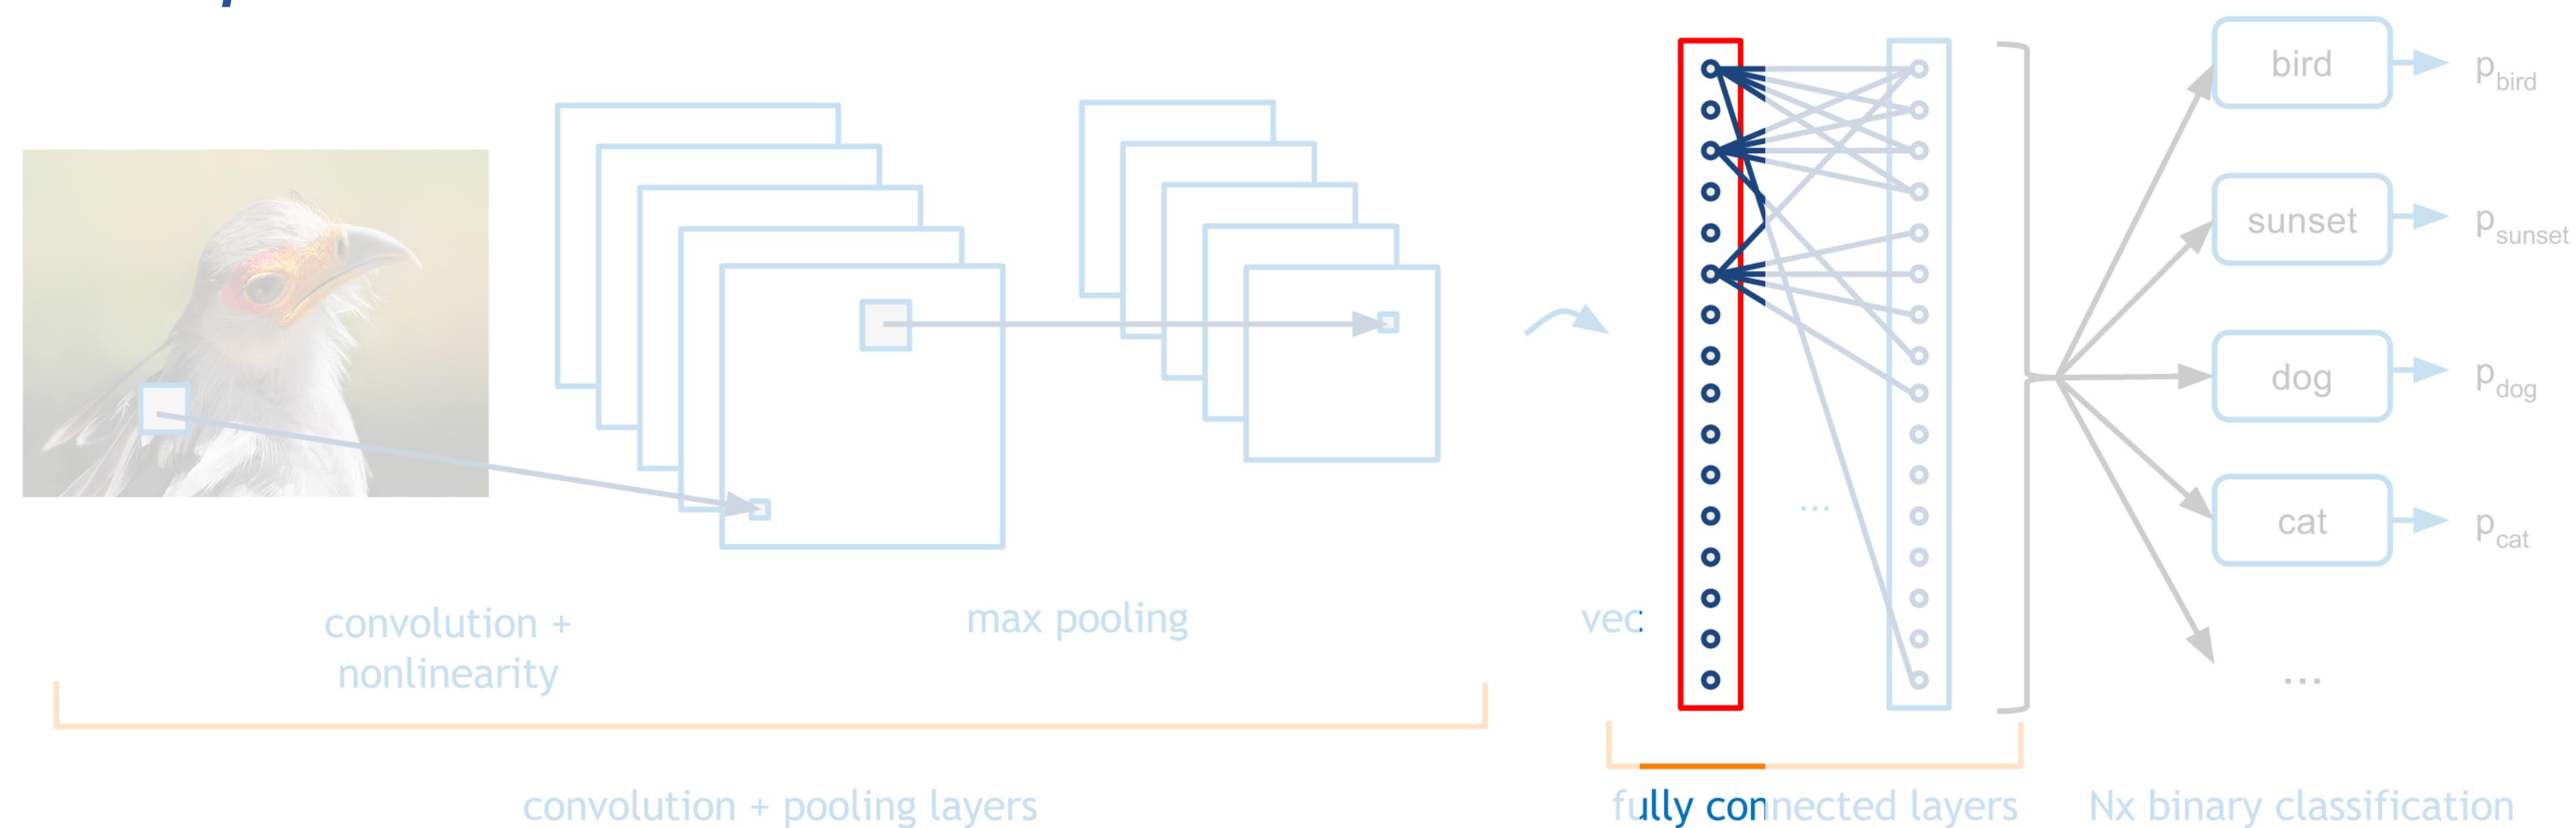
Two tasks in one network

- *The Conv layer starts from RAW data and defines interesting quantities (high level features)*
- *The HLFs replace the physics-motivated inputs of a DNN*
- *The DNN at the end exploits the engineered features to accomplish the task*



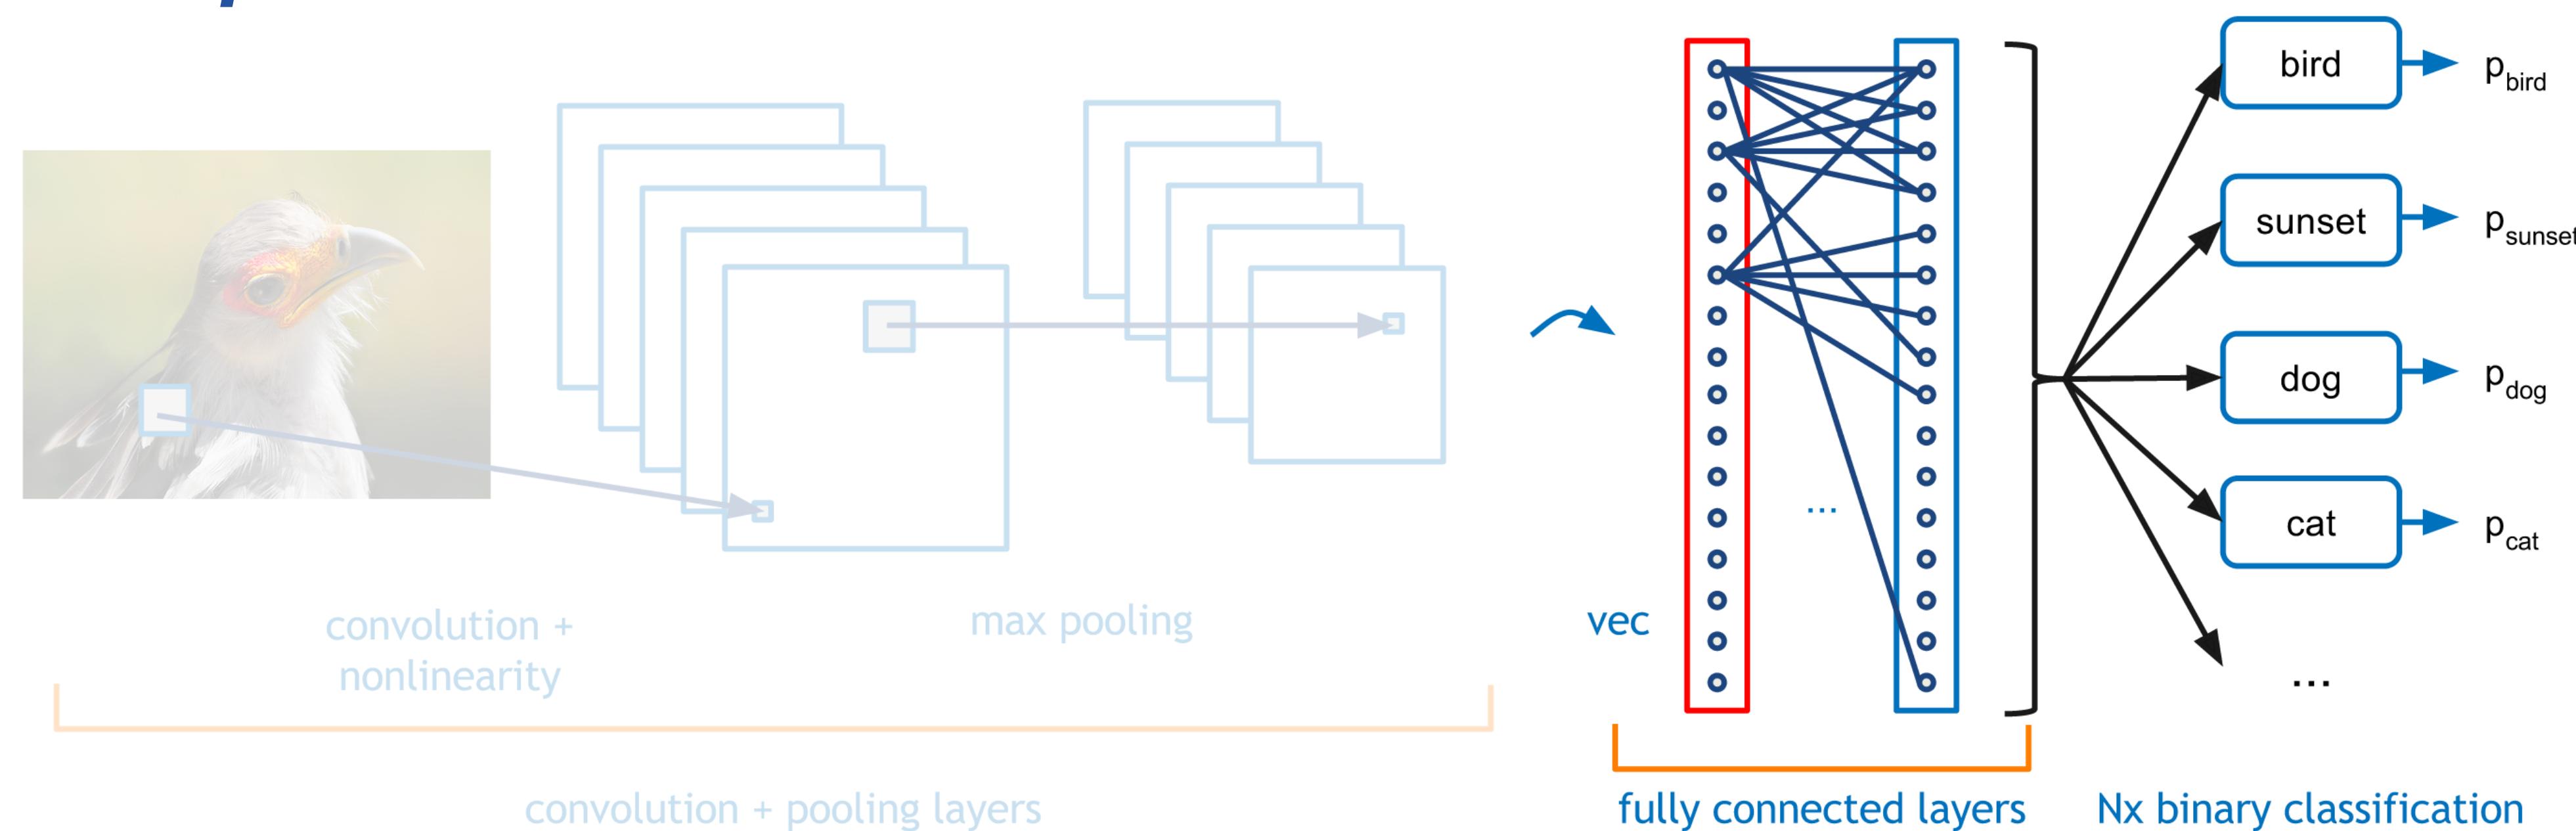
Two tasks in one network

- *The Conv layer starts from RAW data and defines interesting quantities (high level features)*
- *The HLFs replace the physics-motivated inputs of a DNN*
- *The DNN at the end exploits the engineered features to accomplish the task*



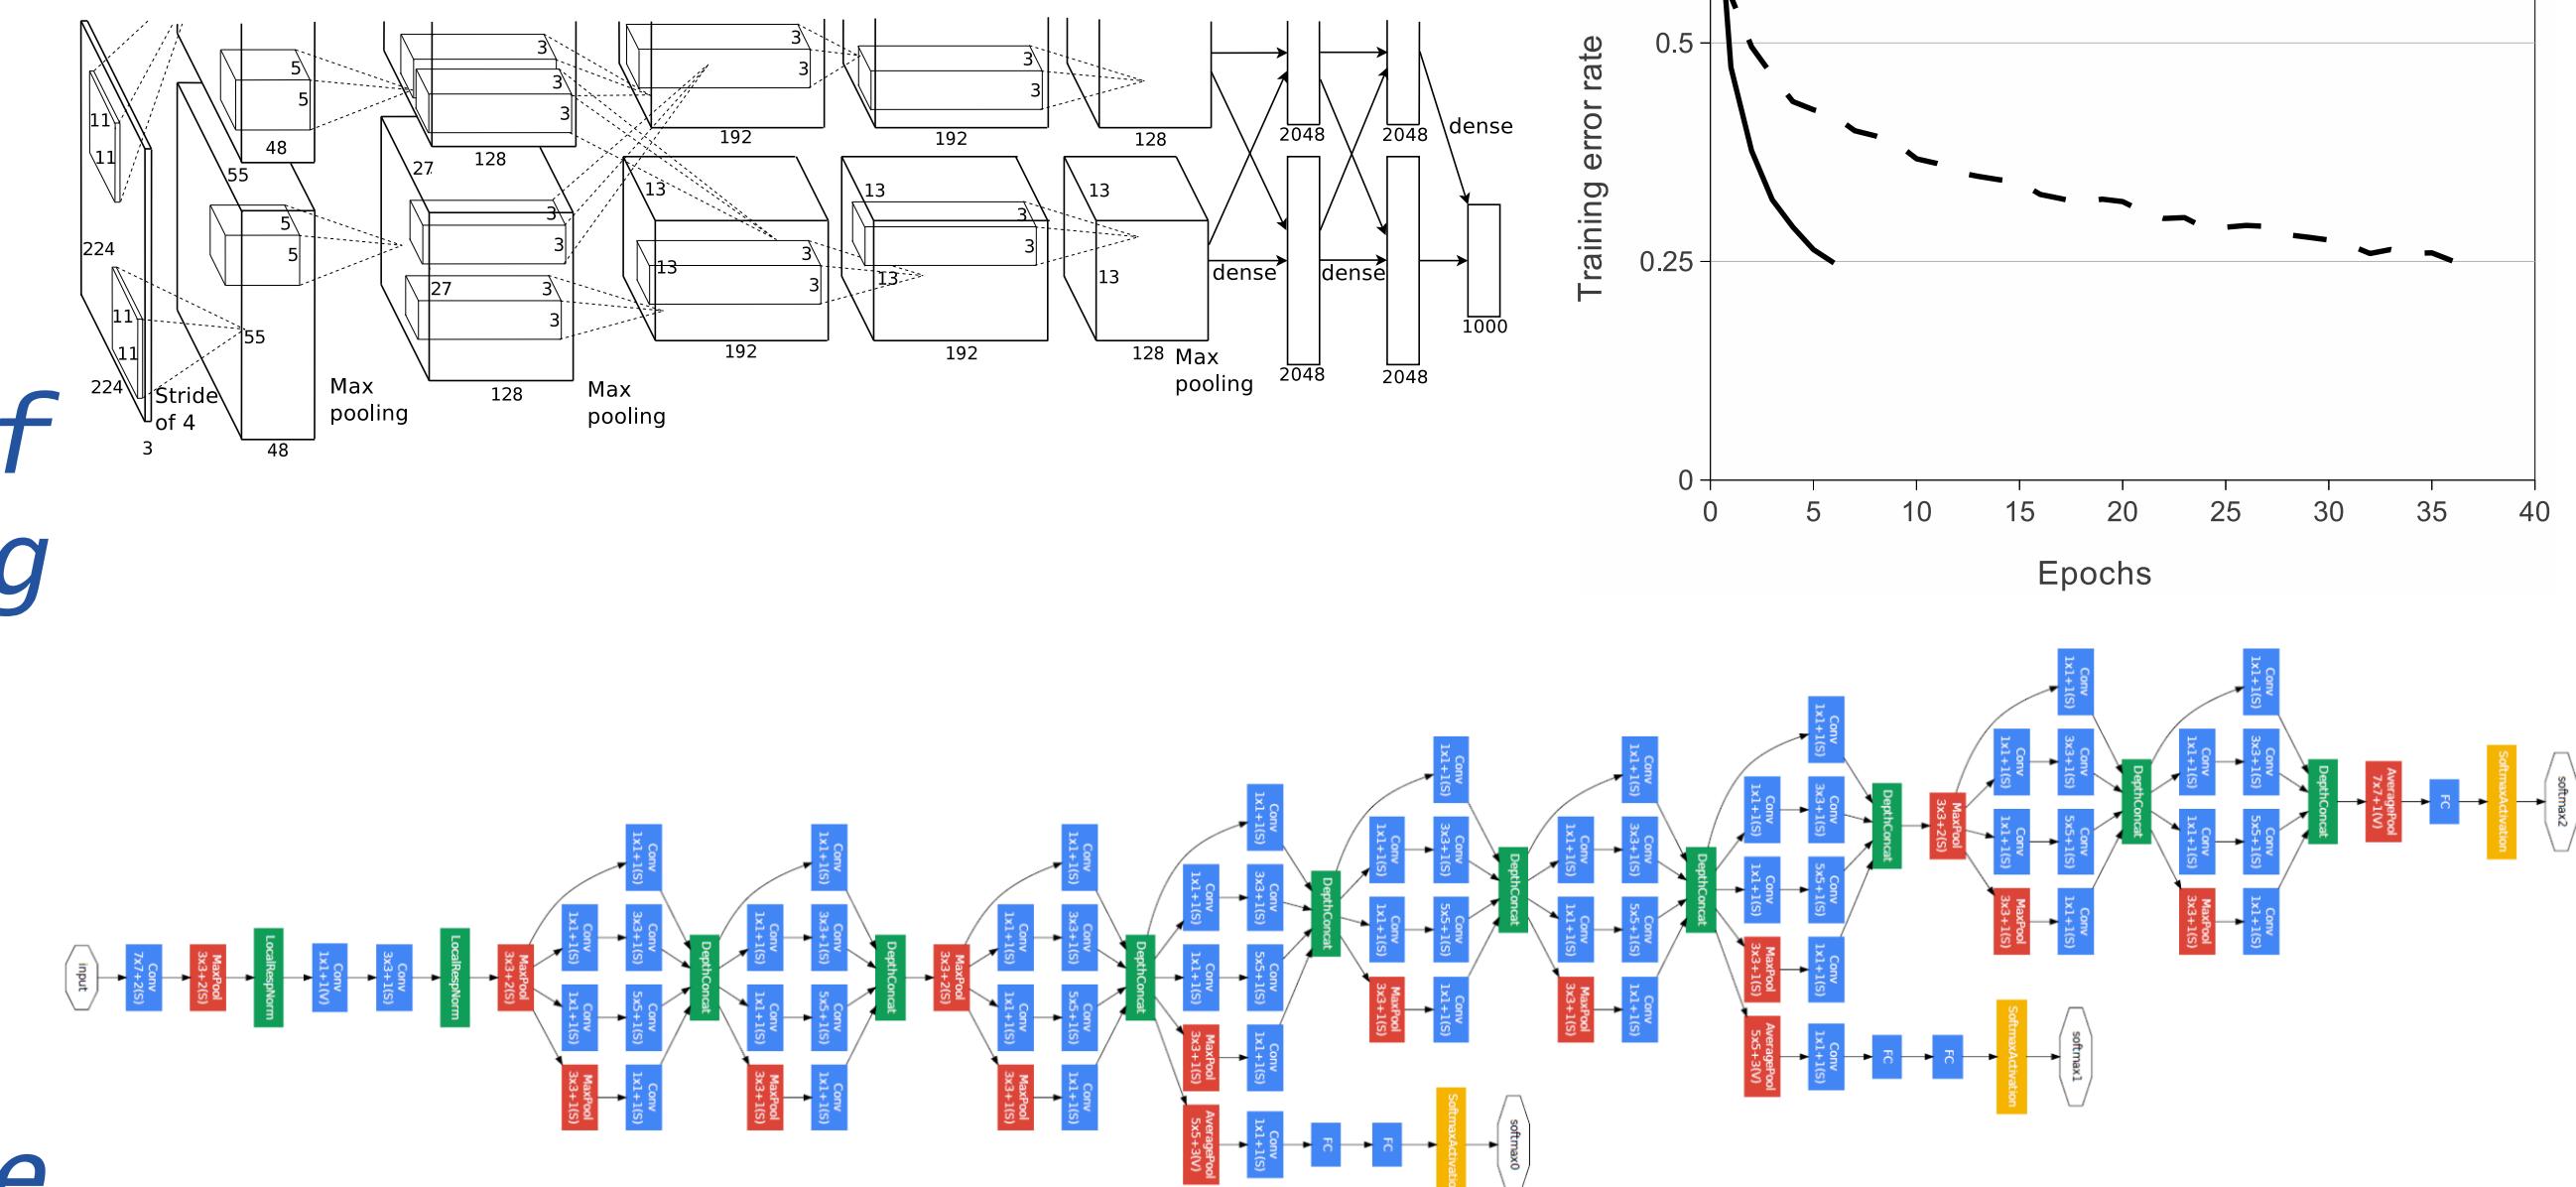
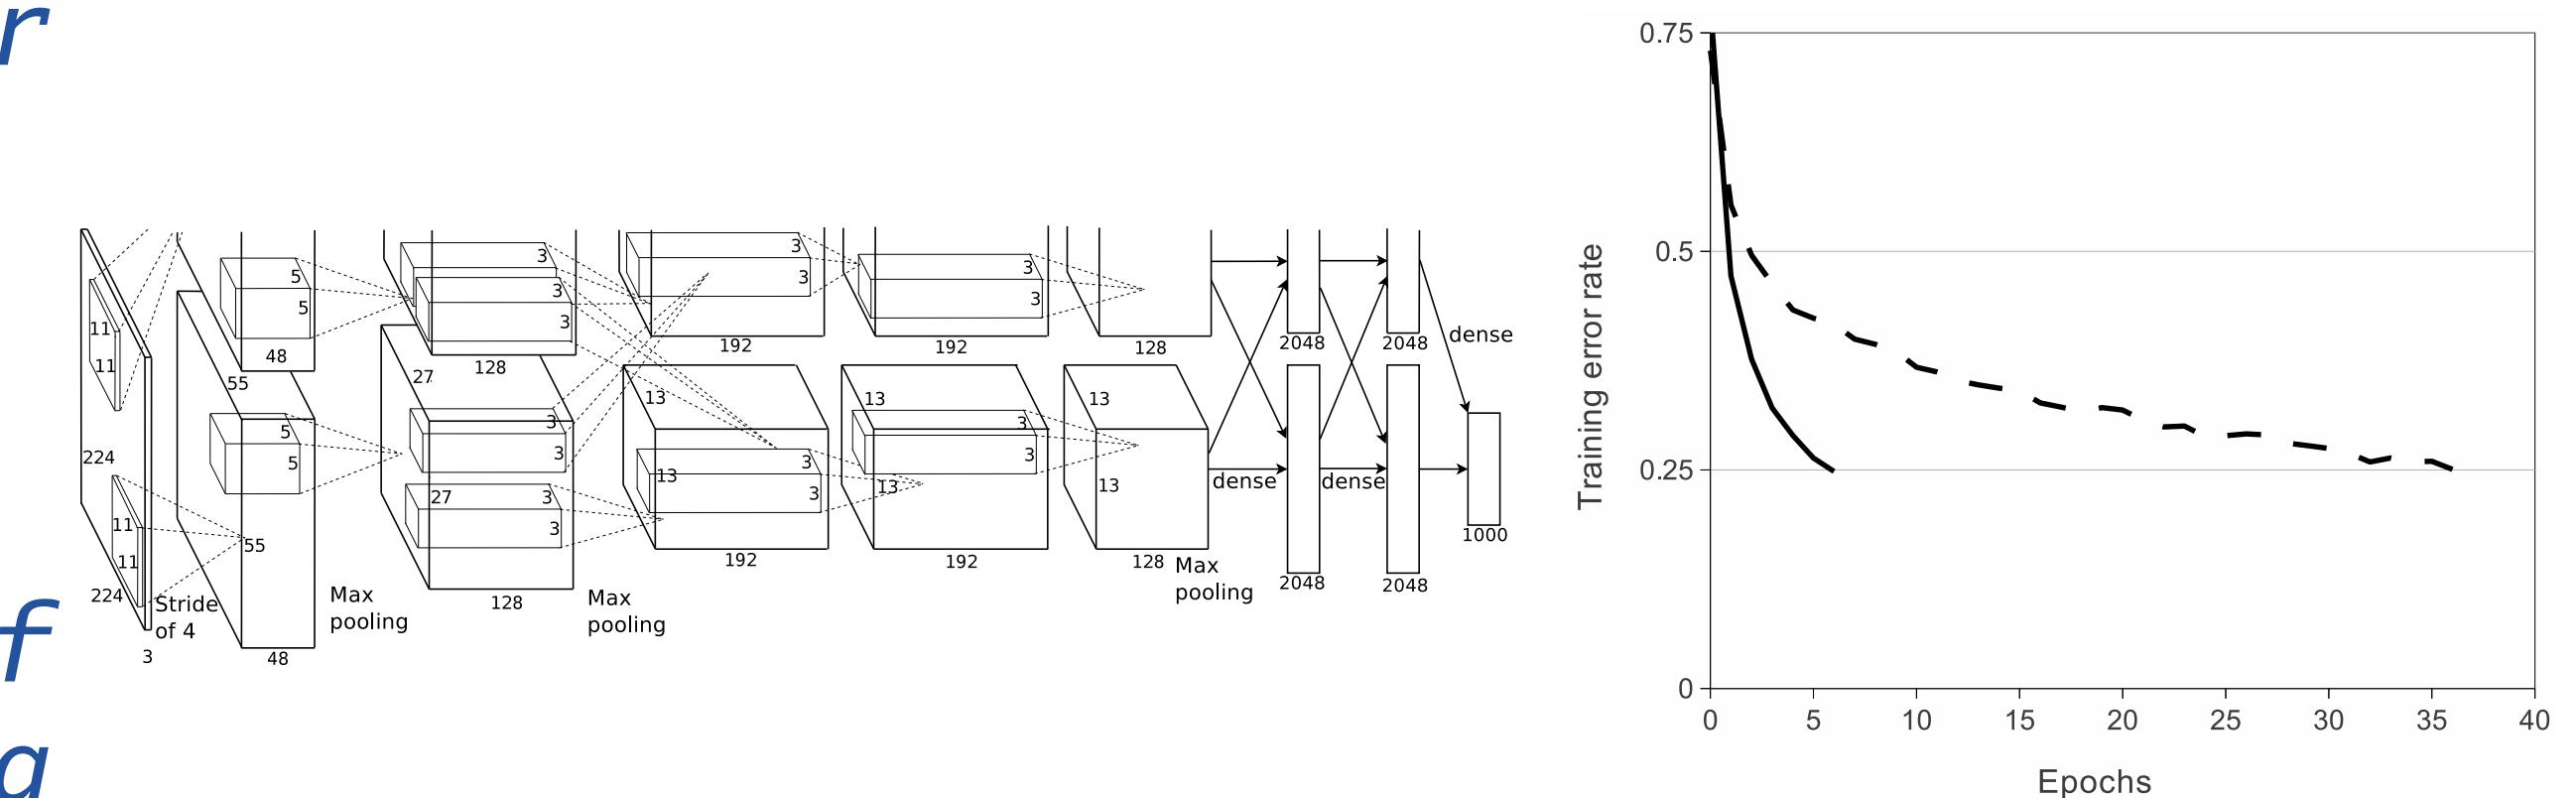
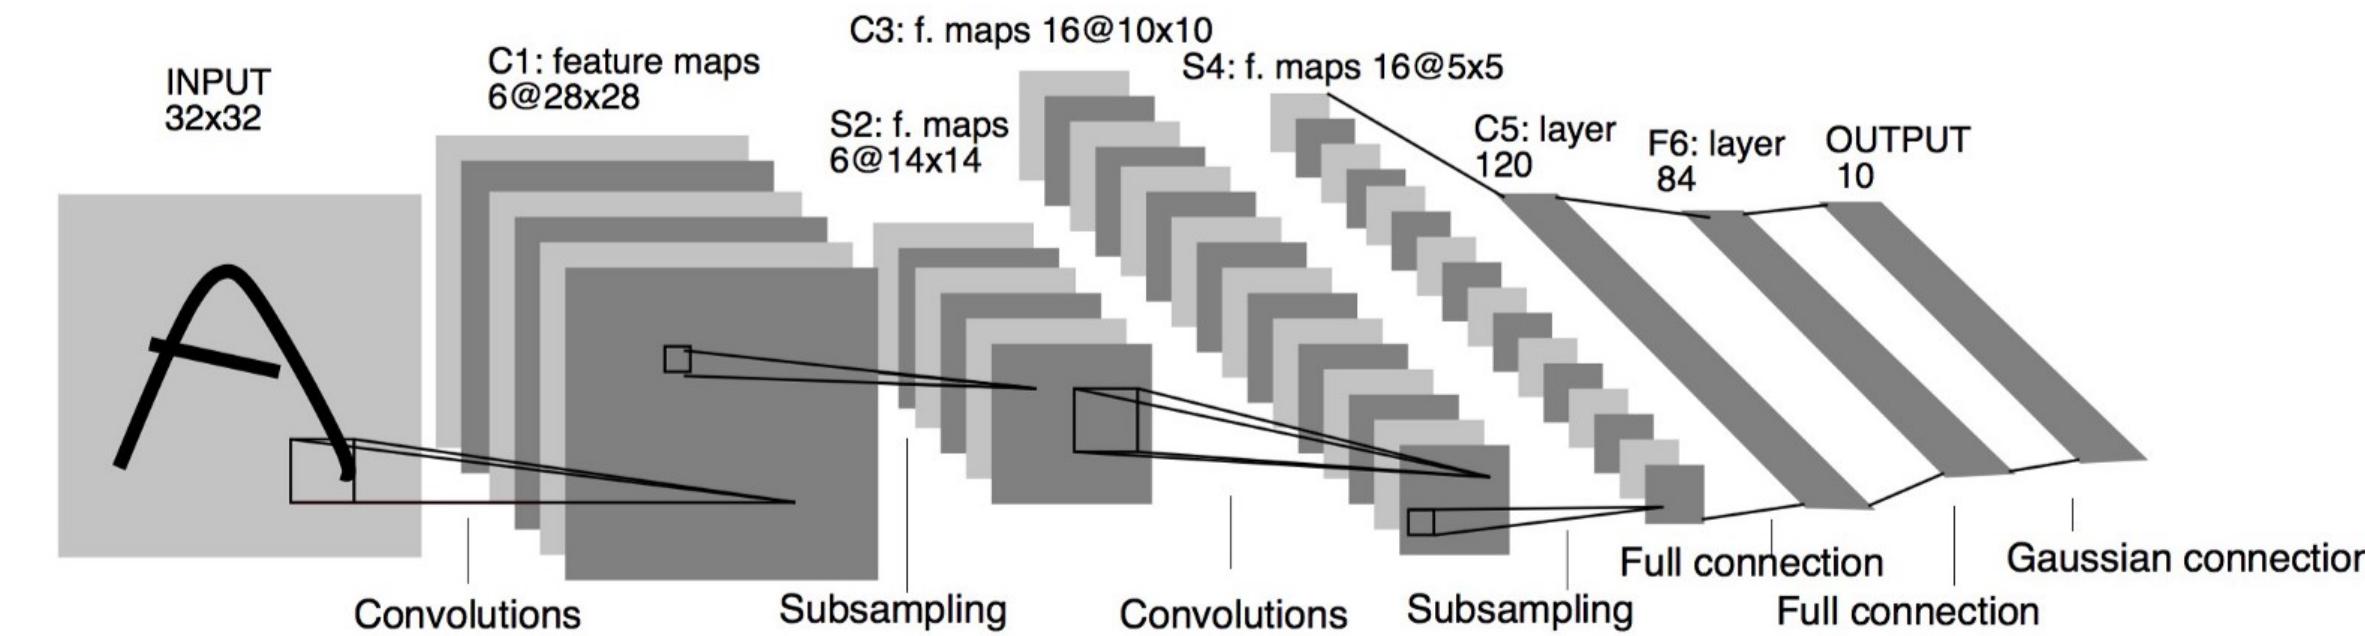
Two tasks in one network

- The Conv layer starts from RAW data and defines interesting quantities (high level features)
- The HLFs replace the physics-motivated inputs of a DNN
- The DNN at the end exploits the engineered features to accomplish the task



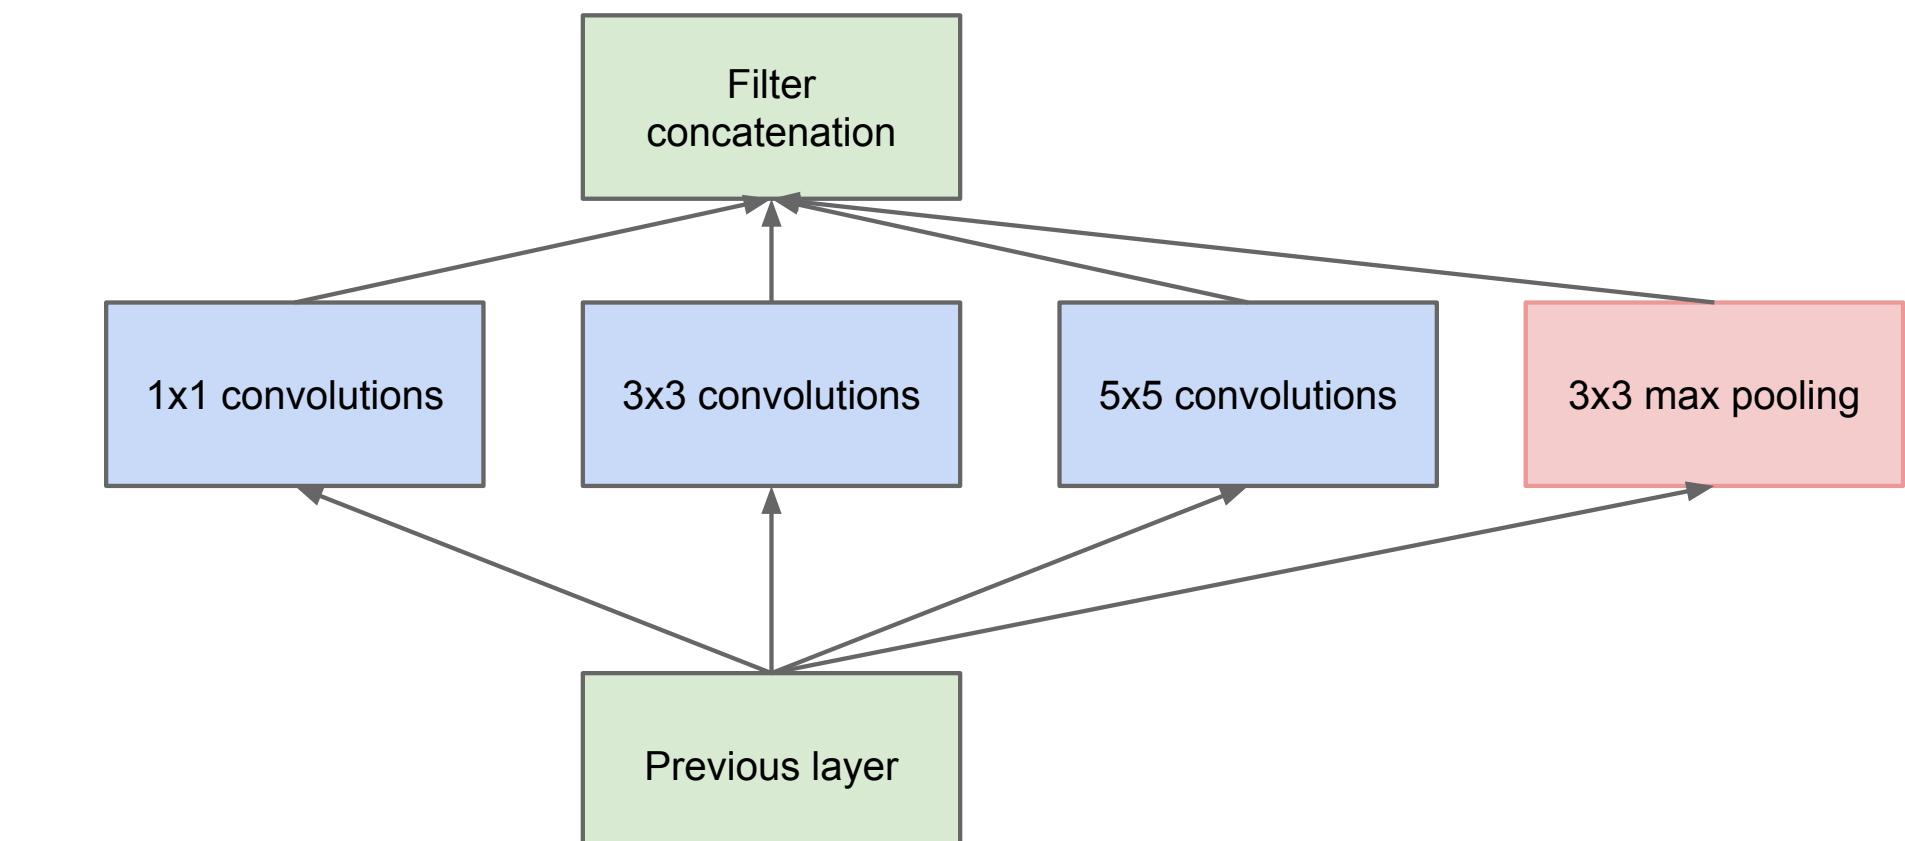
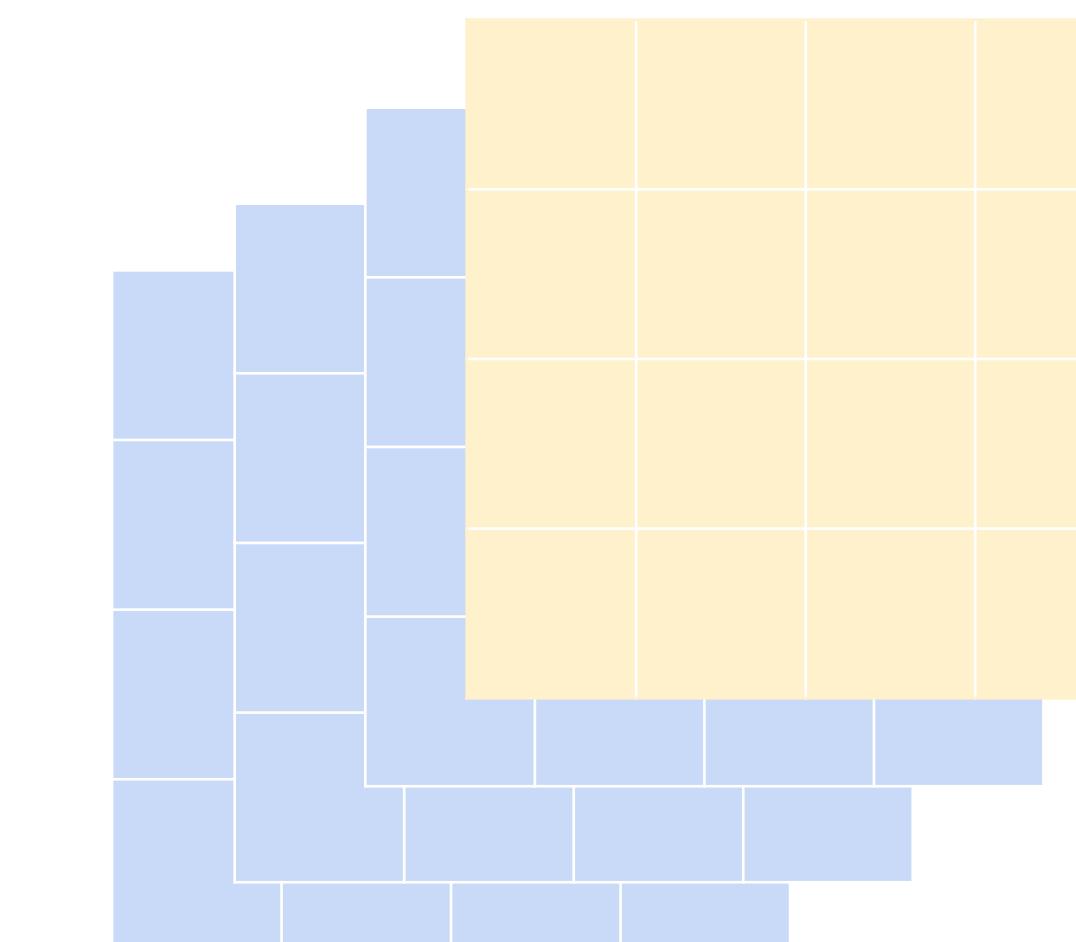
A history of ConvNns

- Neocognitron (1980): translation-invariant image processing with NNs
- LeNet (1989): considered the very first ConvNN, designer for digit recognition (ZIP codes)
- AlexNet (2012): the first big ConvNN (60M parameters, 650K neurons), setting the state of the art: trained on GPUs, using ReLU and Dropout
- GoogLeNet (2014): built on AlexNet, introduced an inception model to reduce the number of parameters

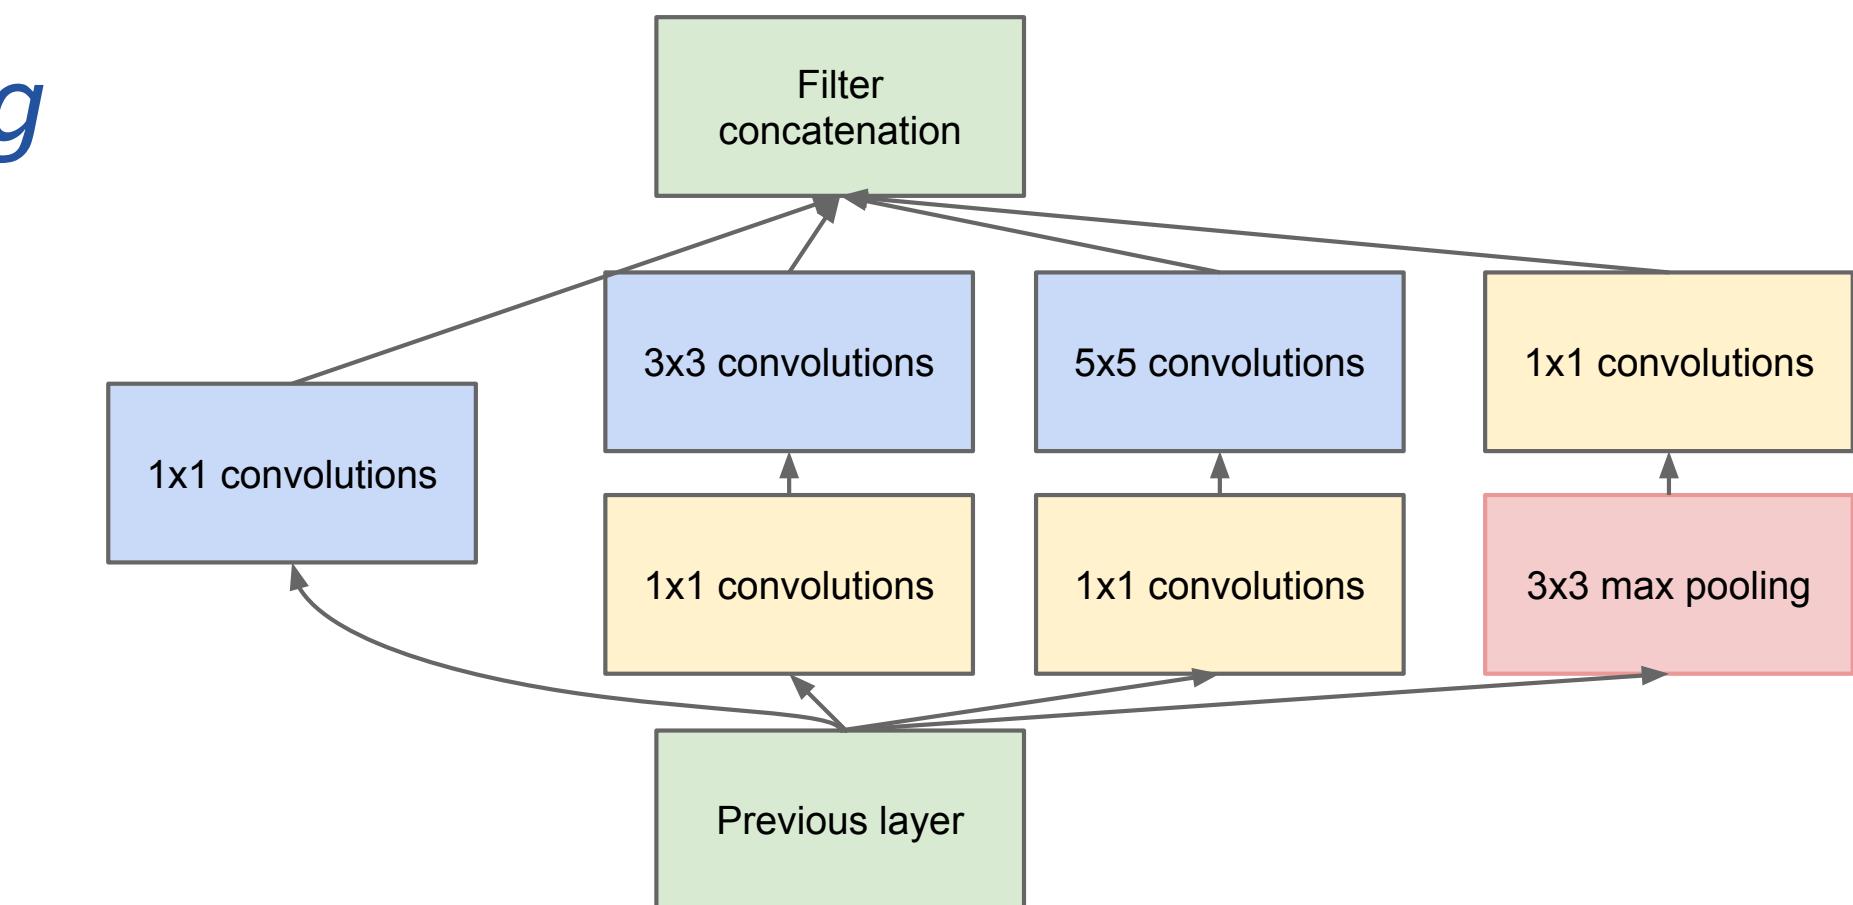


Inception module

- Rather than going deeper and deeper, inception architecture go wider
- Several conv layers, with different filter size, process the same inputs
- This way, more features can be detected from the same image
- The outcome of this parallel processing is then recombined through a concatenation step ass channels of an image



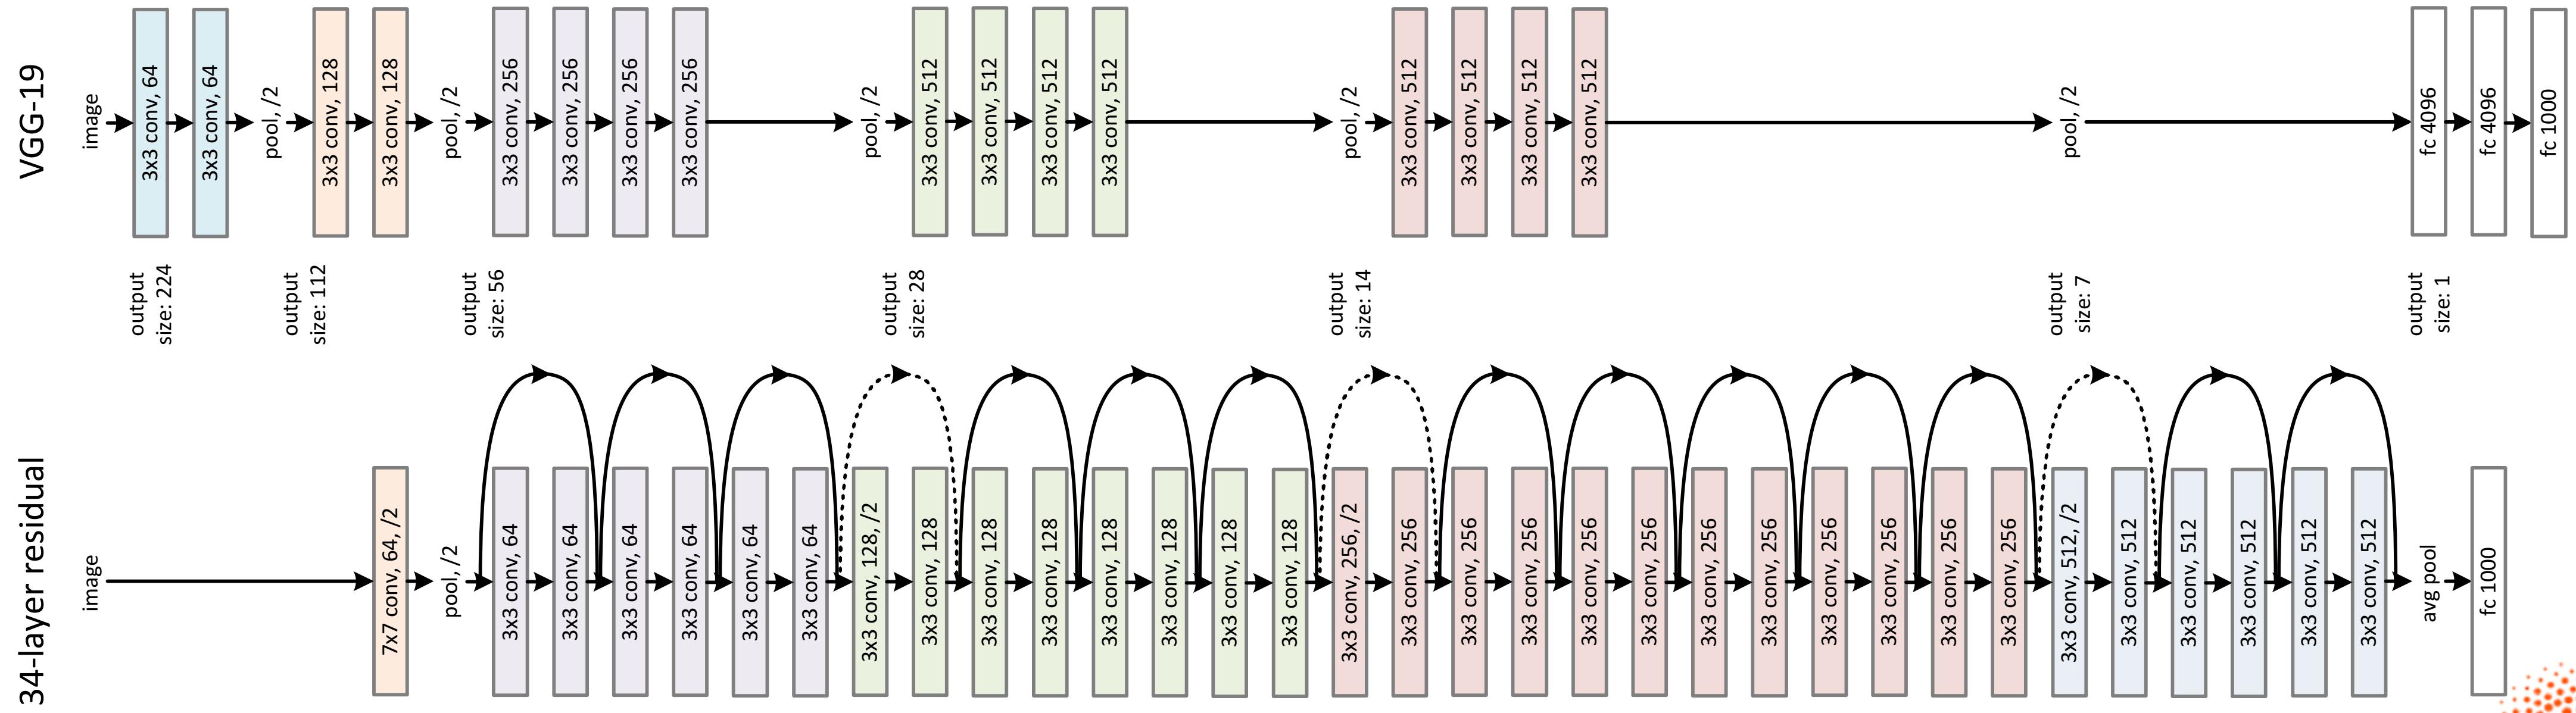
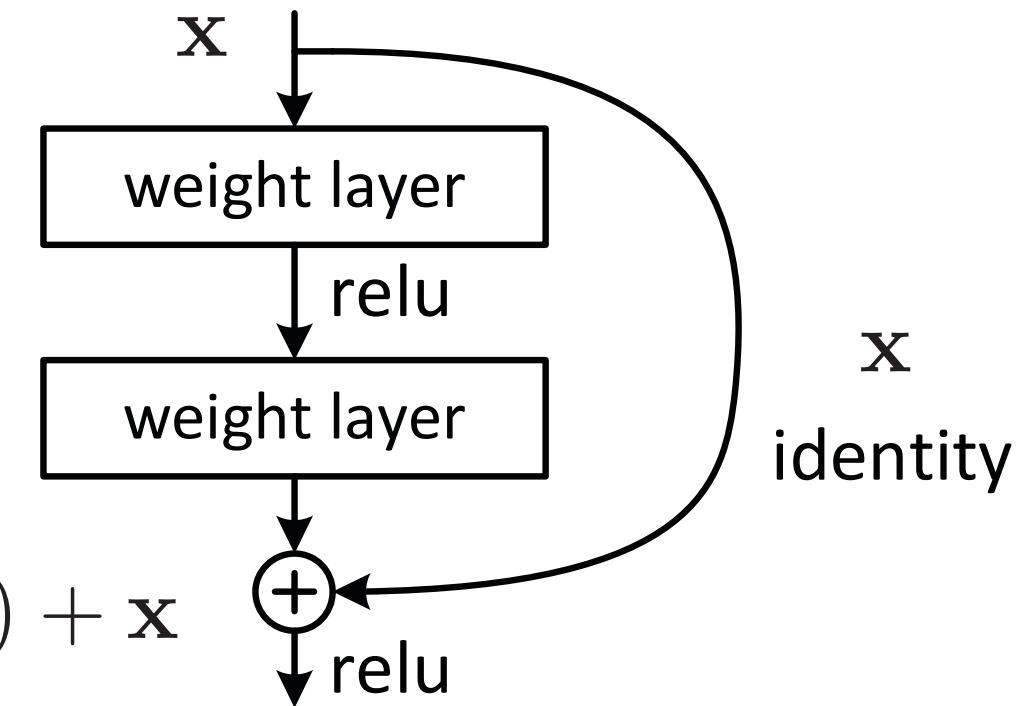
(a) Inception module, naïve version



(b) Inception module with dimension reductions

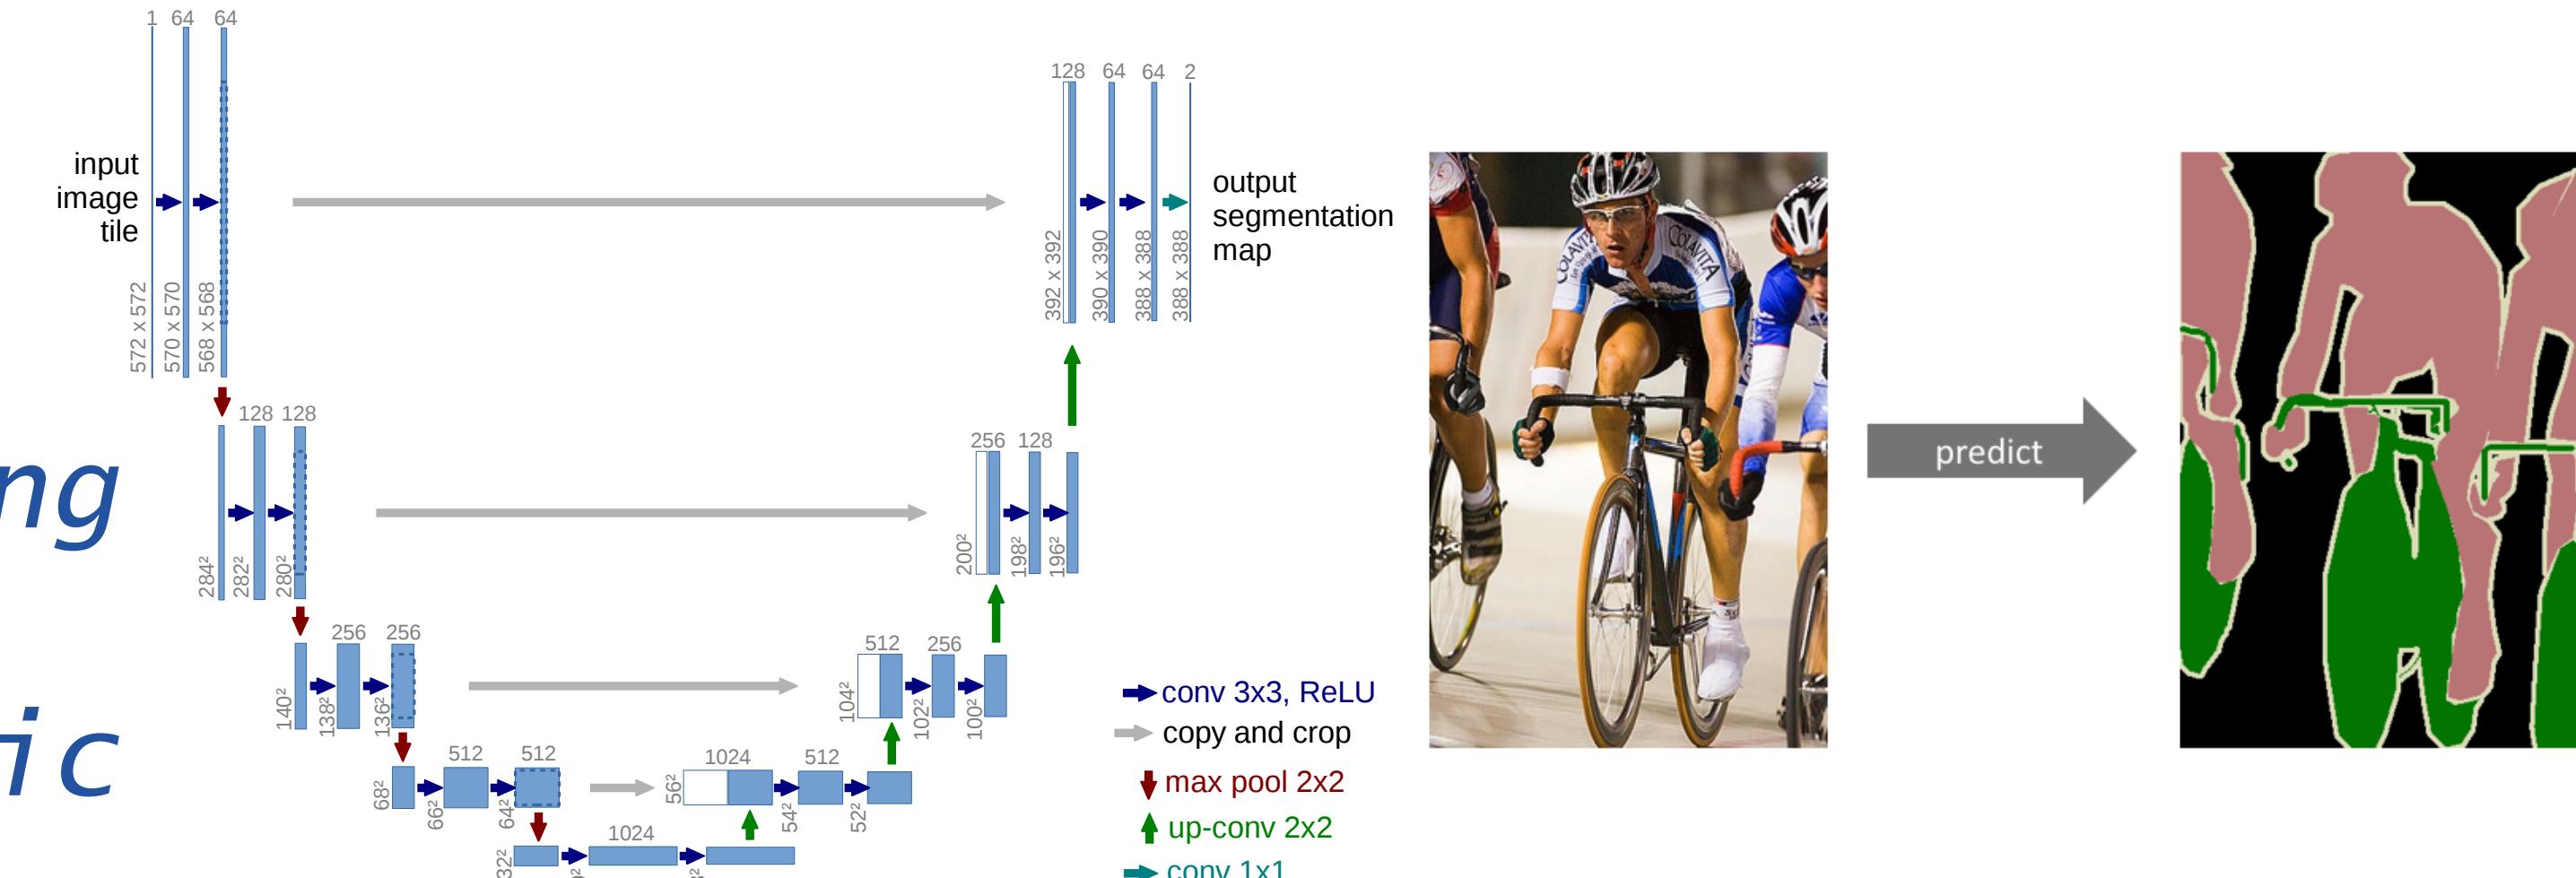
A history of ConvNNS

- VGGNet (2014): exploit small filters (3×3 , 1×1), previously considered not optimal in a stack of Conv layers (rather than 1Conv+1Pooling) $\mathcal{F}(x)$
- ResNet (2015): implemented skip connections, which were proven to boost performances

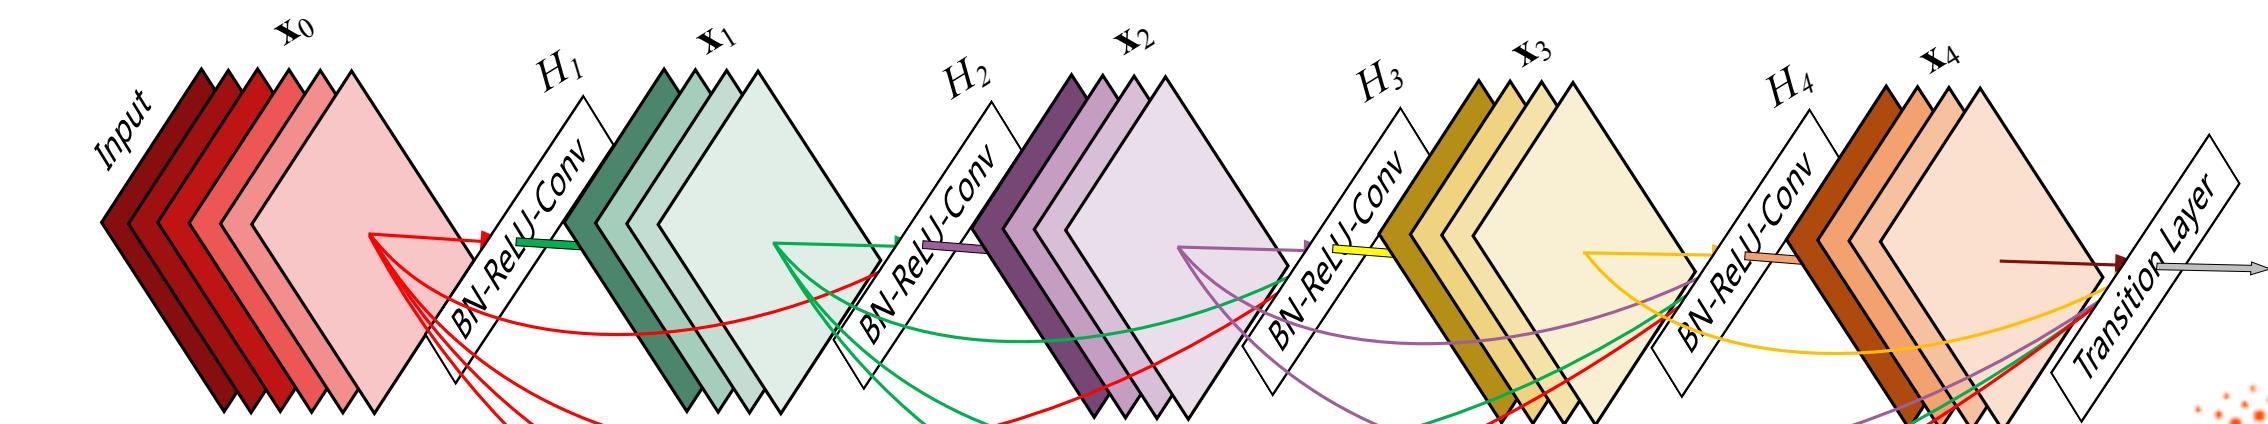
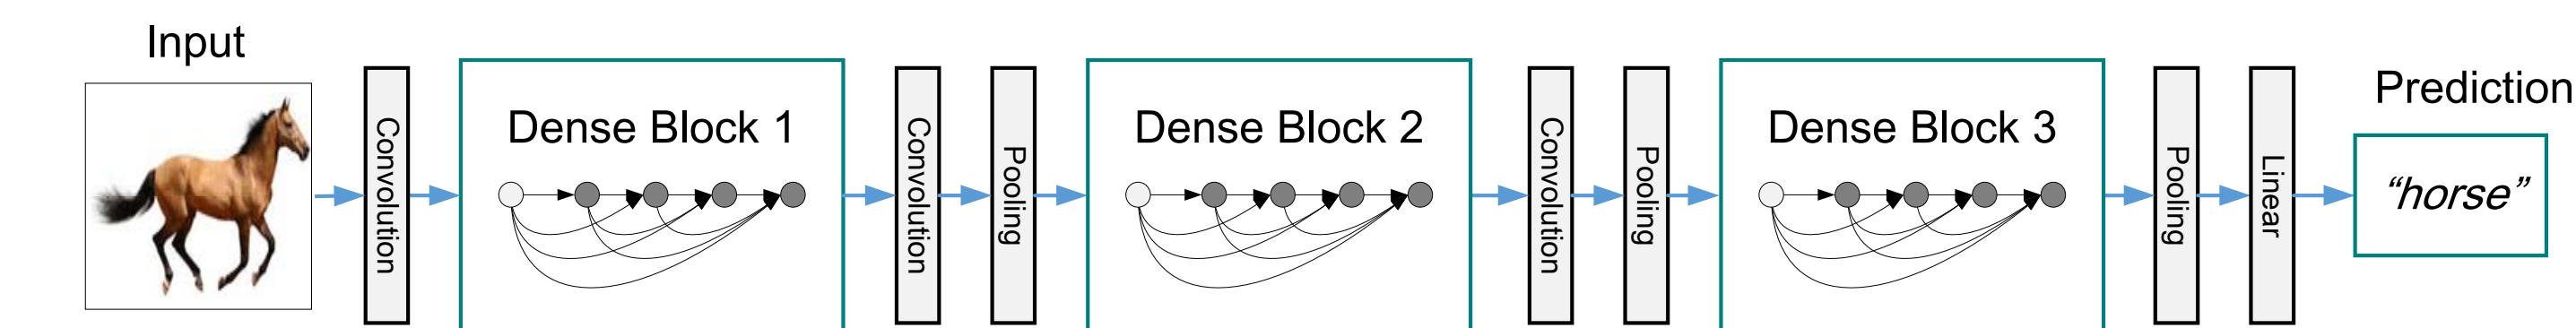


A history of ConvNets

- [U-net \(2015\)](#): conv layers with skip connections, in a downsampling+upsampling U-shaped sequence.
Introduced for semantic segmentation



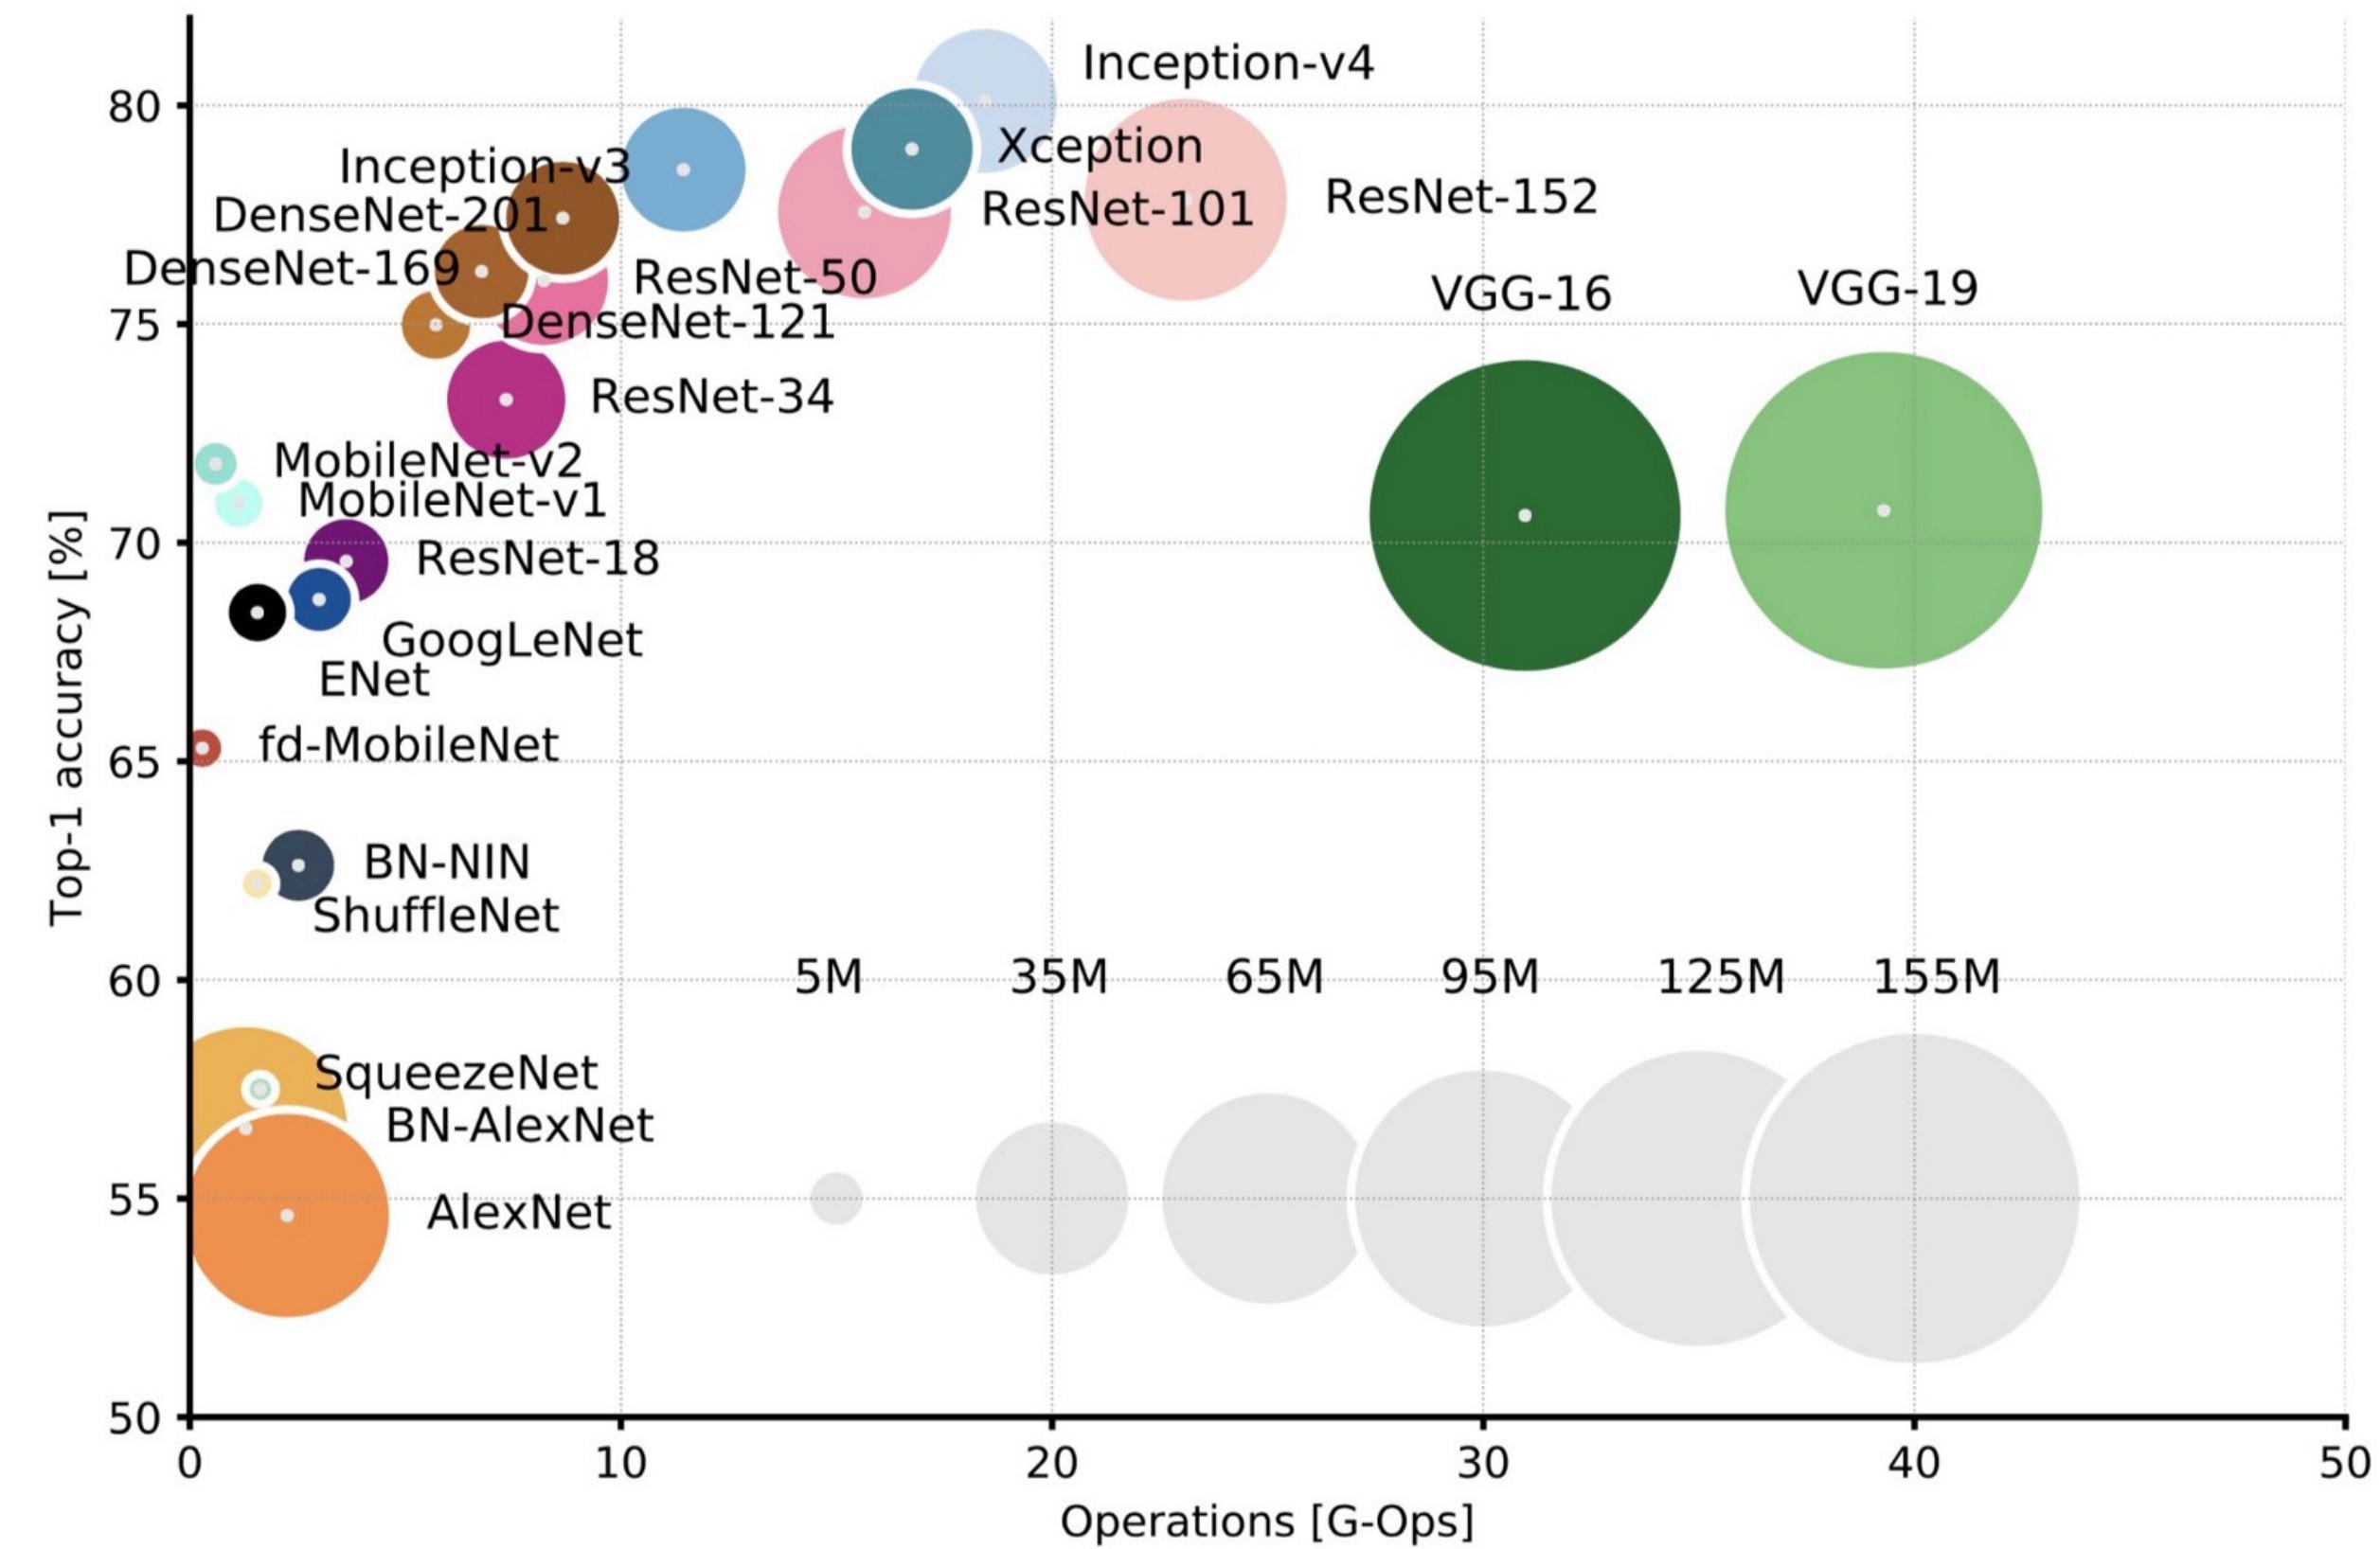
- [DenseNet \(2016\)](#): uses skip connections between a given layer and all the layers downstream in a dense block, with Conv layers in between



The computational cost

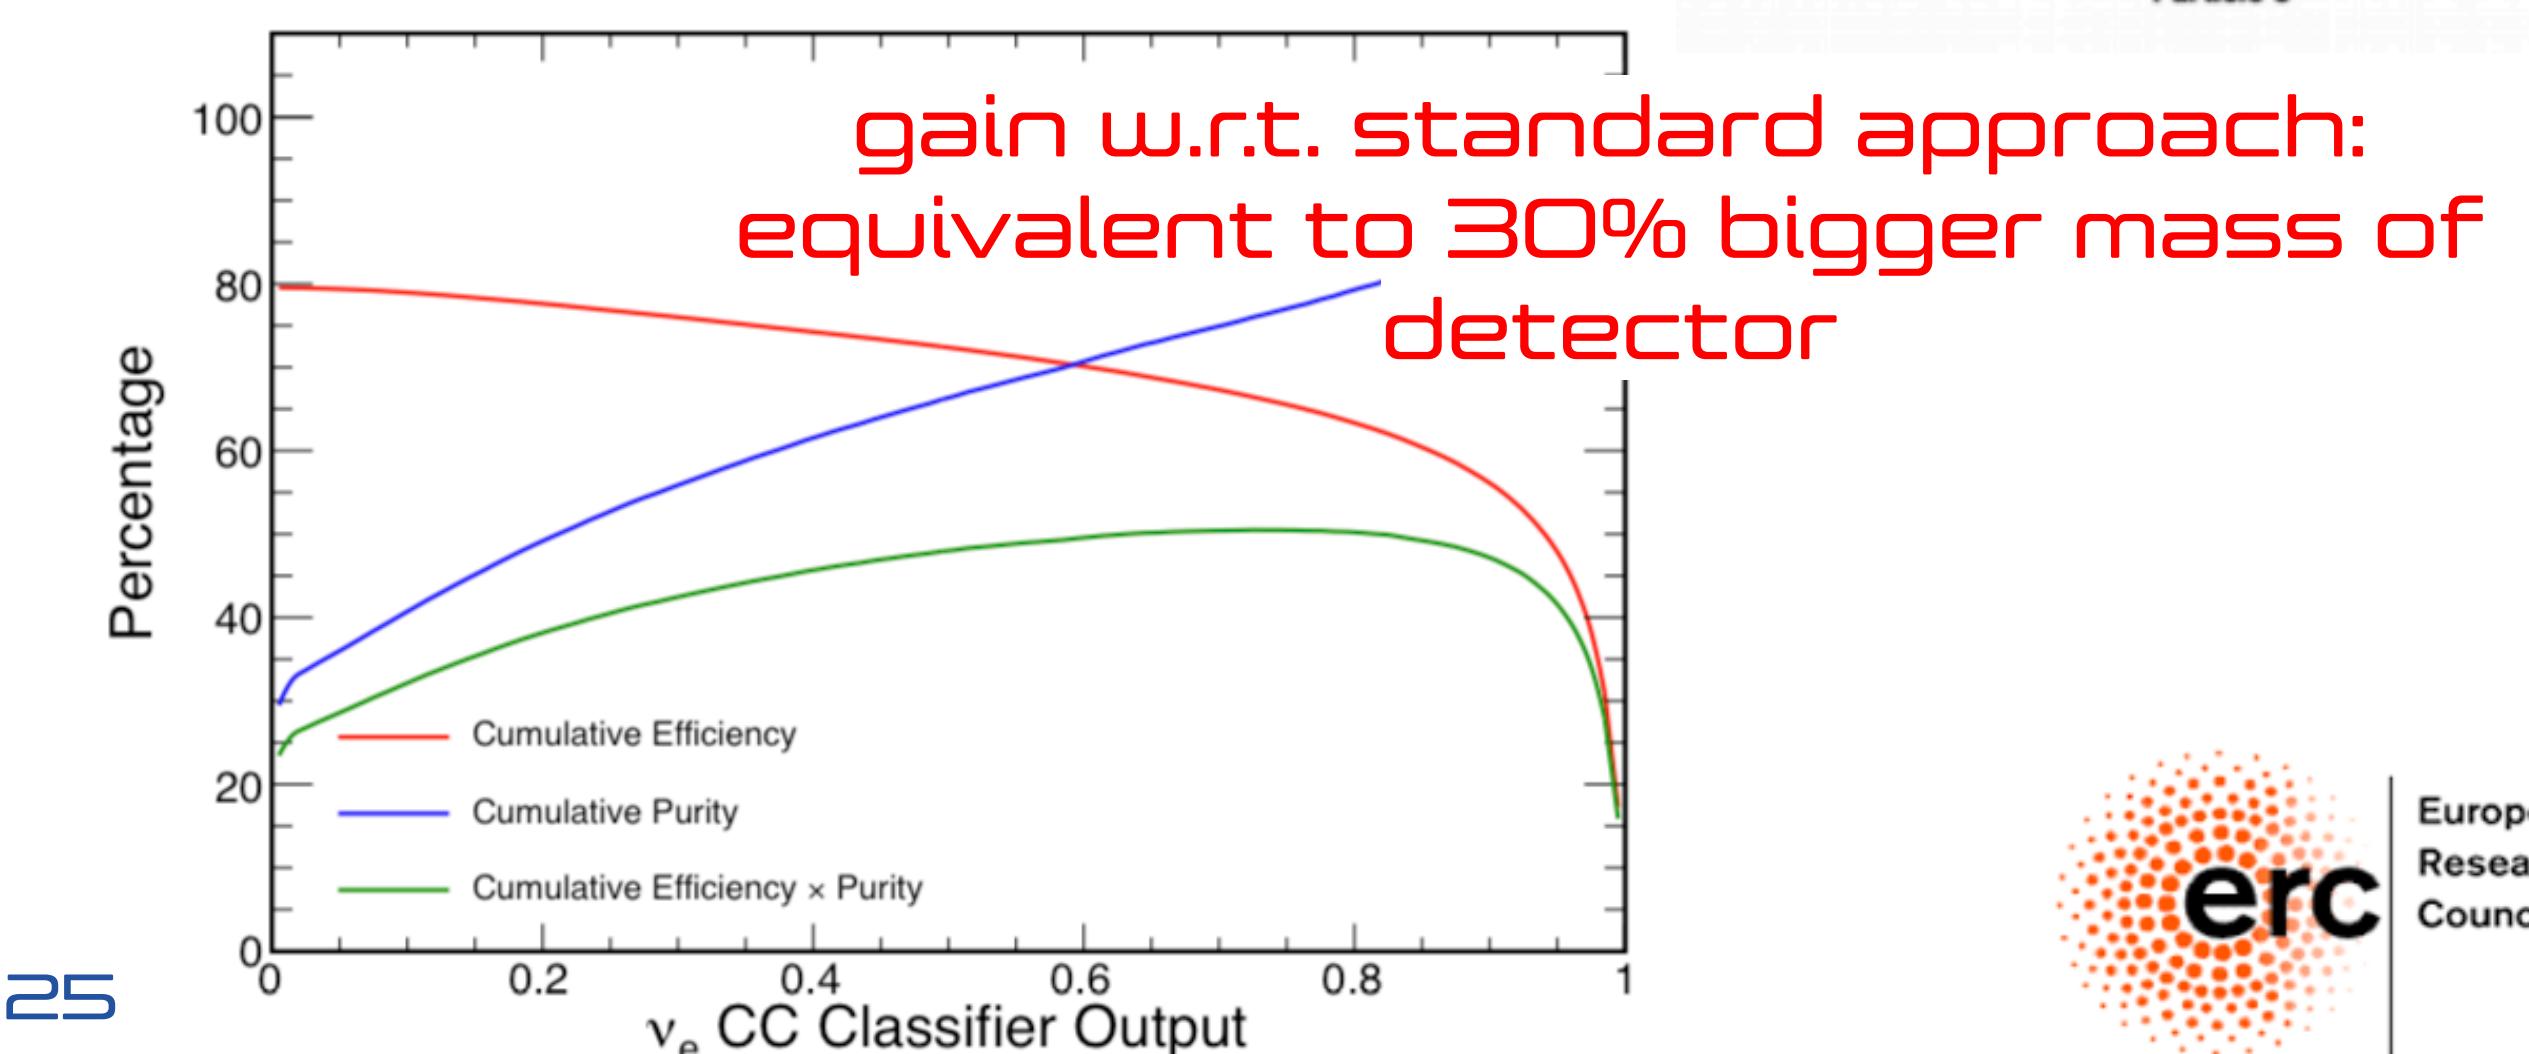
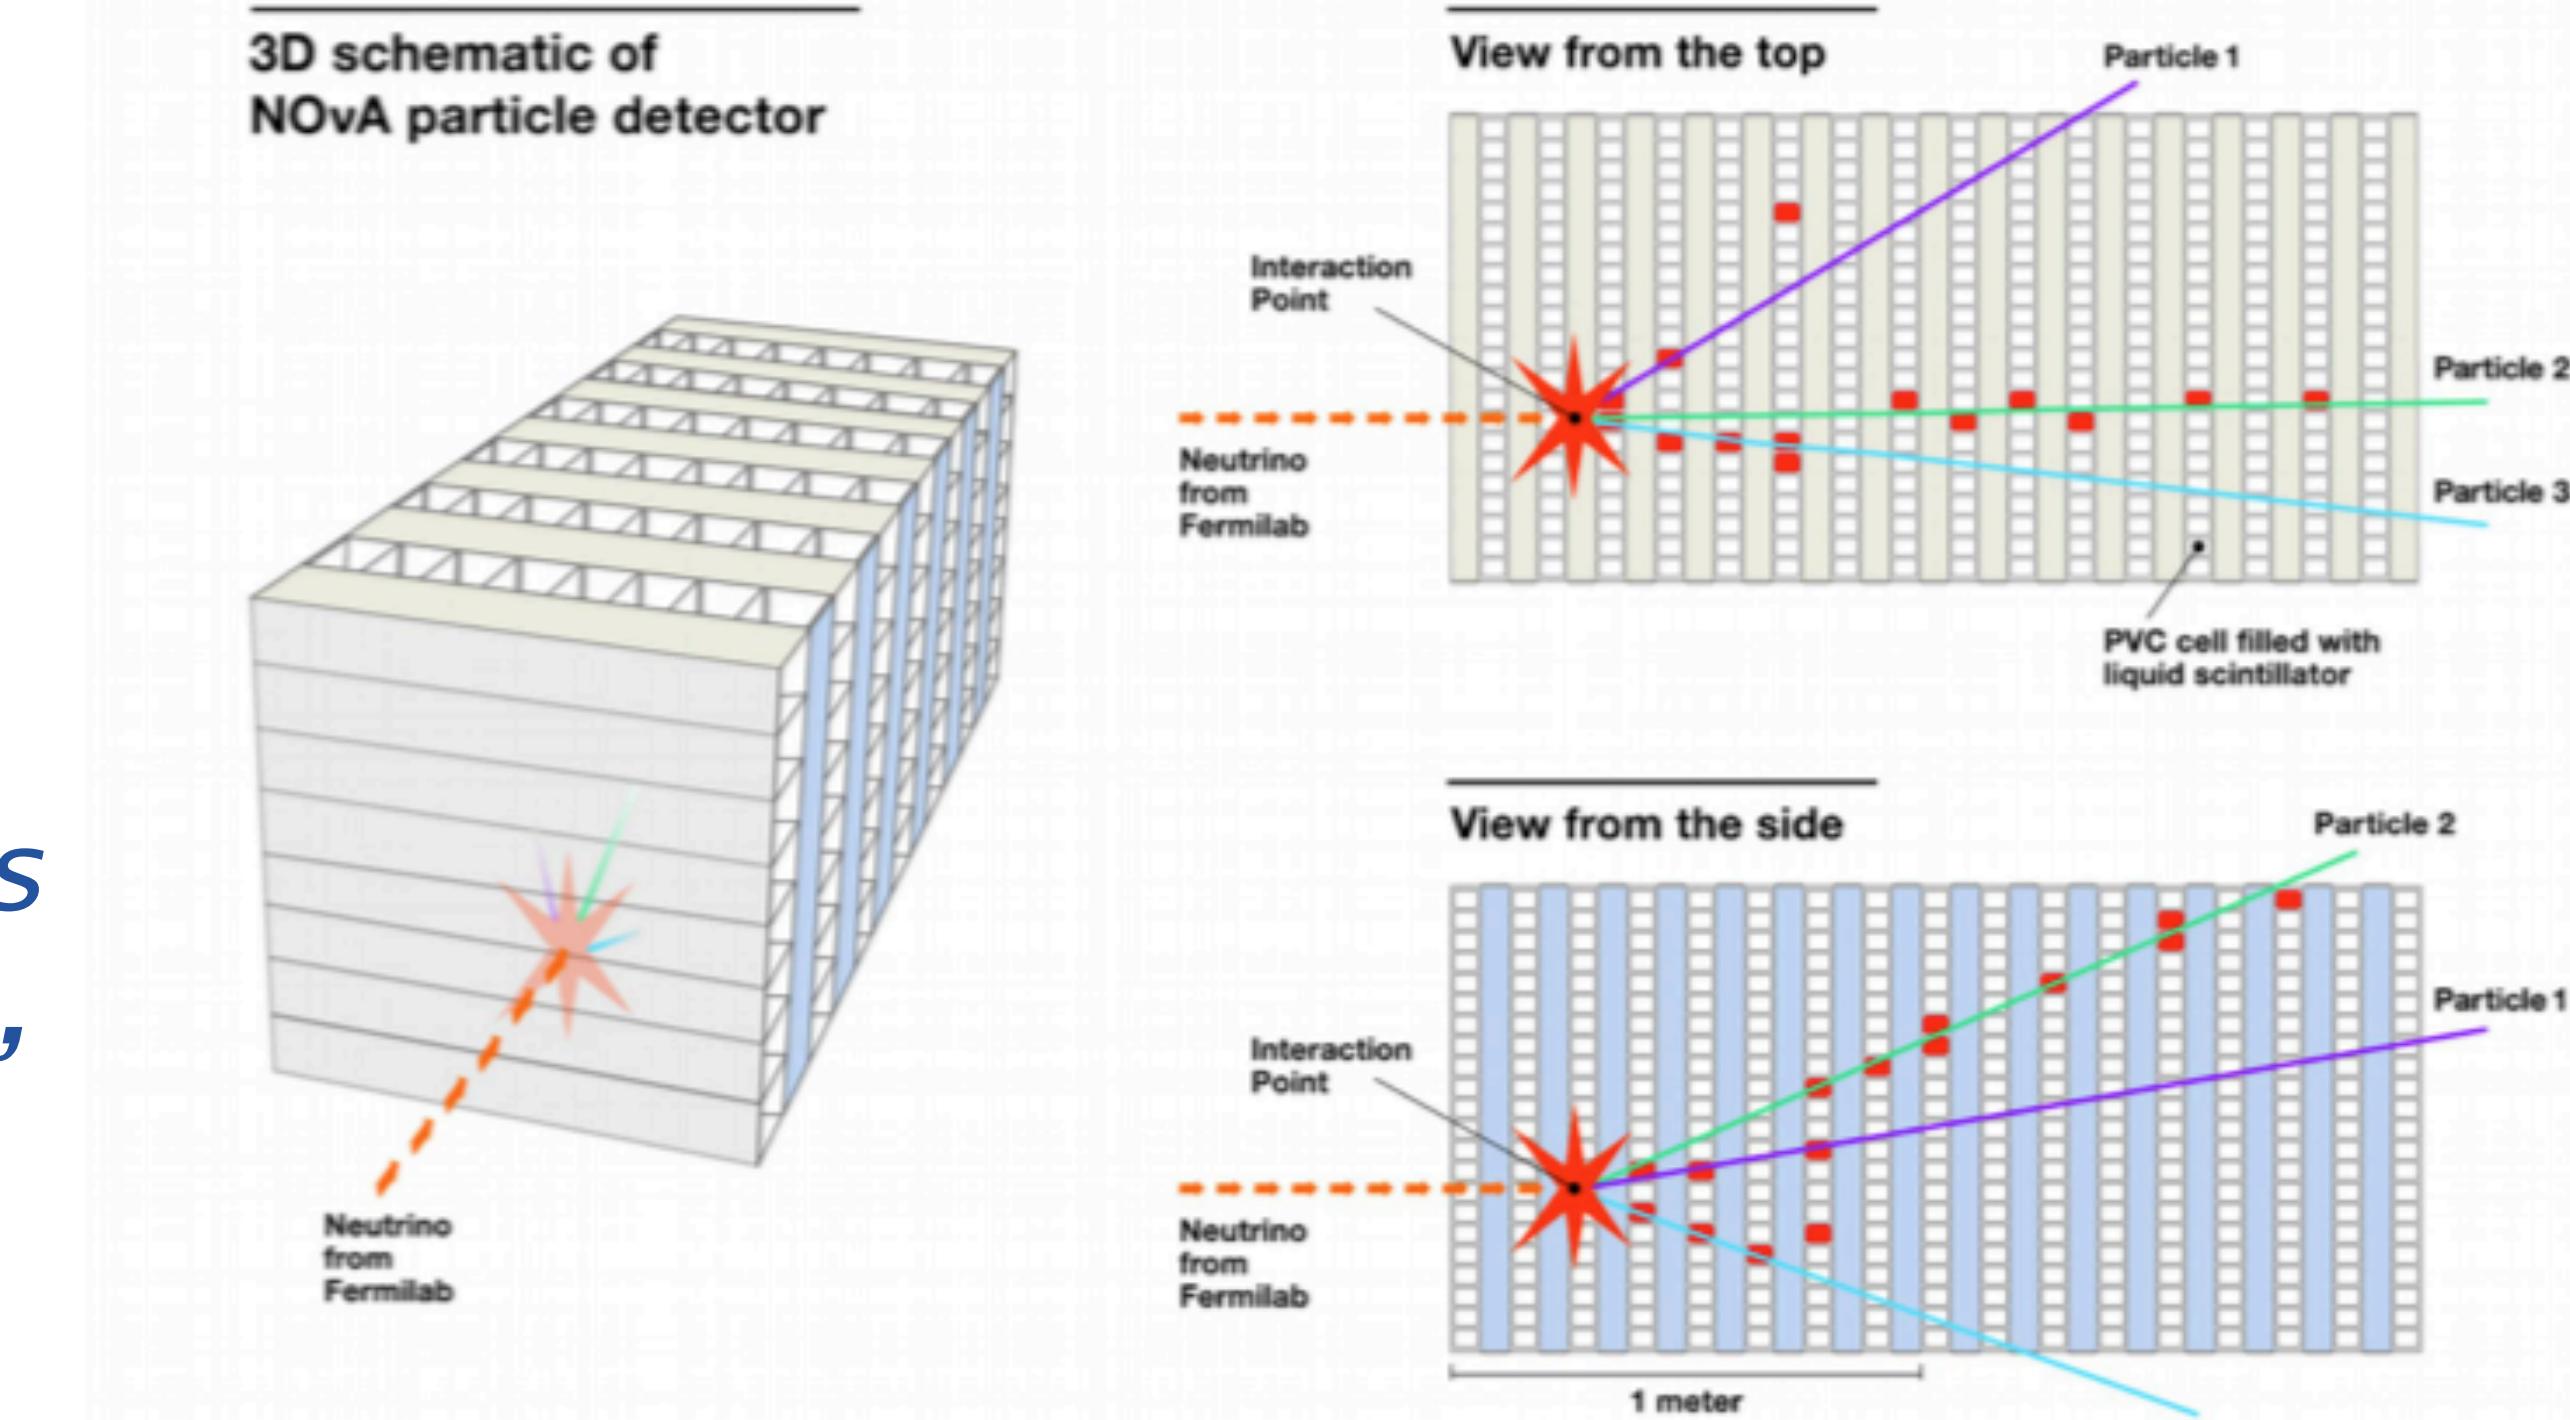
- In this evolution, computing-vision networks drastically improved in performance
- This came at a big cost in complexity (number of parameters and operations)
- The inference with this network became particularly slow
 - big interest in optimize these networks on dedicated resources, e.g. FPGAs
- Many cloud providers provide optimized versions of these networks:
 - you can just use them (re-training if needed), rather than inventing your own one

<https://towardsdatascience.com/neural-network-architectures-156e5bad51ba>



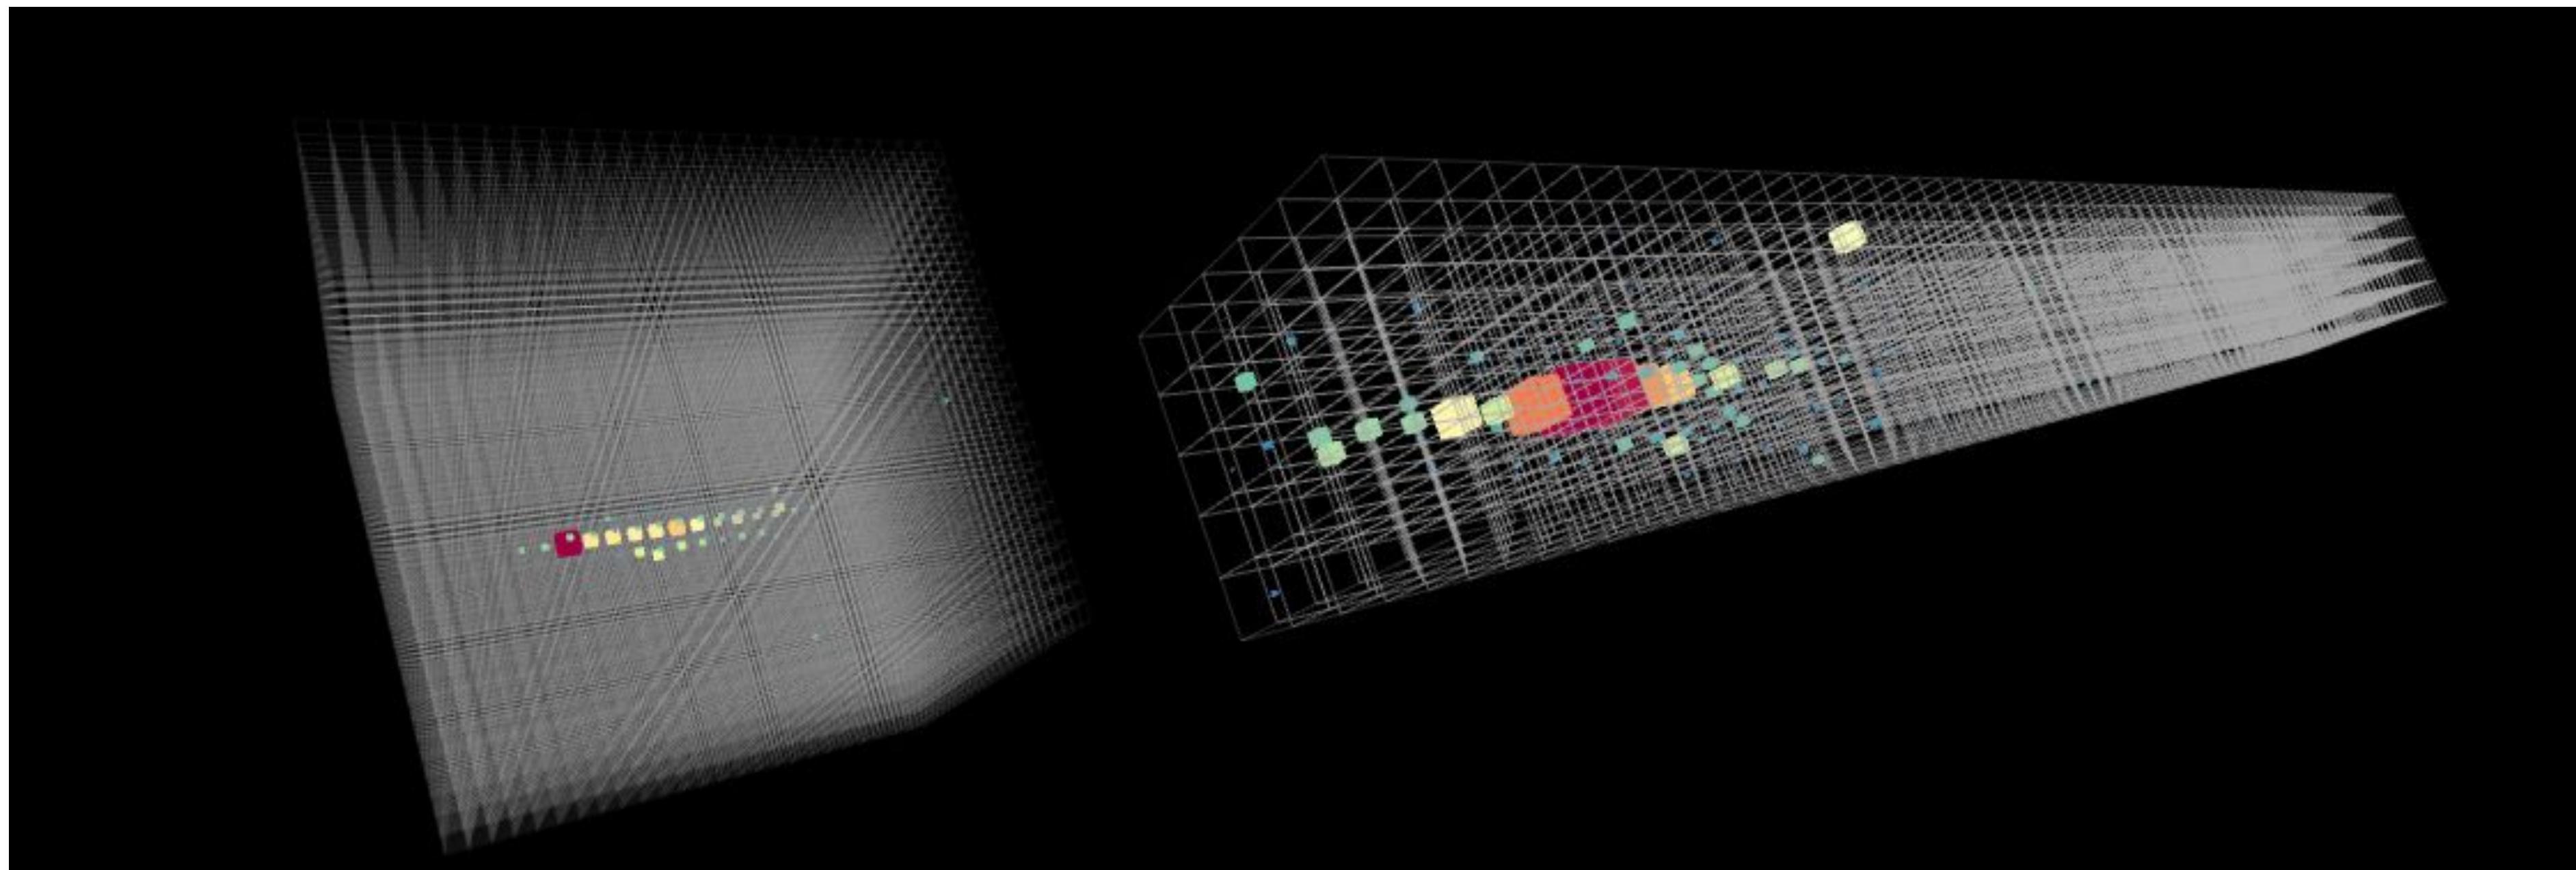
Example: PID for ν experiments

- Many HEP detectors (particularly underground) more and more structured as regular arrays of sensors
- Modern computer-vision techniques work with images as arrays of pixel sensor (in 1D, 2D, and 3D)
- These techniques were applied by Nova on electron and muon ID
- Impressive gain over traditional techniques (comparable to +30% detector == \$\$\$ saved)



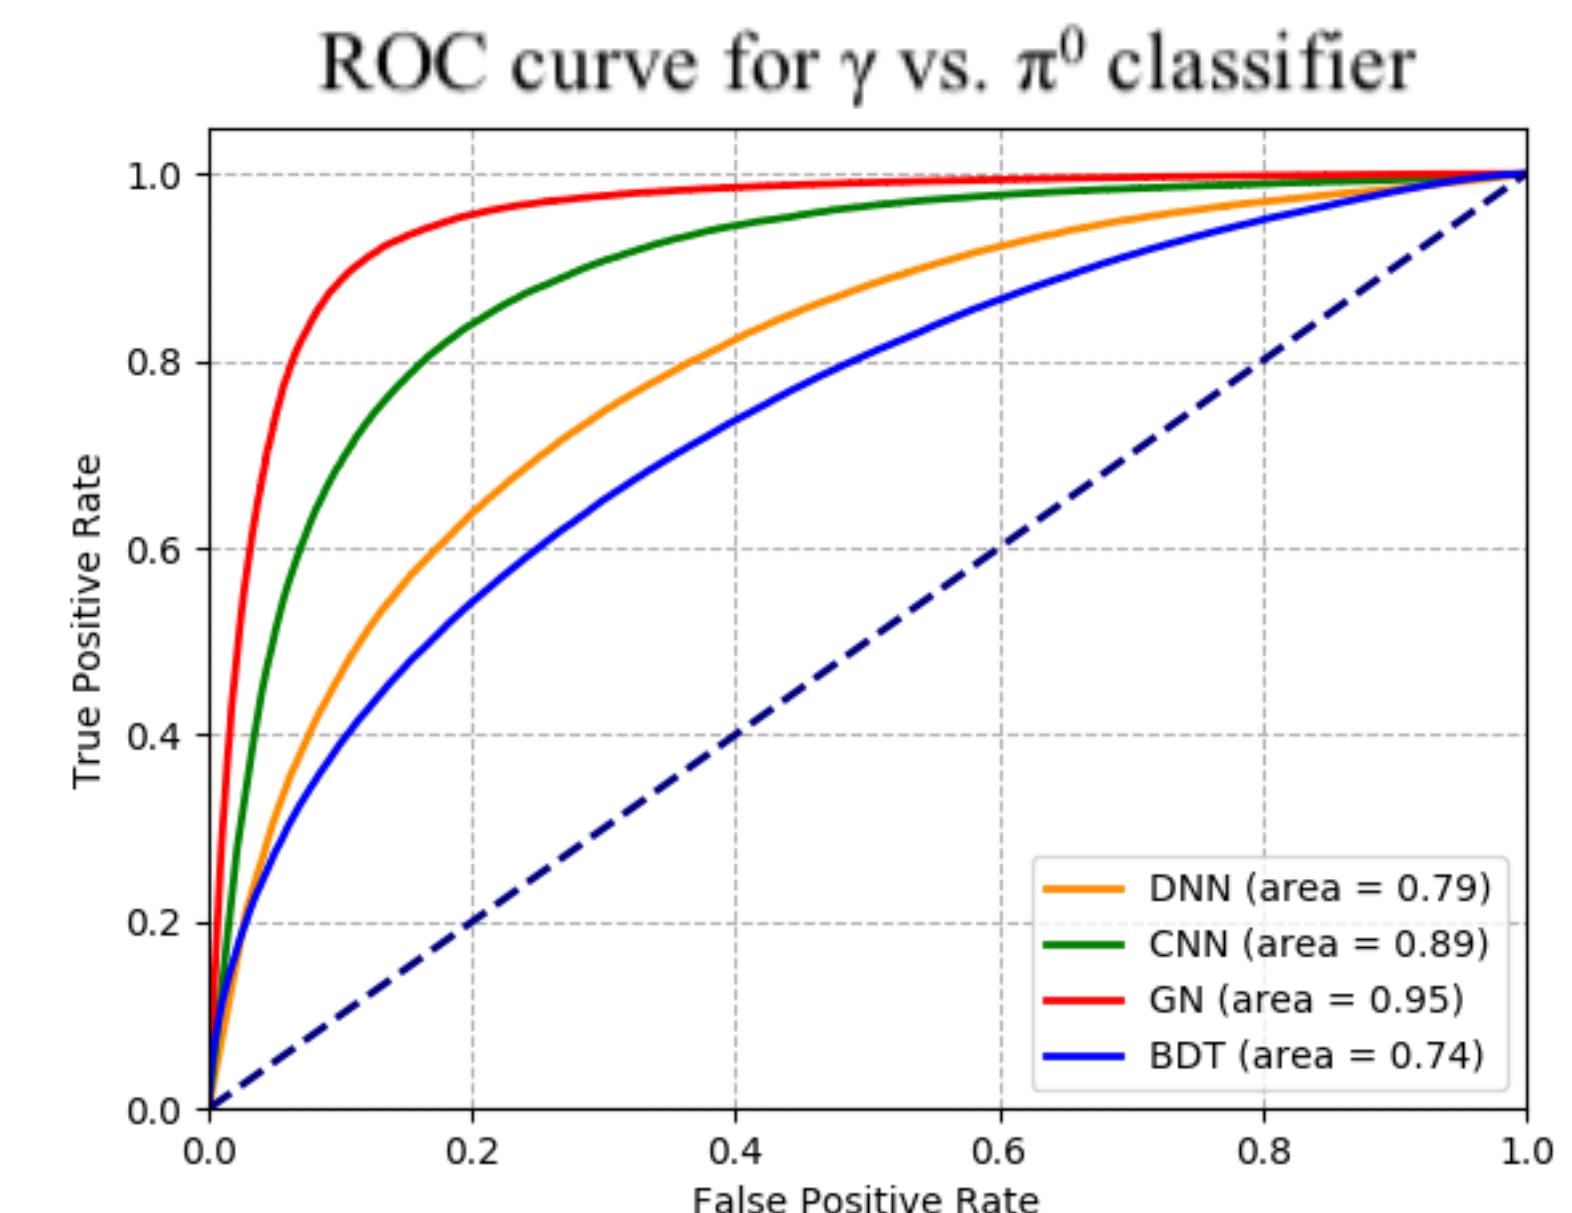
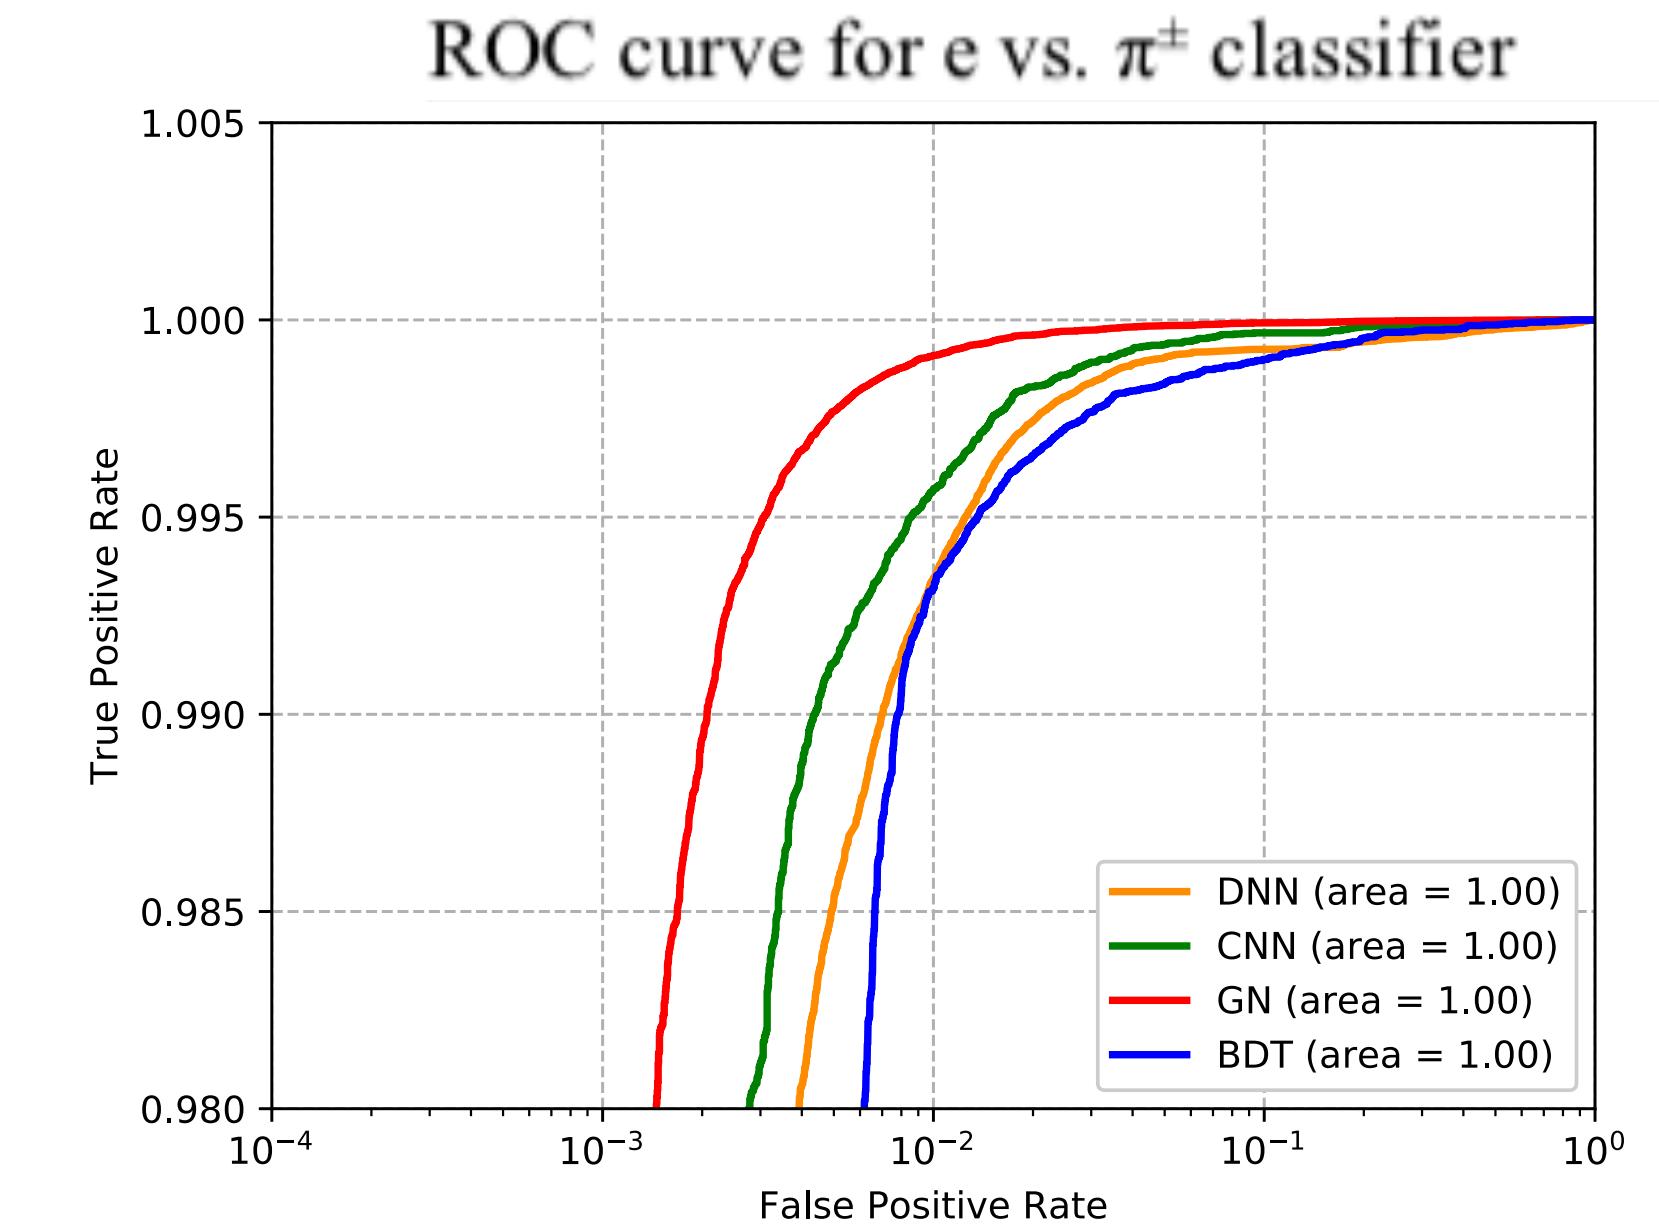
Calorimetry & Computer Vision

- (next generation) digital calorimeters: 3D arrays of sensors with more regular geometry
- Ideal configuration to apply Convolutional Neural Network
- speed up reconstruction at similar performances
- and possibly improve performances



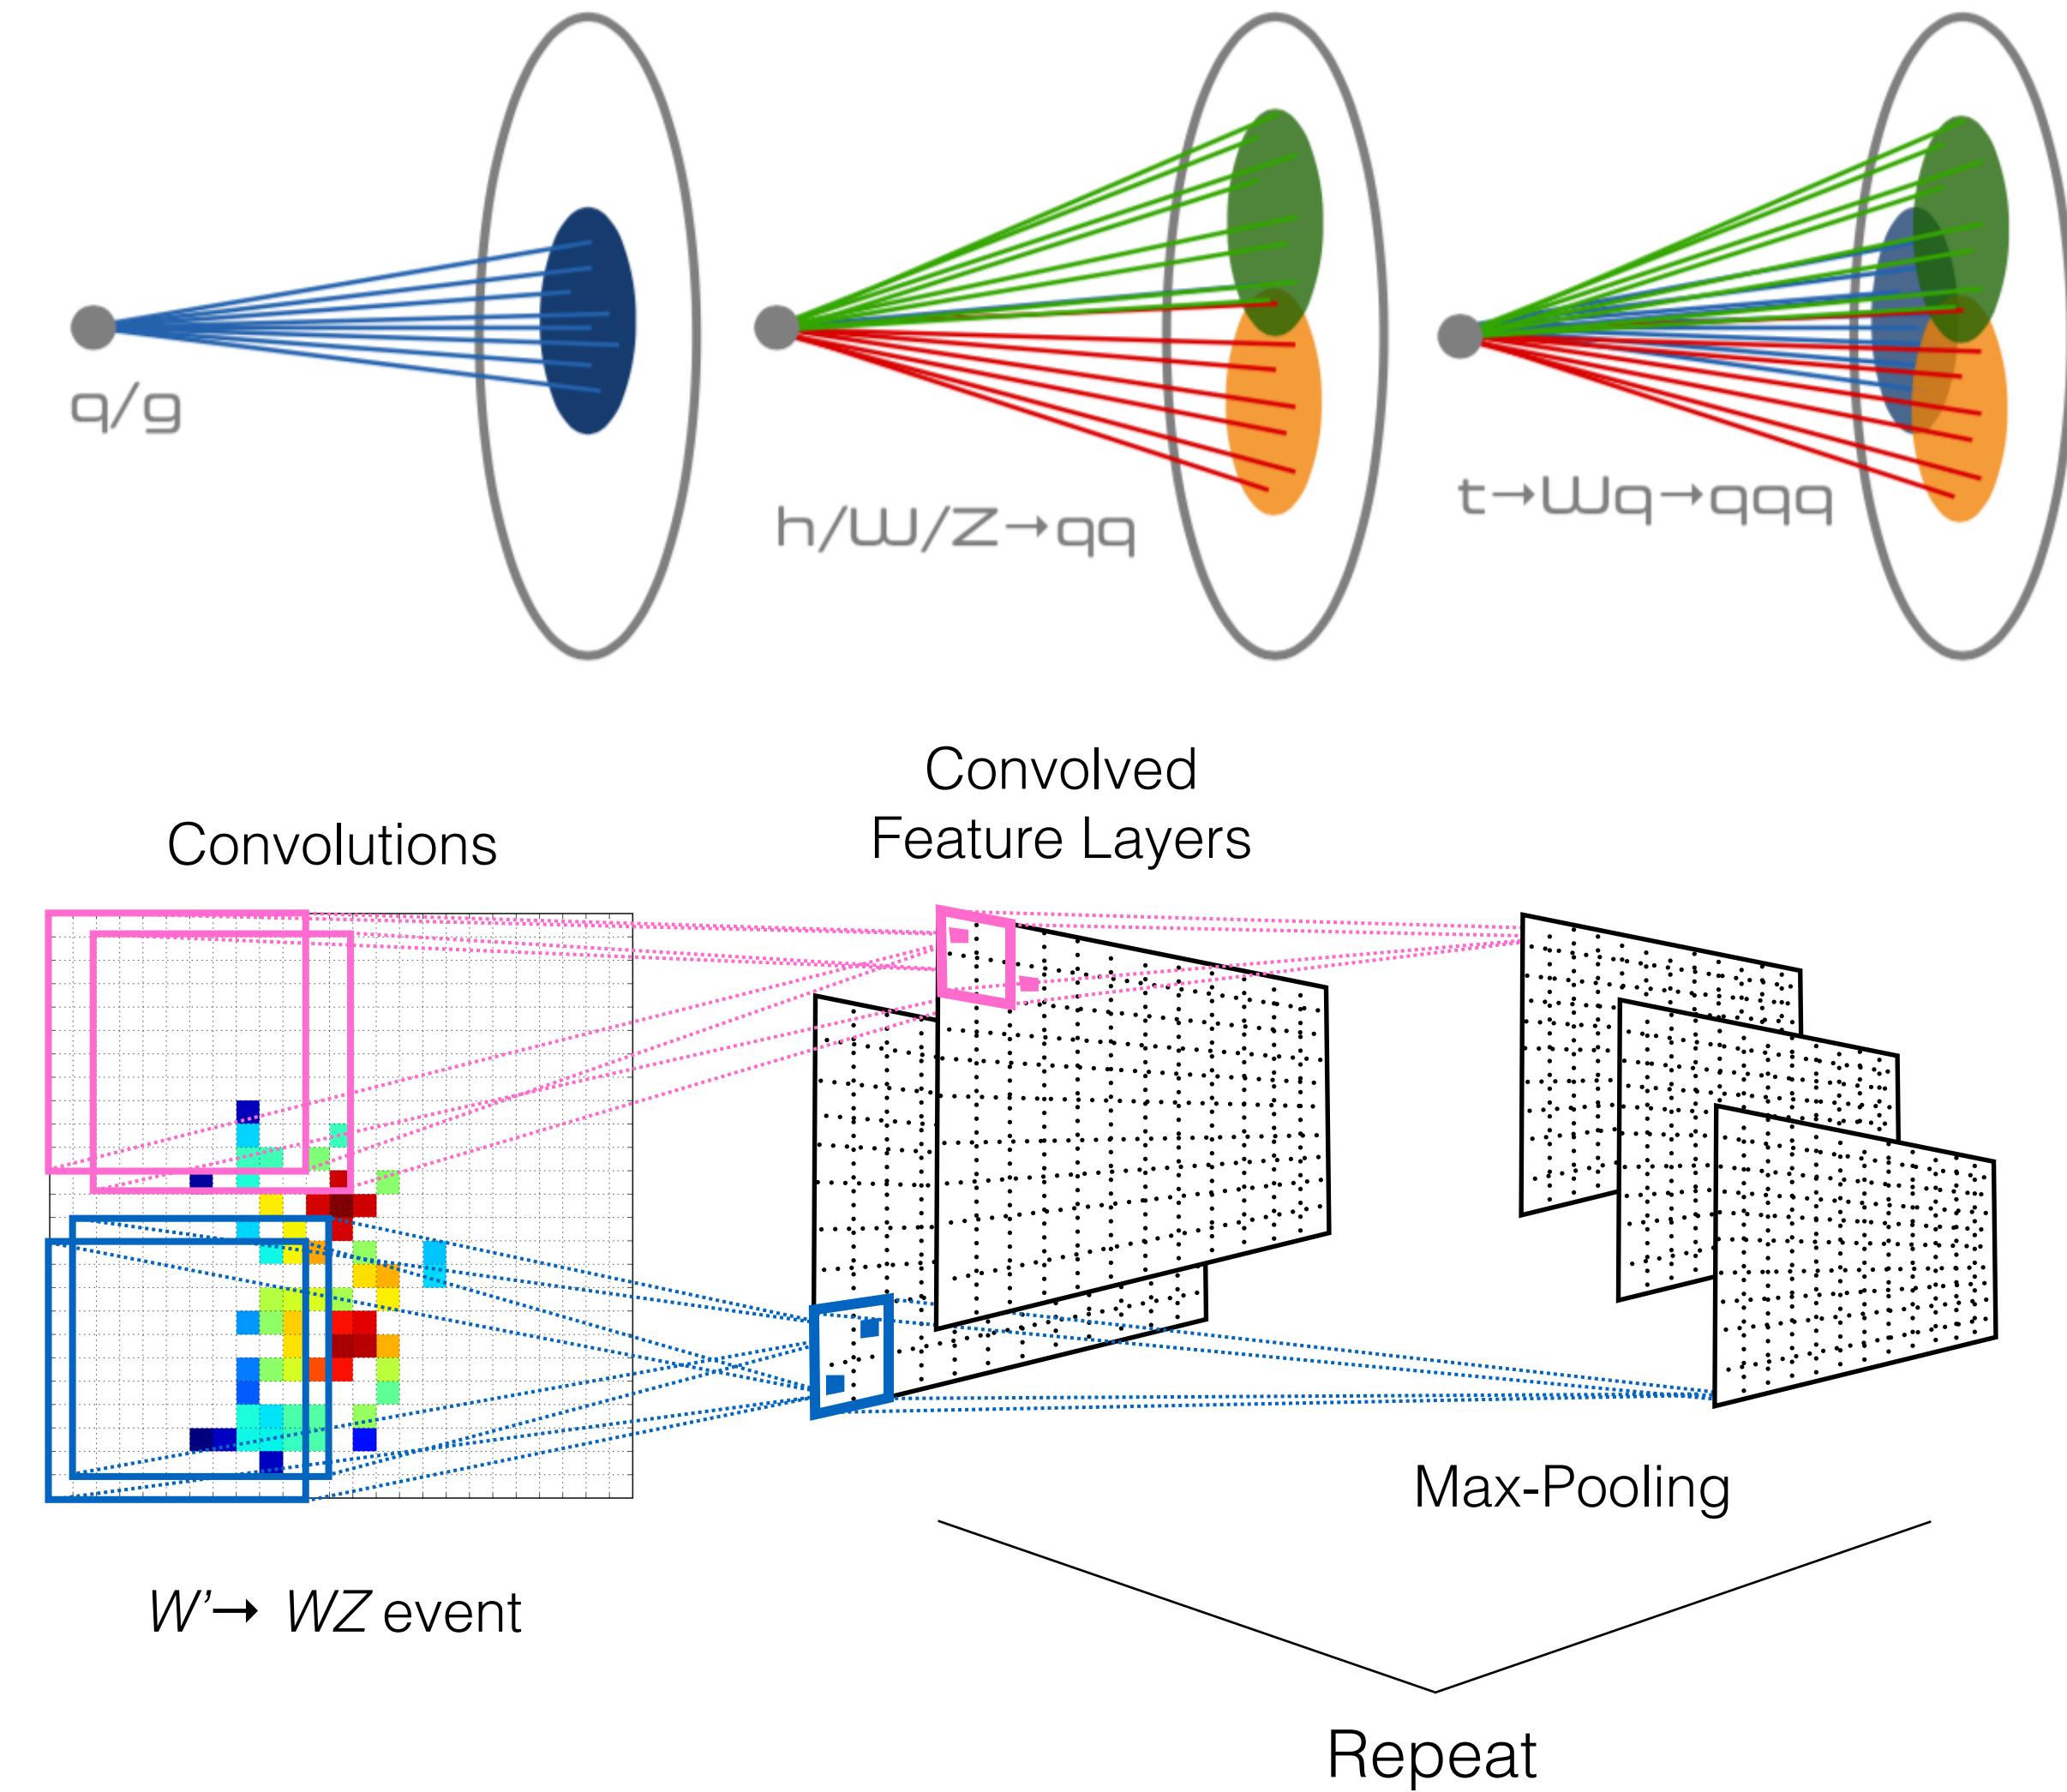
Example: Particle ID

- We tried particle ID on a sample of simulated events
 - one particle/event (e , γ , π^0 , π)
 - Different event representations
 - high-level features related to event shape (moments of X, Y , and Z projections, etc)
 - raw data (energy recorded in each cell)
 - Pre-filtered pion events to select the nasty ones and make the problem harder



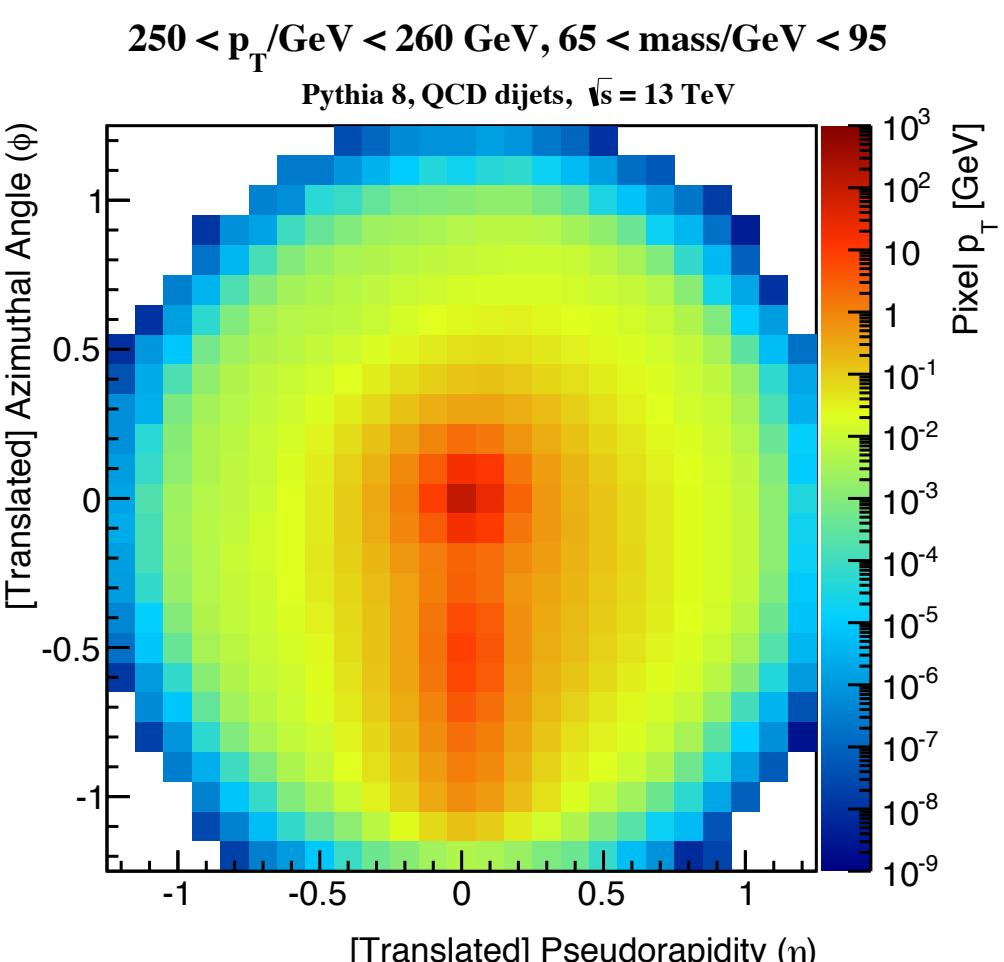
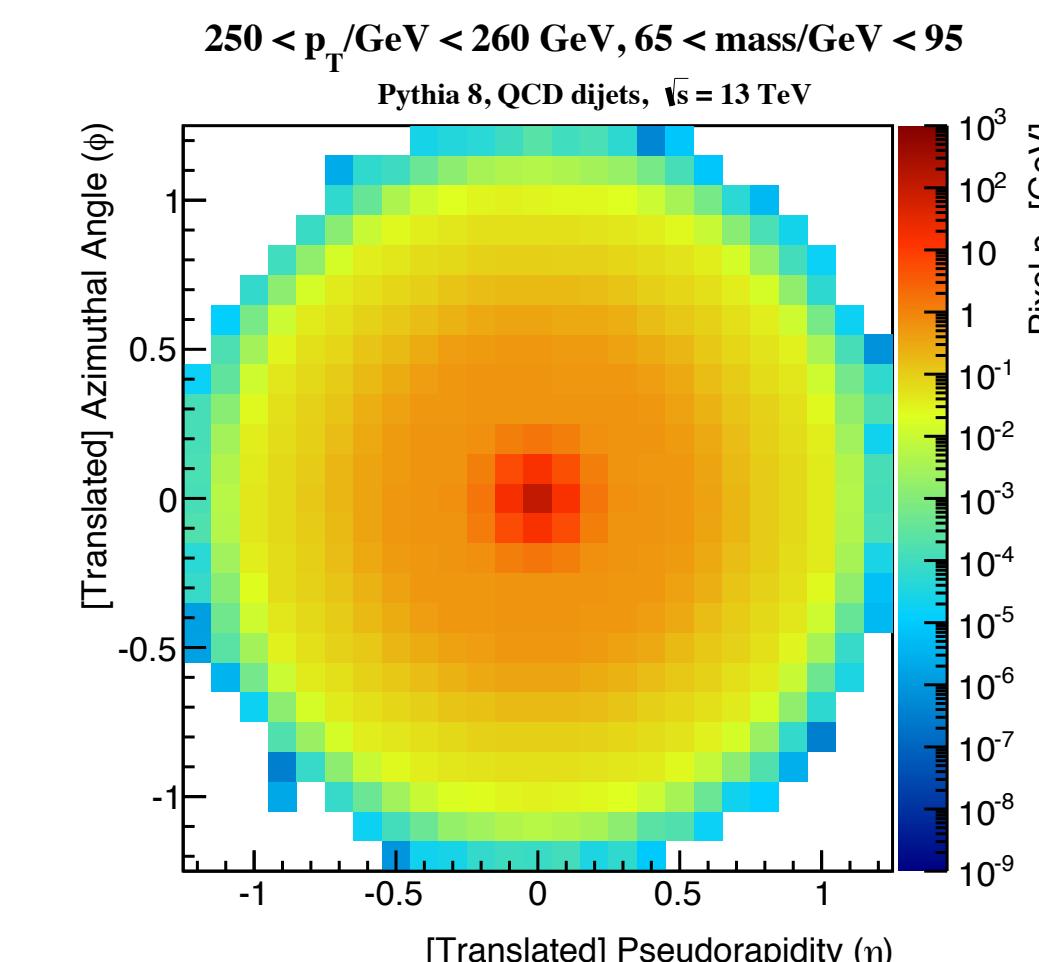
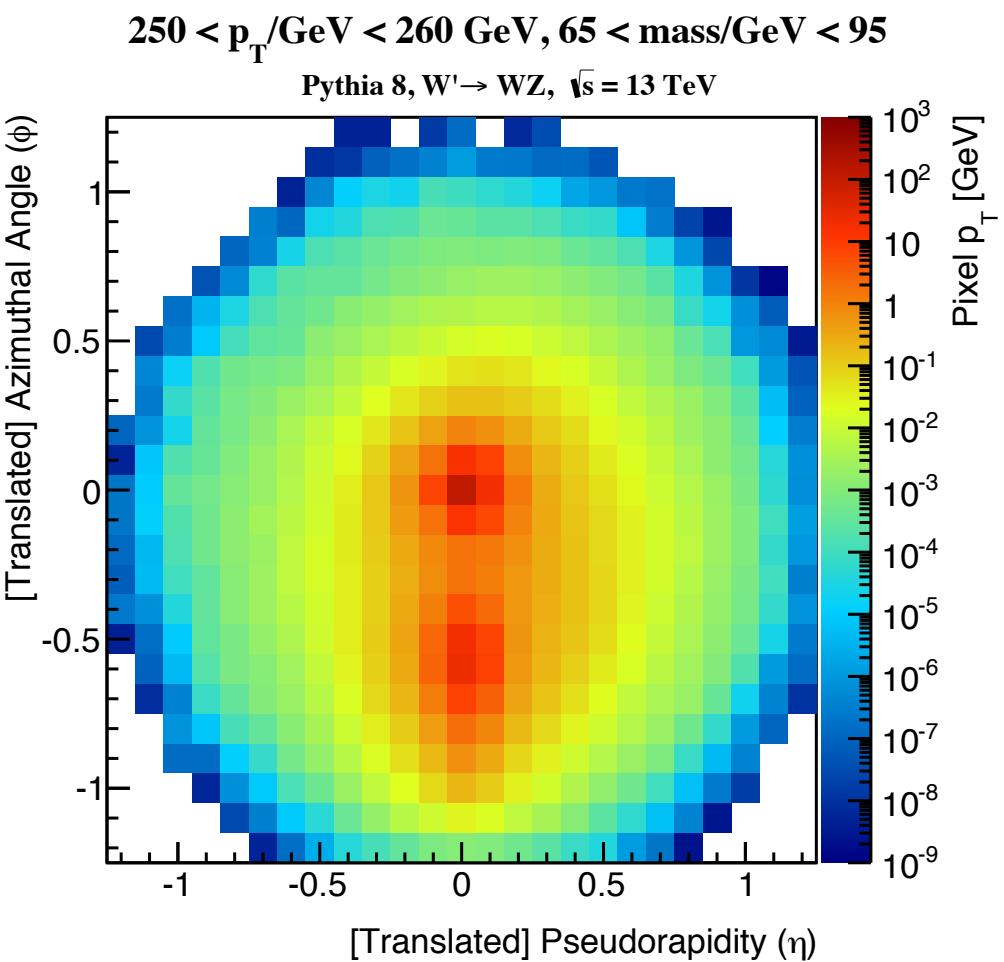
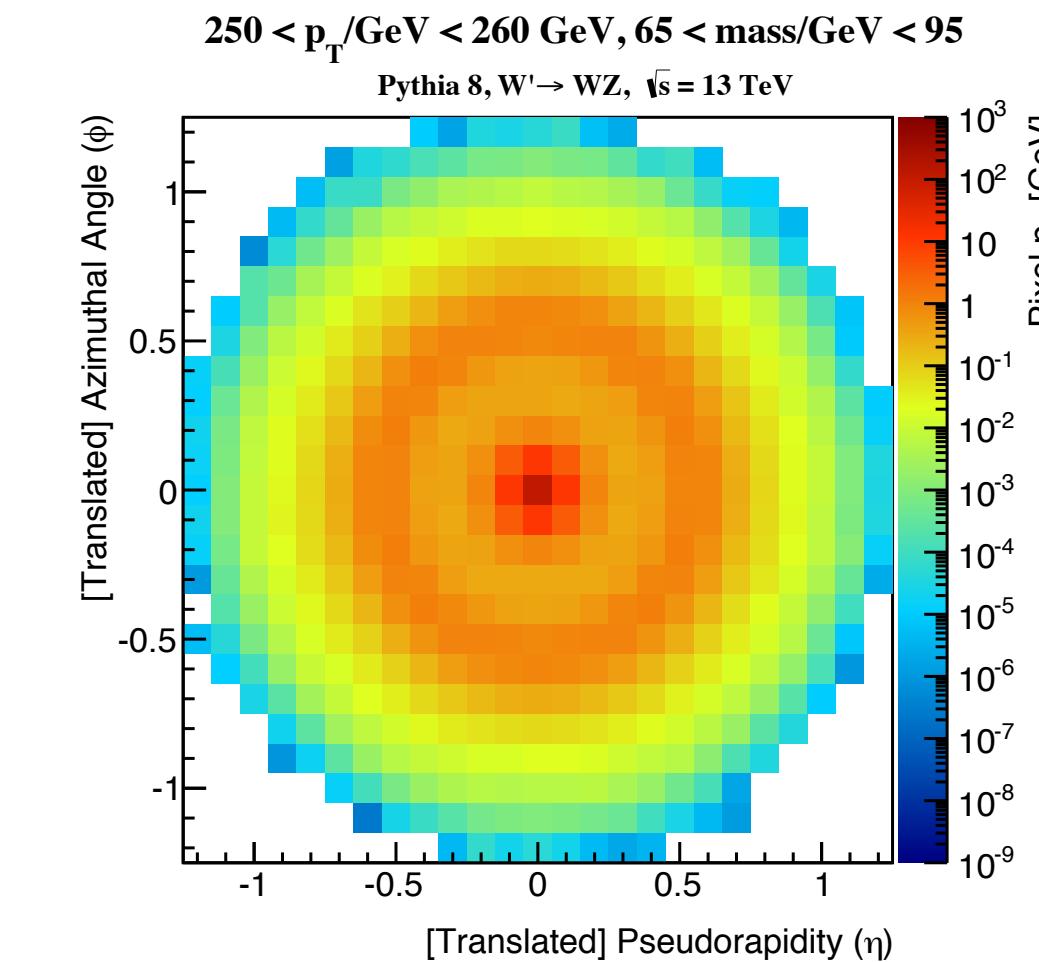
Jet as images (for ConvNN)

- One can pixelate the surface crossed by the jet and create an image with the momentum deposited in each cell
- Such an image can then be processed with computing-vision techniques
- Pros: can benefit of the progresses made in optimizing computing vision
- Cons: underlying assumption on detector geometry (regular array of pixels) made sacrificing information of the actual detector



Jet as images (for ConvNN)

- One can pixelate the surface crossed by the jet and create an image with the momentum deposited in each cell
- Such an image can then be processed with computing-vision techniques
- Pros: can benefit of the progresses made in optimizing computing vision
- Cons: underlying assumption on detector geometry (regular array of pixels) made sacrificing information of the actual detector



Summary

- Conv Networks are the most striking example of the success of Deep Learning
- A story of improvement from the 90s to now
- Very effective as a fast tool for particle reconstruction
- Same geometry can be used for image generation with adversarial training
- Can speed up one of the heaviest tasks in particle physics